

Performance

Admin:

Lab this week: discuss projects.

Ordering parts: keep all receipts! If paying with a credit card, you must have a receipt that shows at least the last four digits of the card.

If you need a particular part, ask and we can either order it for you or see if we have it already (e.g., ADC and DAC chips).



Falcon 9

Elon Musk





**Why is performance
optimization important?**

```
Welcome to the CS107E shell. Remember to  
type on your PS/2 keyboard!  
Pi> hi slow shell r
```

User experience matters

BIG DATA





Demo: clear

SW Optimization Strategies

- **Compiler flags** - leverage automatic optimizations
 - `gcc -O0 [source files] [object files] -o output file`
 - Common options: -O0, -O1, -O2, -O3
 - Other options: -Ofast, -Og, -Os
 - Ofast May produce inaccurate math results


```

draw_pixel:
    8010:    f8 40 2d e9    push    {r3, r4, r5, r6,
r7, lr}
    8014:    00 40 a0 e1    mov     r4, r0
    8018:    01 50 a0 e1    mov     r5, r1
    801c:    02 70 a0 e1    mov     r7, r2
    8020:    c9 01 00 eb    bl     #1828 <fb_get_width>
    8024:    00 60 a0 e1    mov     r6, r0
    8028:    c7 01 00 eb    bl     #1820 <fb_get_width>
    802c:    28 30 9f e5    ldr     r3, [pc, #40]
    8030:    00 30 93 e5    ldr     r3, [r3]
    8034:    97    <unknown>
    8035:    06 02 e0 02    rsceq   r0, r0,

...
    8050:    03 20 c3 e5    strb    r2, [r3, #3]
    8054:    f8 40 bd e8    pop     {r3, r4, r5, r6, r7,
lr}

```

draw_pixel:

8010:	f8 40 2d e9	push	{r3, r4, r5, r6, r7, lr}
8014:	00 40 a0 e1	mov	r4, r0
8018:	01 50 a0 e1	mov	r5, r1
801c:	02 70 a0 e1	mov	r7, r2
8020:	c9 01 00 eb	bl	#1828 <fb_get_width>
8024:	00 60 a0 e1	mov	r6, r0
8028:	c7 01 00 eb	bl	#1820 <fb_get_width>
802c:	28 30 9f e5	ldr	r3, [pc, #40]
8030:	00 30 93 e5	ldr	r3, [r3]
8034:	97 <unknown>		
8035:	06 02 e0 02	rsceq	r0, r0,
...			
8050:	03 20 c3 e5	strb	r2, [r3, #3]
8054:	f8 40 bd e8	pop	{r3, r4, r5, r6, r7, lr}

Overhead for creating stack frame, saving regs

draw_pixel:

8010:	f8 40 2d e9	push	{r3, r4, r5, r6,
r7, lr}			
8014:	00 40 a0 e1	mov	r4, r0
8018:	01 50 a0 e1	mov	r5, r1
801c:	02 70 a0 e1	mov	r7, r2
8020:	c9 01 00 eb	b1	#1828 <fb_get_width>
8024:	00 60 a0 e1	mov	r6, r0
8028:	c7 01 00 eb	b1	#1820 <fb_get_width>
802c:	28 30 9f e5	ldr	r3, [pc, #40]
8030:	00 30 93 e5	ldr	r3, [r3]
8034:	97 <unknown>		
8035:	06 02 e0 02	rsceq	r0, r0,
...			
8050:	03 20 c3 e5	strb	r2, [r3, #3]
8054:	f8 40 bd e8	pop	{r3, r4, r5, r6, r7,
lr}			

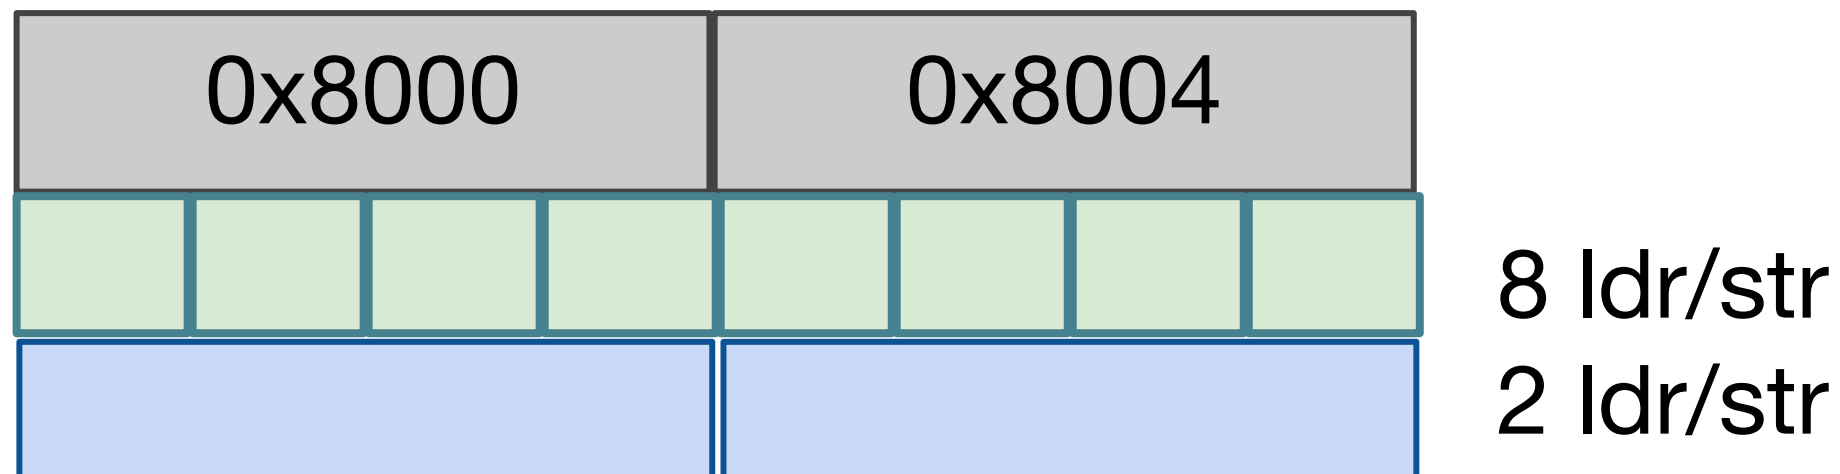
Overhead to call functions

SW Optimization Strategies

- **Compiler flags**
- **Inlining/Remove repetitive function calls**
 - Reduce function calls overhead
 - Can use the `inline` keyword to inline short helper function

SW Optimization Strategies

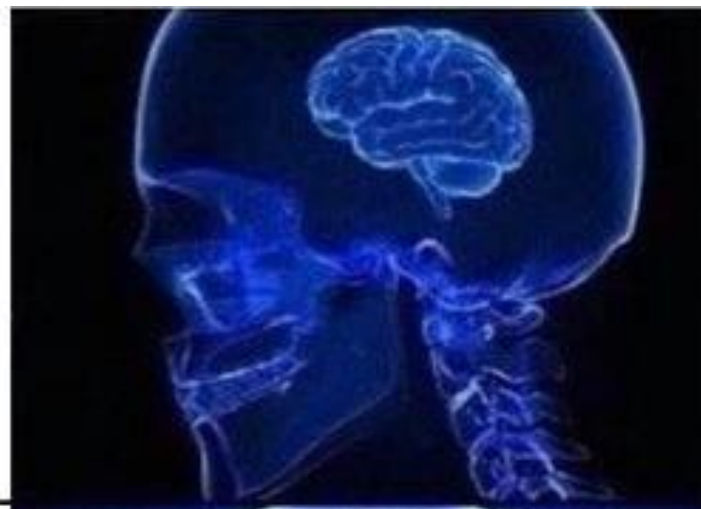
- **Compiler flags**
- **Inlining/Remove repetitive function calls**
- **Aggregate loads/stores**
 - can load/store a word once instead of 4 bytes



SW Optimization Strategies

- **Compiler flags**
- **Inlining/Remove repetitive function calls**
- **Aggregate loads/stores**
- **Loop optimization**
 - Code Hoisting: remove loop-invariant ops from loop body
 - Flatten nested loop
 - Loop unrolling: combine operations from multiple iteration

**OPTIMIZING
WITH
COMPILER FLAG**



**REMOVE
FUNCTION
CALL OVERHEAD**



**LOOP
UNROLLING**



**WRITE YOUR
OWN ASM HELPER**

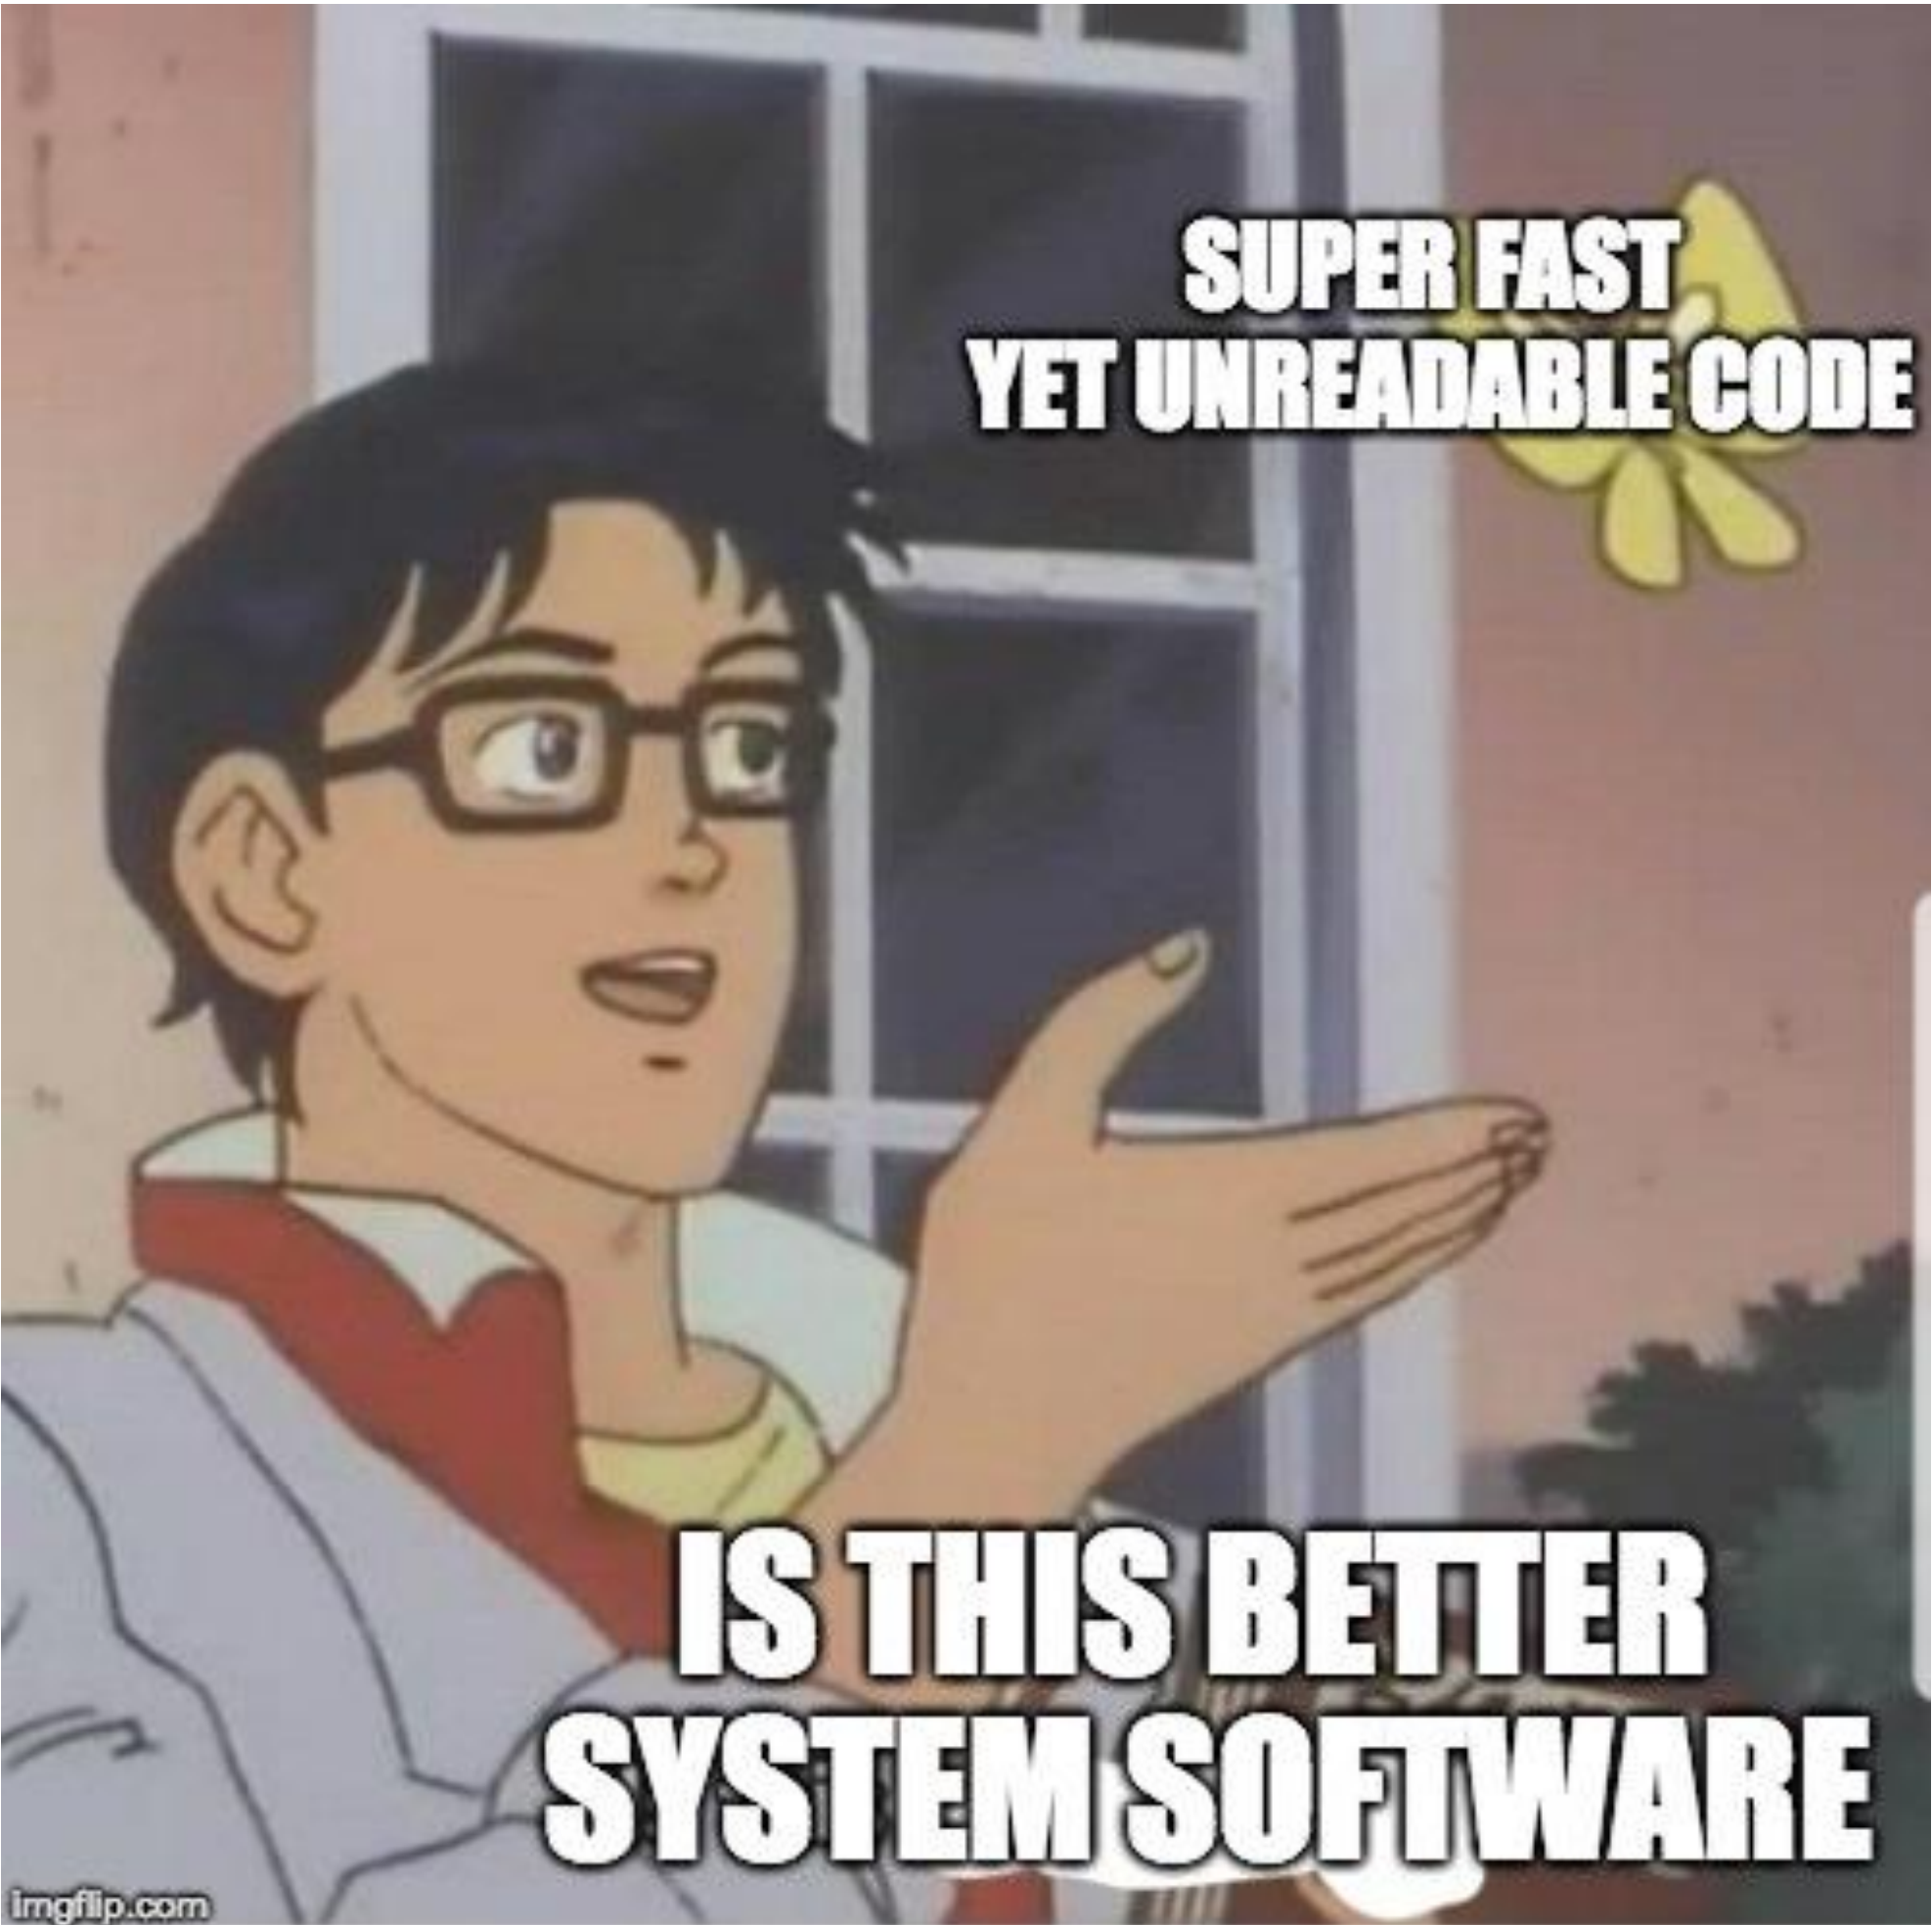


SW Optimization Strategies

- **Compiler flags**
- **Inlining/Remove repetitive function calls**
- **Aggregate loads/stores**
- **Loop optimization**
- **Manual Assembly**

SW Optimization Strategies

- **Compiler flags** - leverage automatic optimizations
- **Inlining** - Reduce function call overhead
- **Aggregate loads/stores** - less instructions per memory operation
- **Loop optimizations** - code hoisting, combination, unrolling
- **Manual assembly** - Be smarter than the compiler



**SUPER FAST
YET UNREADABLE CODE**

**IS THIS BETTER
SYSTEM SOFTWARE**

When to optimize?

Amdahl's Law

$$S_{latency}(s) = \frac{1}{(1 - p) + \frac{p}{s}}$$

- $S_{latency}$ - the theoretical speedup of the execution of the entire program
- s - the speedup of the part of the program you're optimizing
- p - the proportion of execution time that the part the program you're optimizing originally occupied

Amdahl's Law



Suppose your original program took **t** cycles to execute

Amdahl's Law



The program is divided into two distinct portions:

- Part **A** takes 75% of the time
- Part **B** takes 25% of the time

Amdahl's Law



If we optimize part B to make it 5 times faster, this only reduces the overall computation time slightly

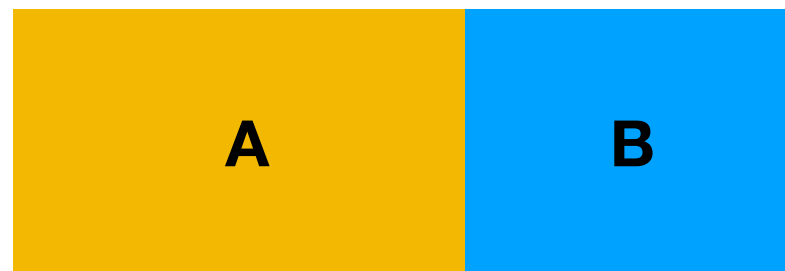
Amdahl's Law

Part B is 25% of the overall program ($p = .25$) and we speed it up by a factor of 5 ($s = 5$)

$$S_{latency} = \frac{1}{1 - .25 + \frac{.25}{5}} = 1.25$$

Overall program speedup is 1.25

Amdahl's Law



If we optimize part A to make it just twice as fast, we get a greater overall speedup

Amdahl's Law

Part A is 75% of the overall program ($p = .75$) and we speed it up by a factor of 2 ($s = 2$)

$$S_{latency} = \frac{1}{1 - .75 + \frac{.75}{2}} = 1.60$$

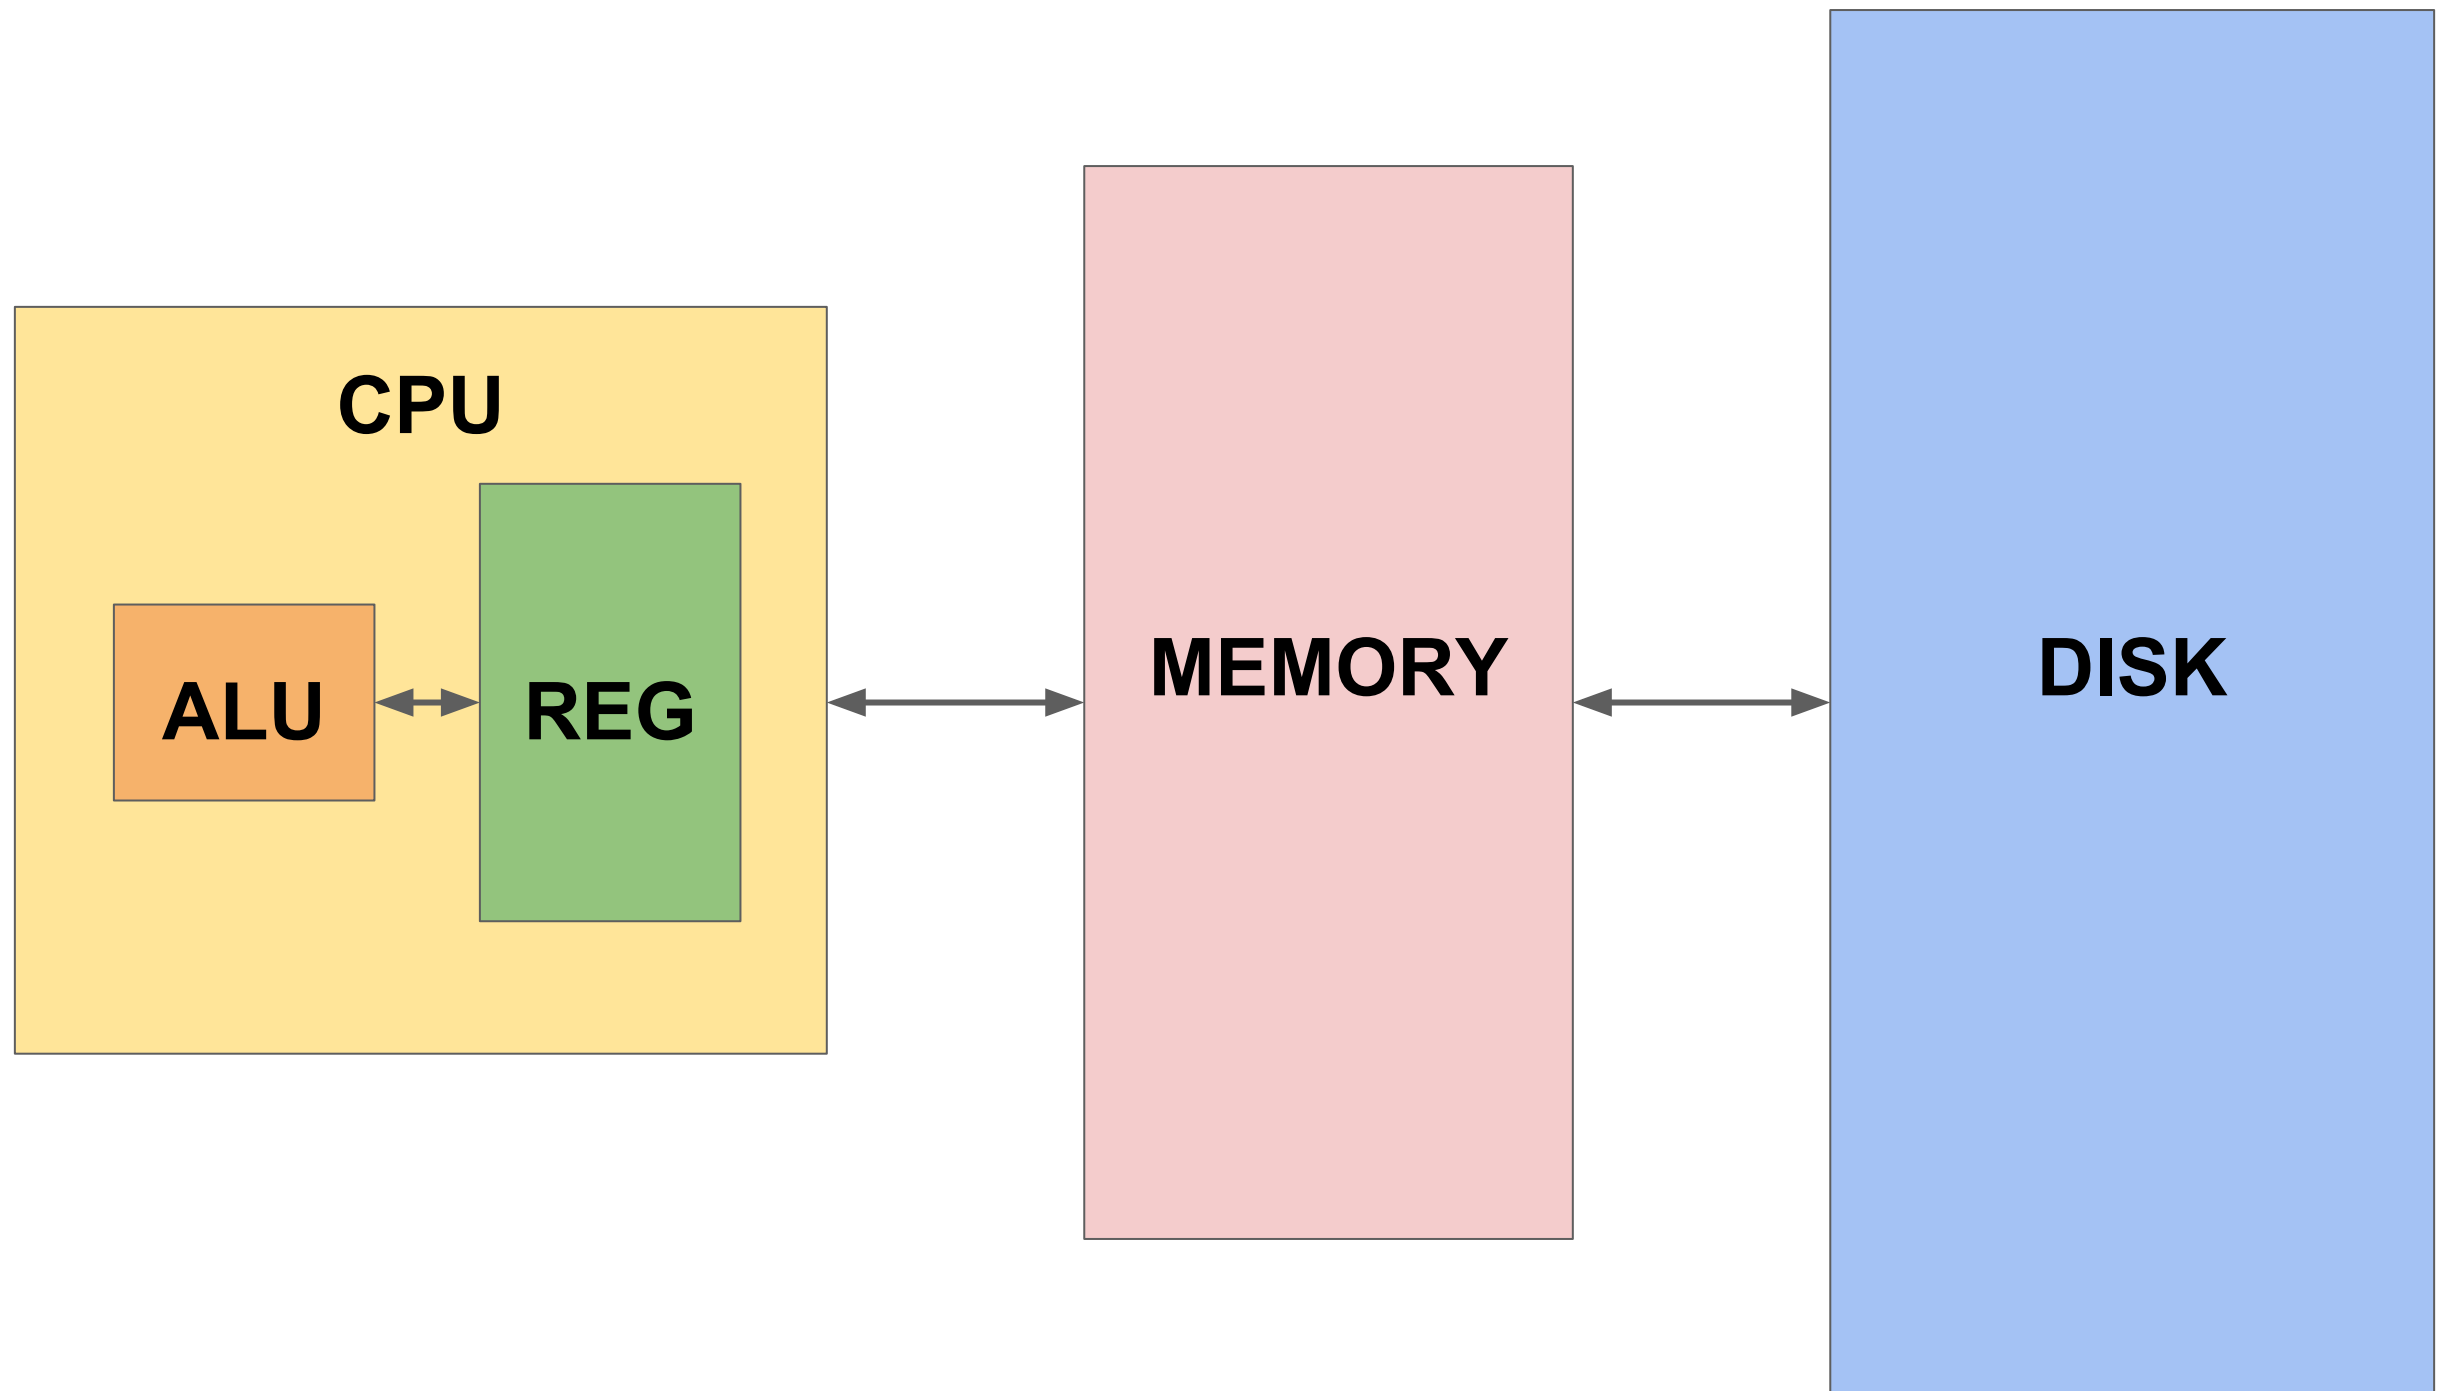
Overall program speedup is 1.60

Profiling

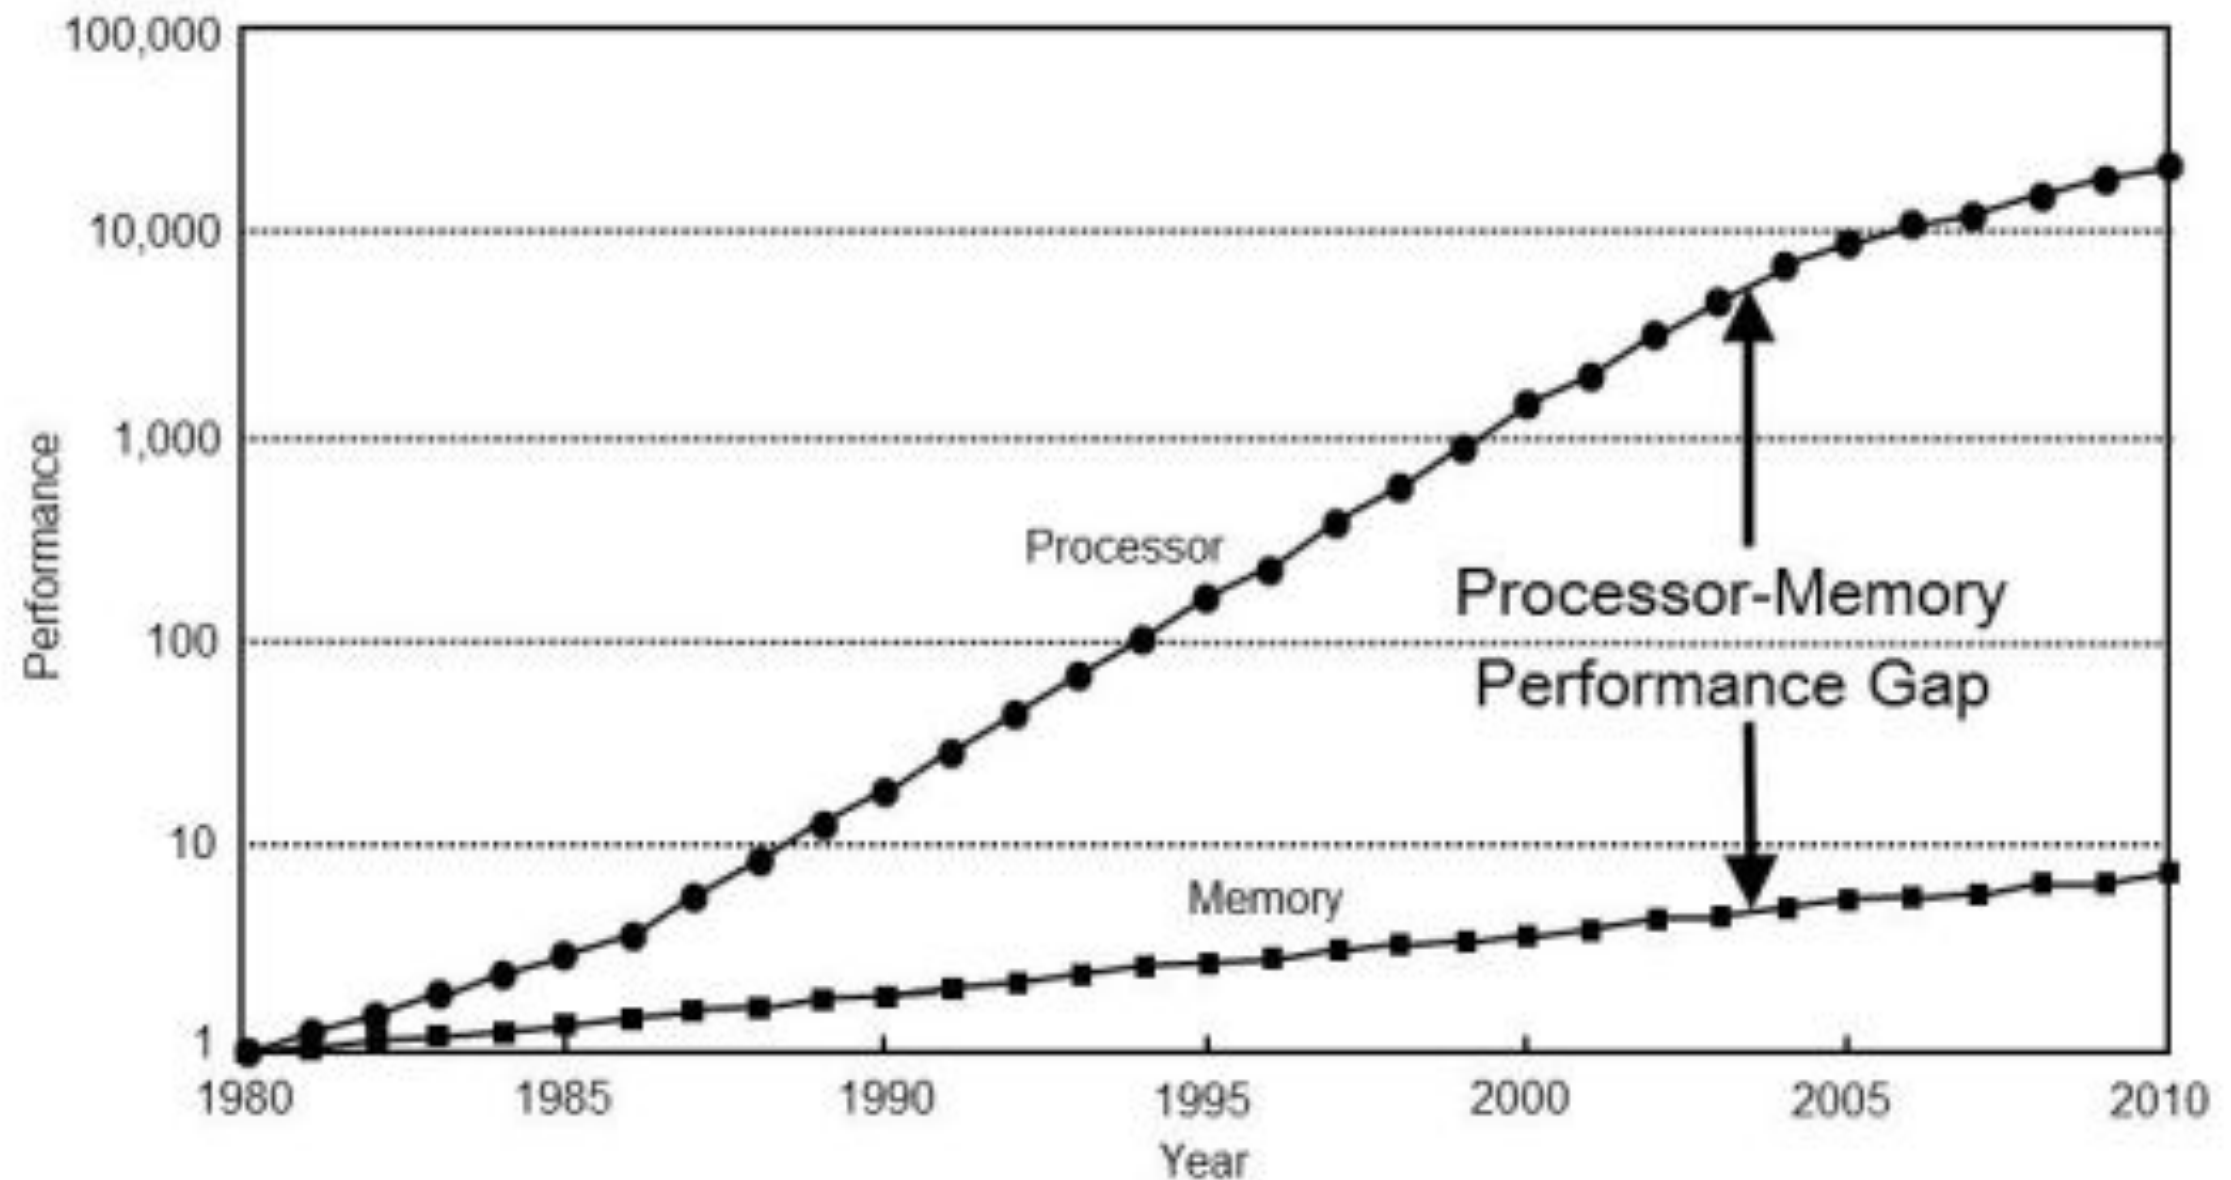
- Analyze your program at runtime to measure characteristics of interest
 - Space/time complexity
 - Frequency of certain instructions
 - Frequency and Duration of Function Calls
- Useful for guiding optimization or debug

Demo: *stackprof*

Beyond Software Optimization

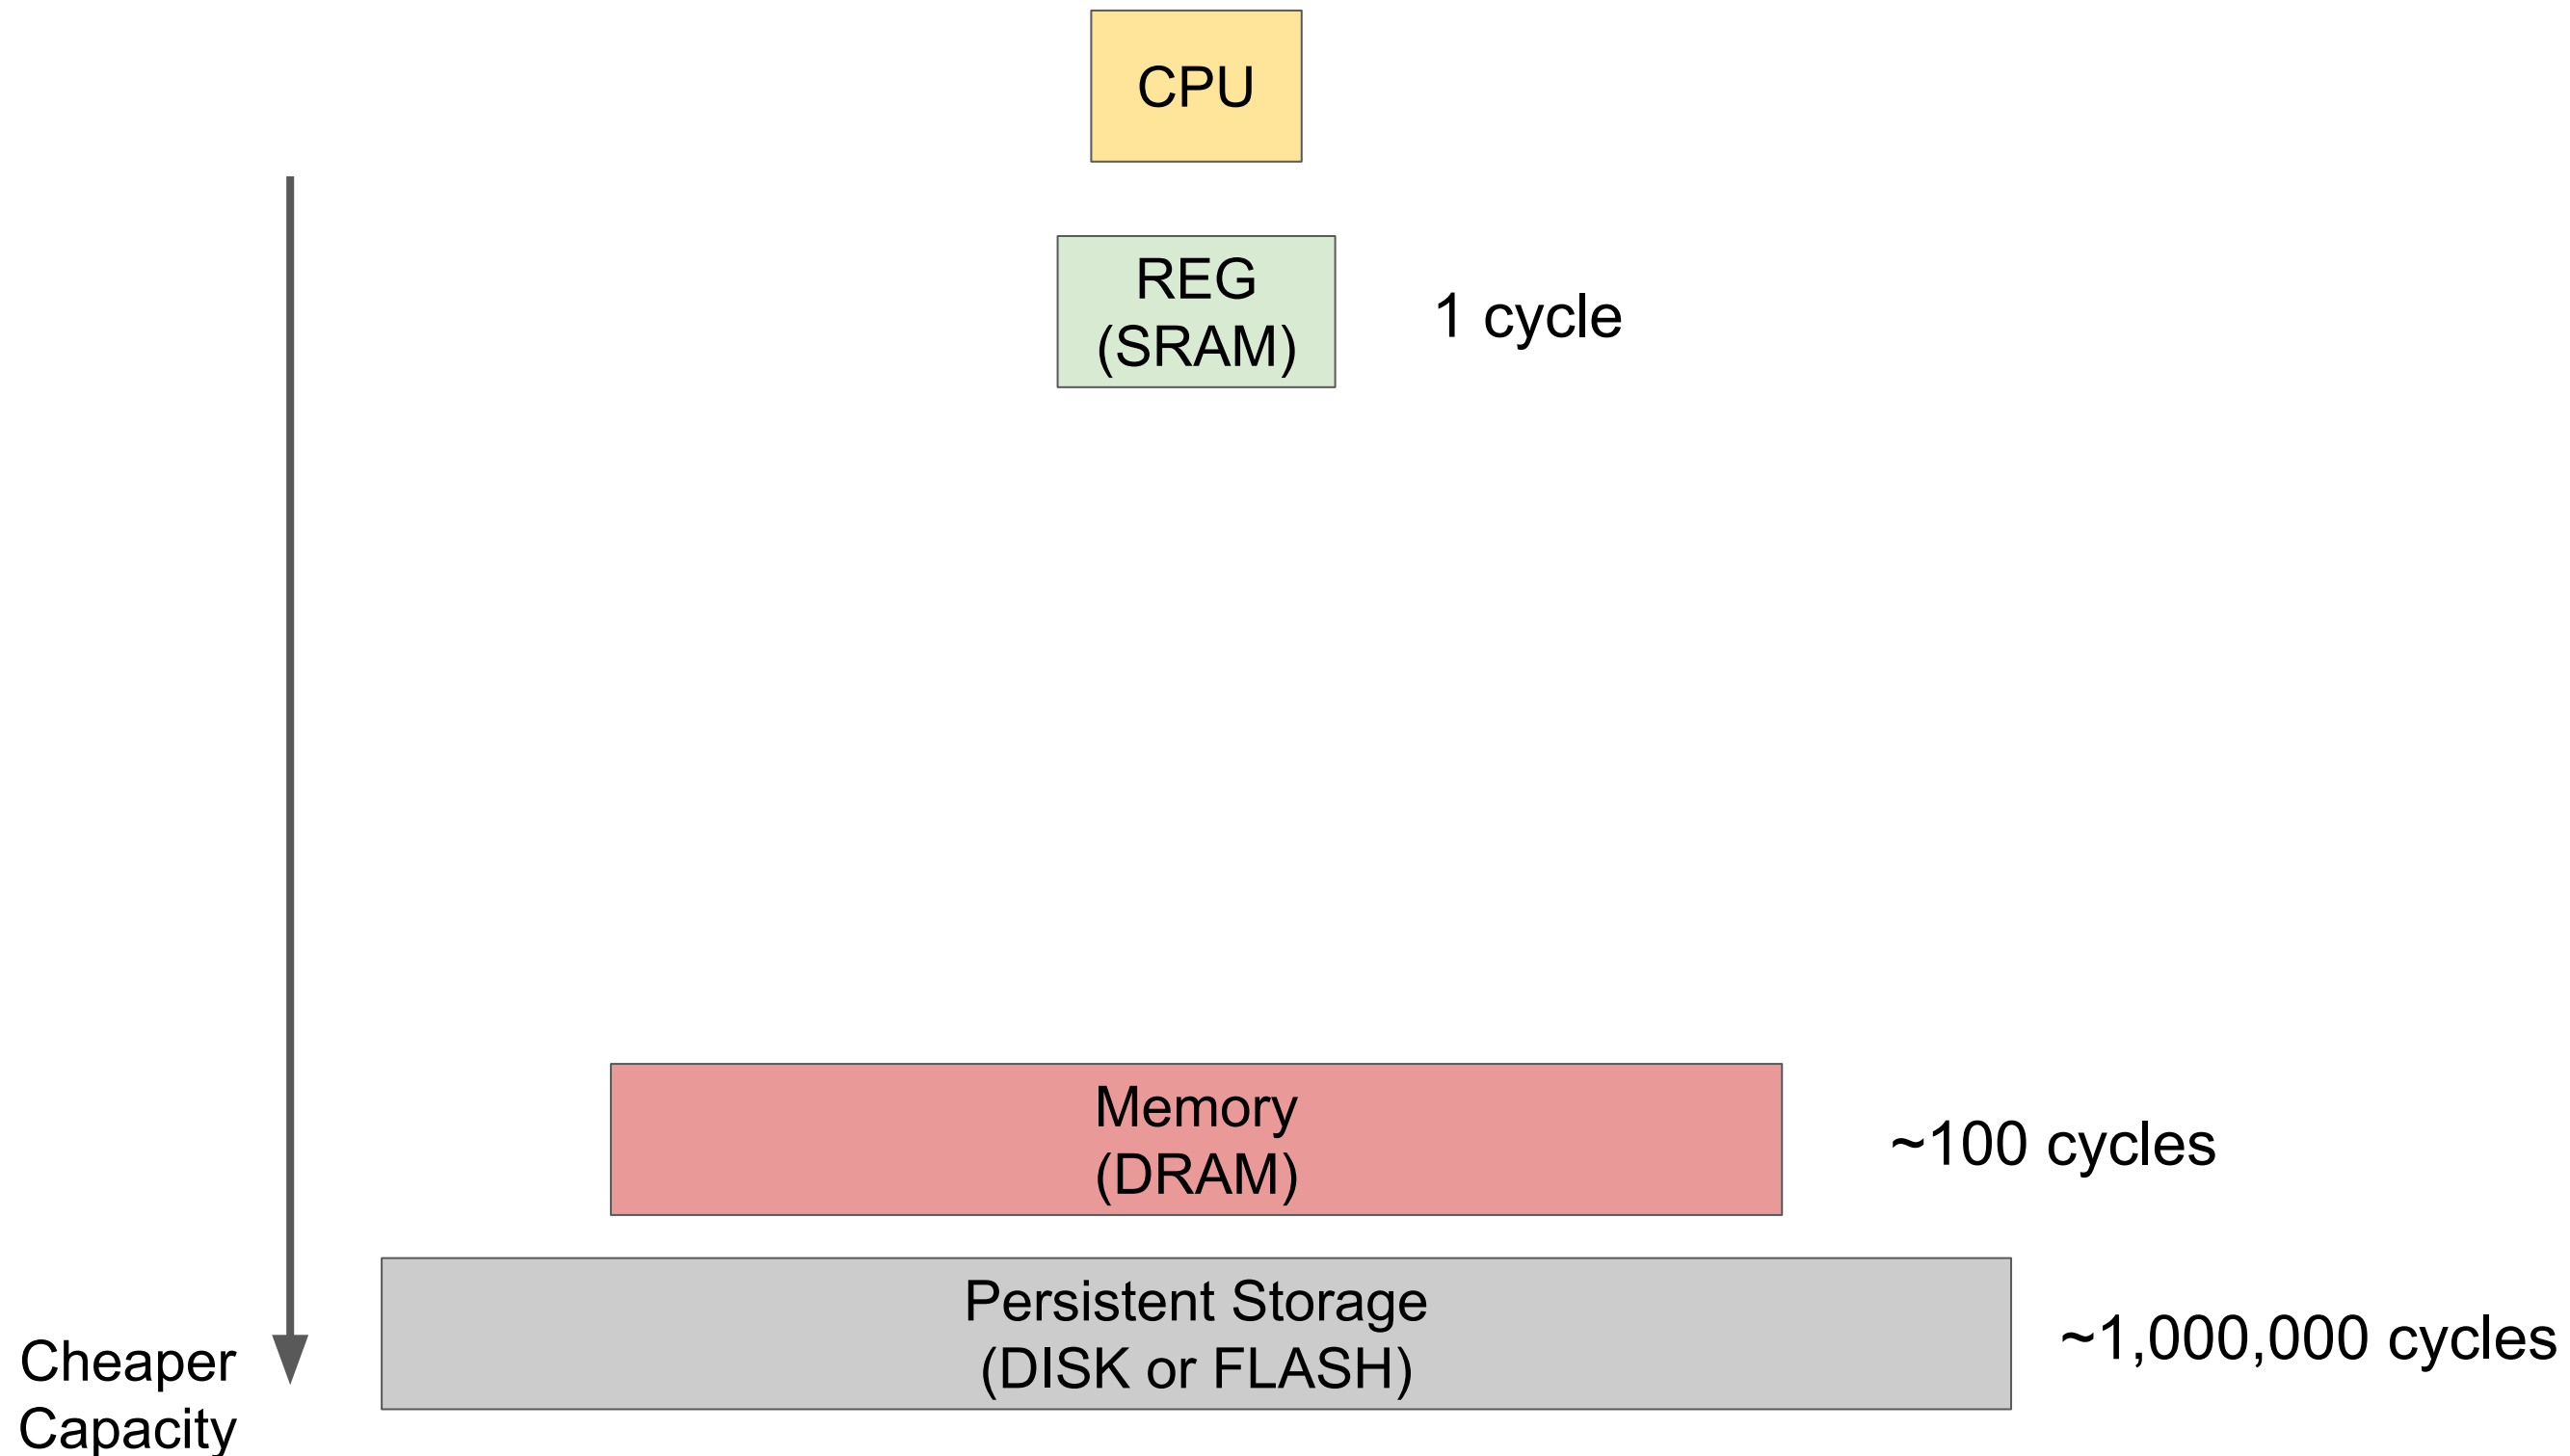


Processor-Memory Performance Gap



[Hennessy, J.L.; Patterson, D.A. Computer Organization and Design, 2nd ed.]

Memory Hierarchy



```
// [r0]++;  
ldr    r1, [r0]  
add    r1, #1  
str    r1, [r0]
```



**Wait for 100
clocks?!**

Me: optimizes assembly code

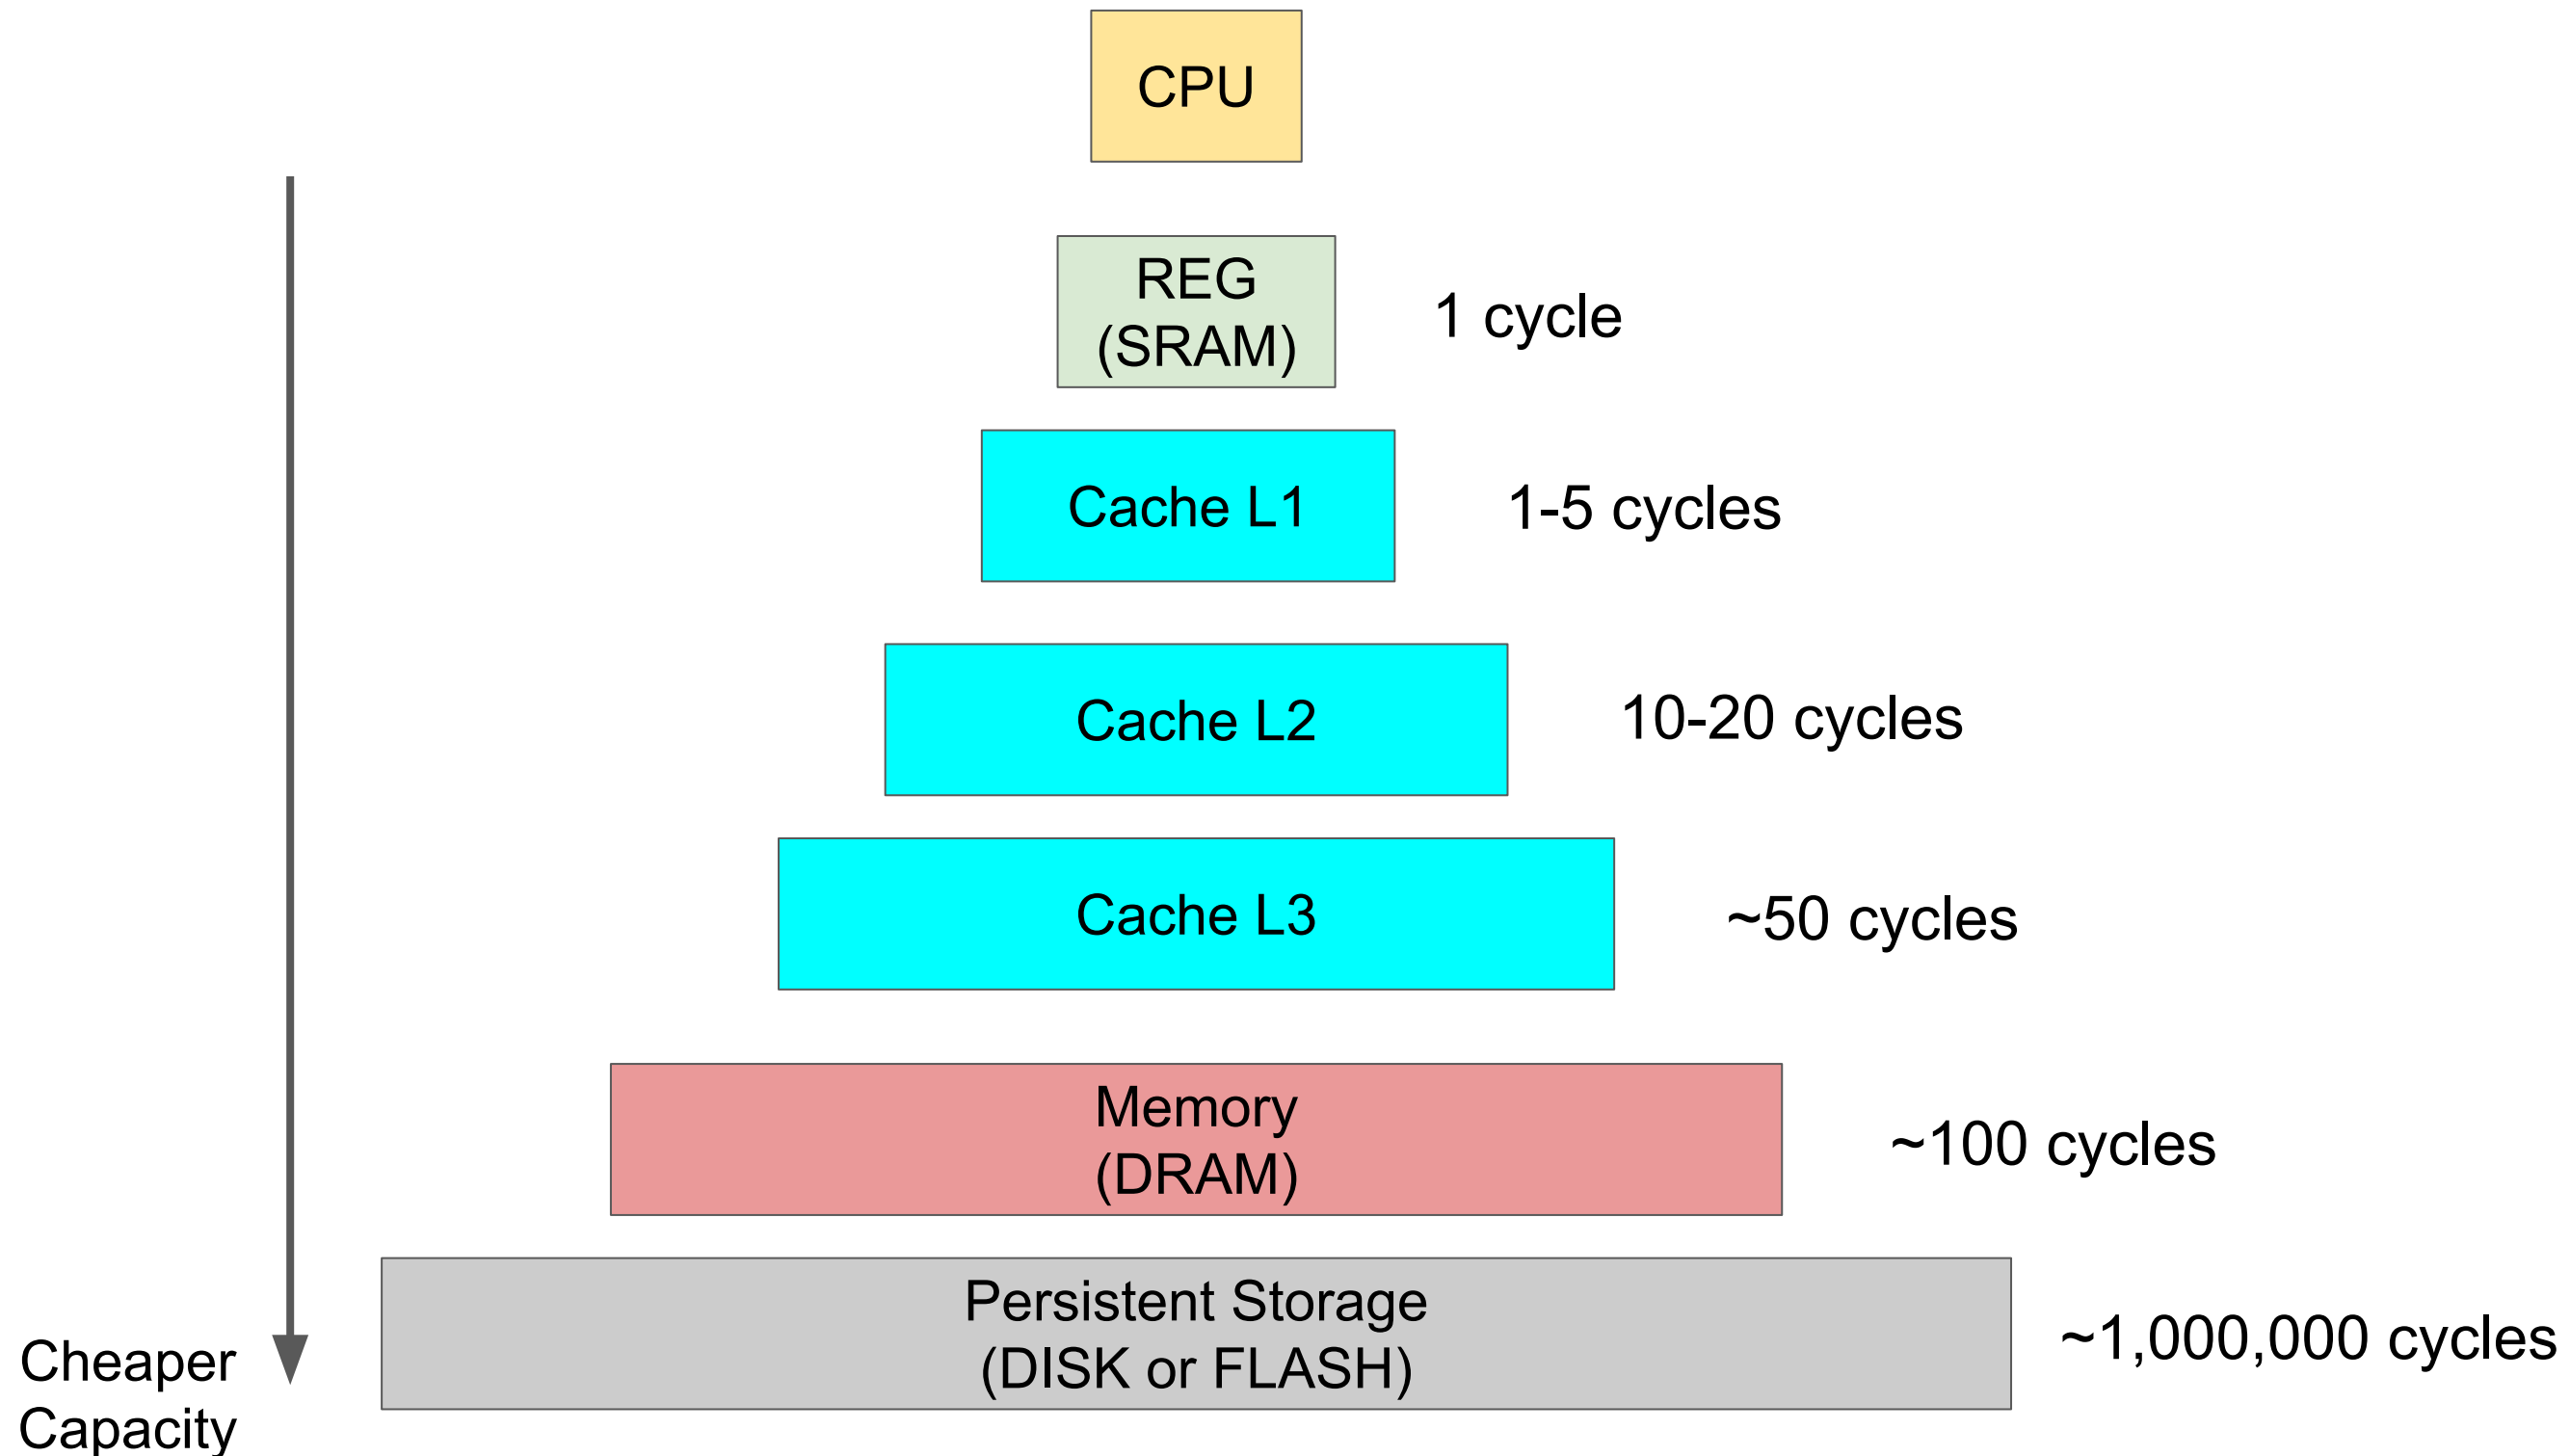
/ 1. Memory: stalls

a. Me: for 100 cycles

)



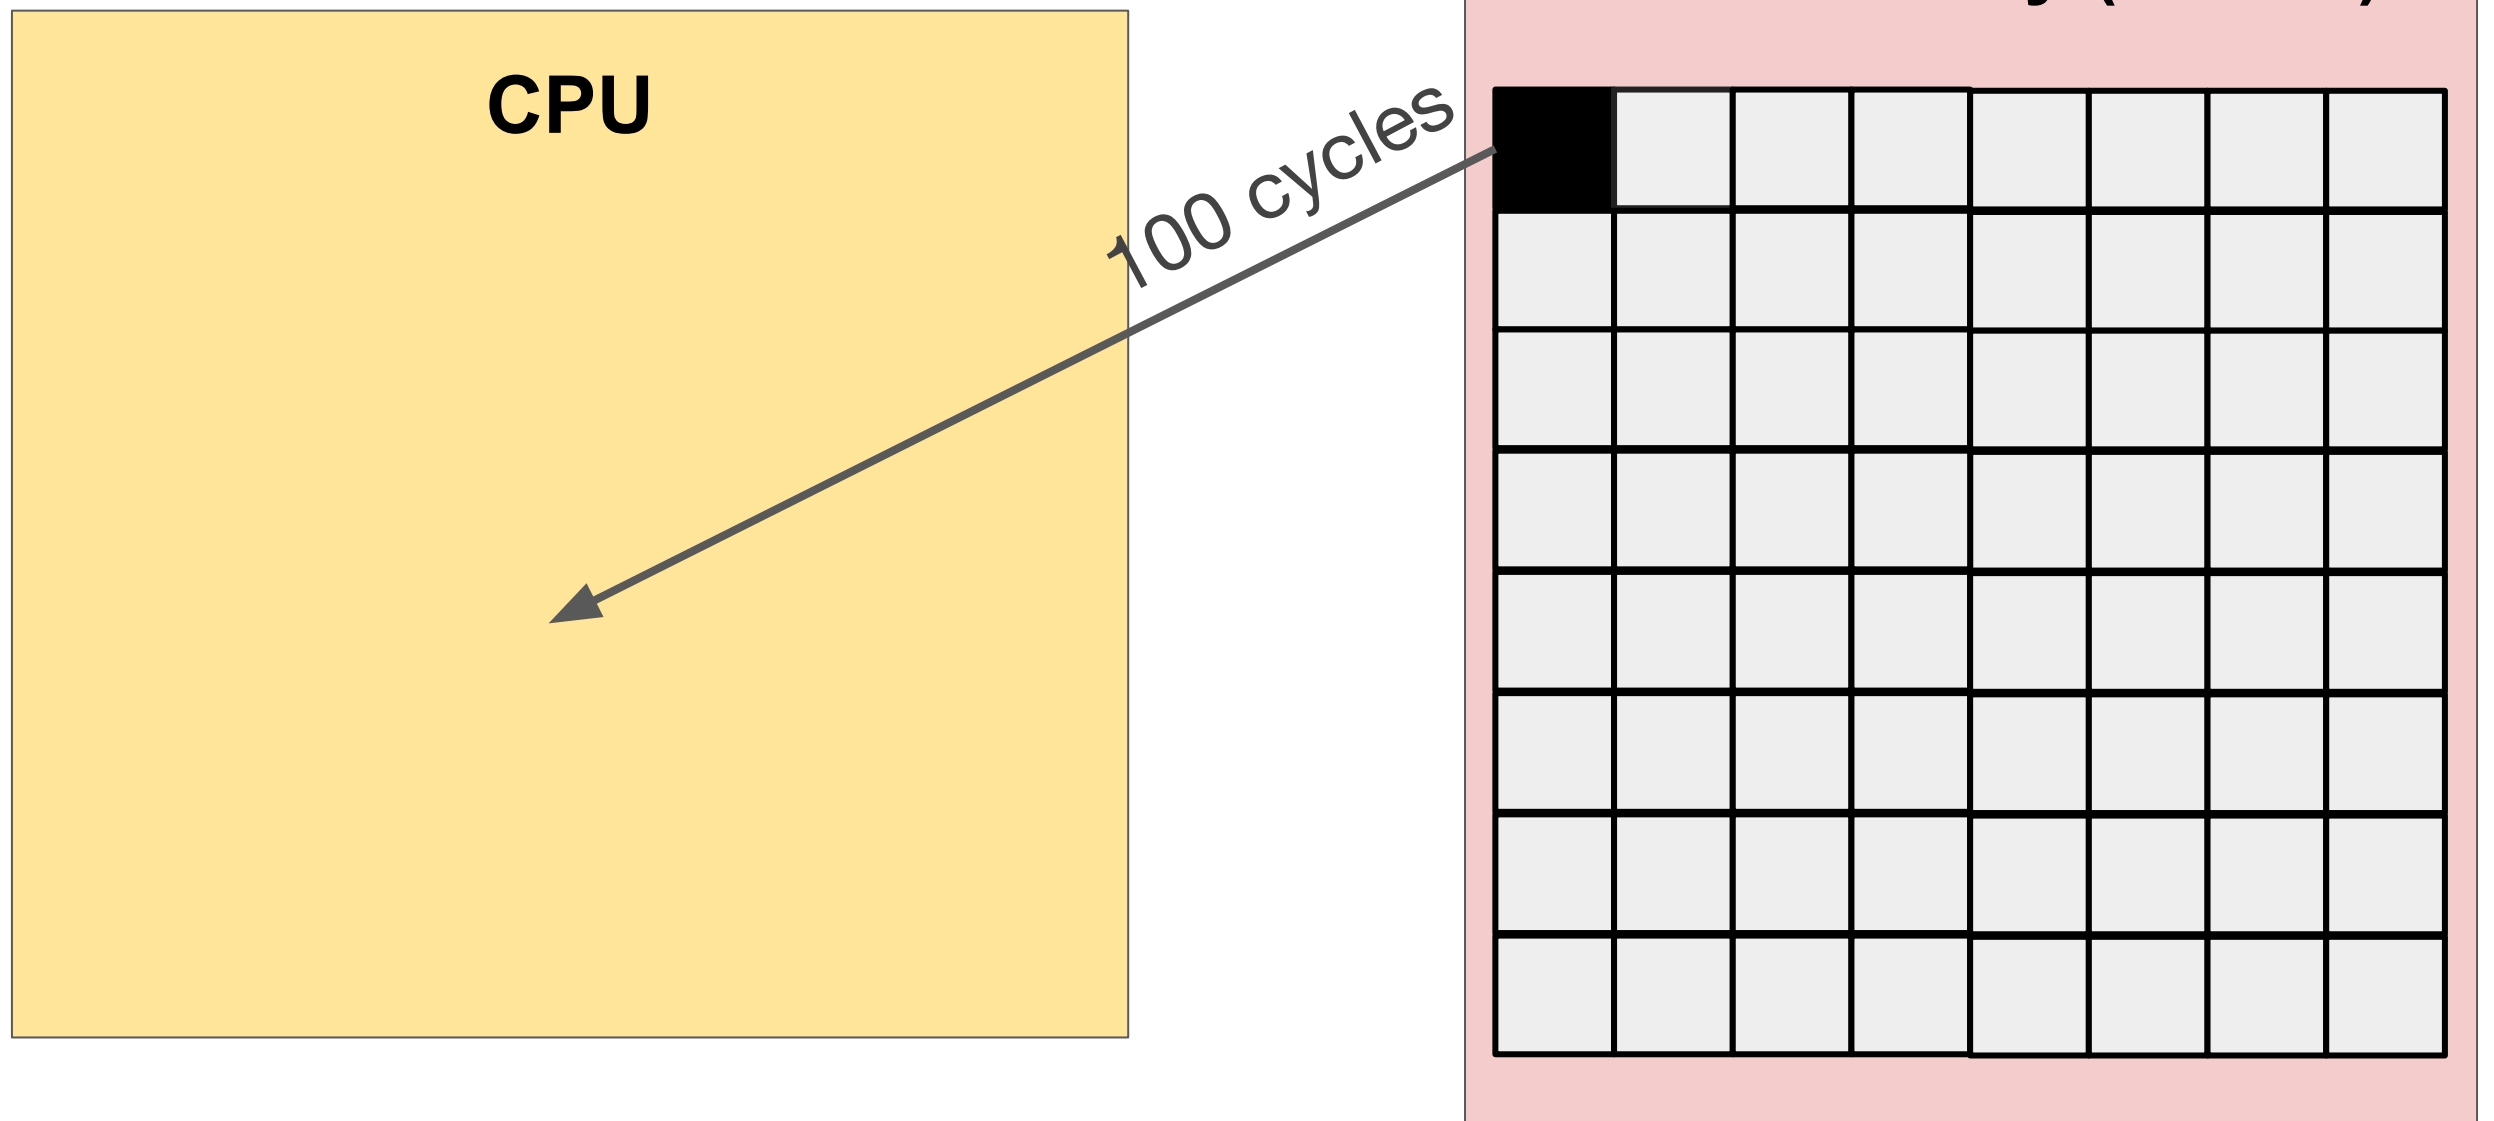
Memory Hierarchy



w/o caches

CPU wants to access memory location in black.

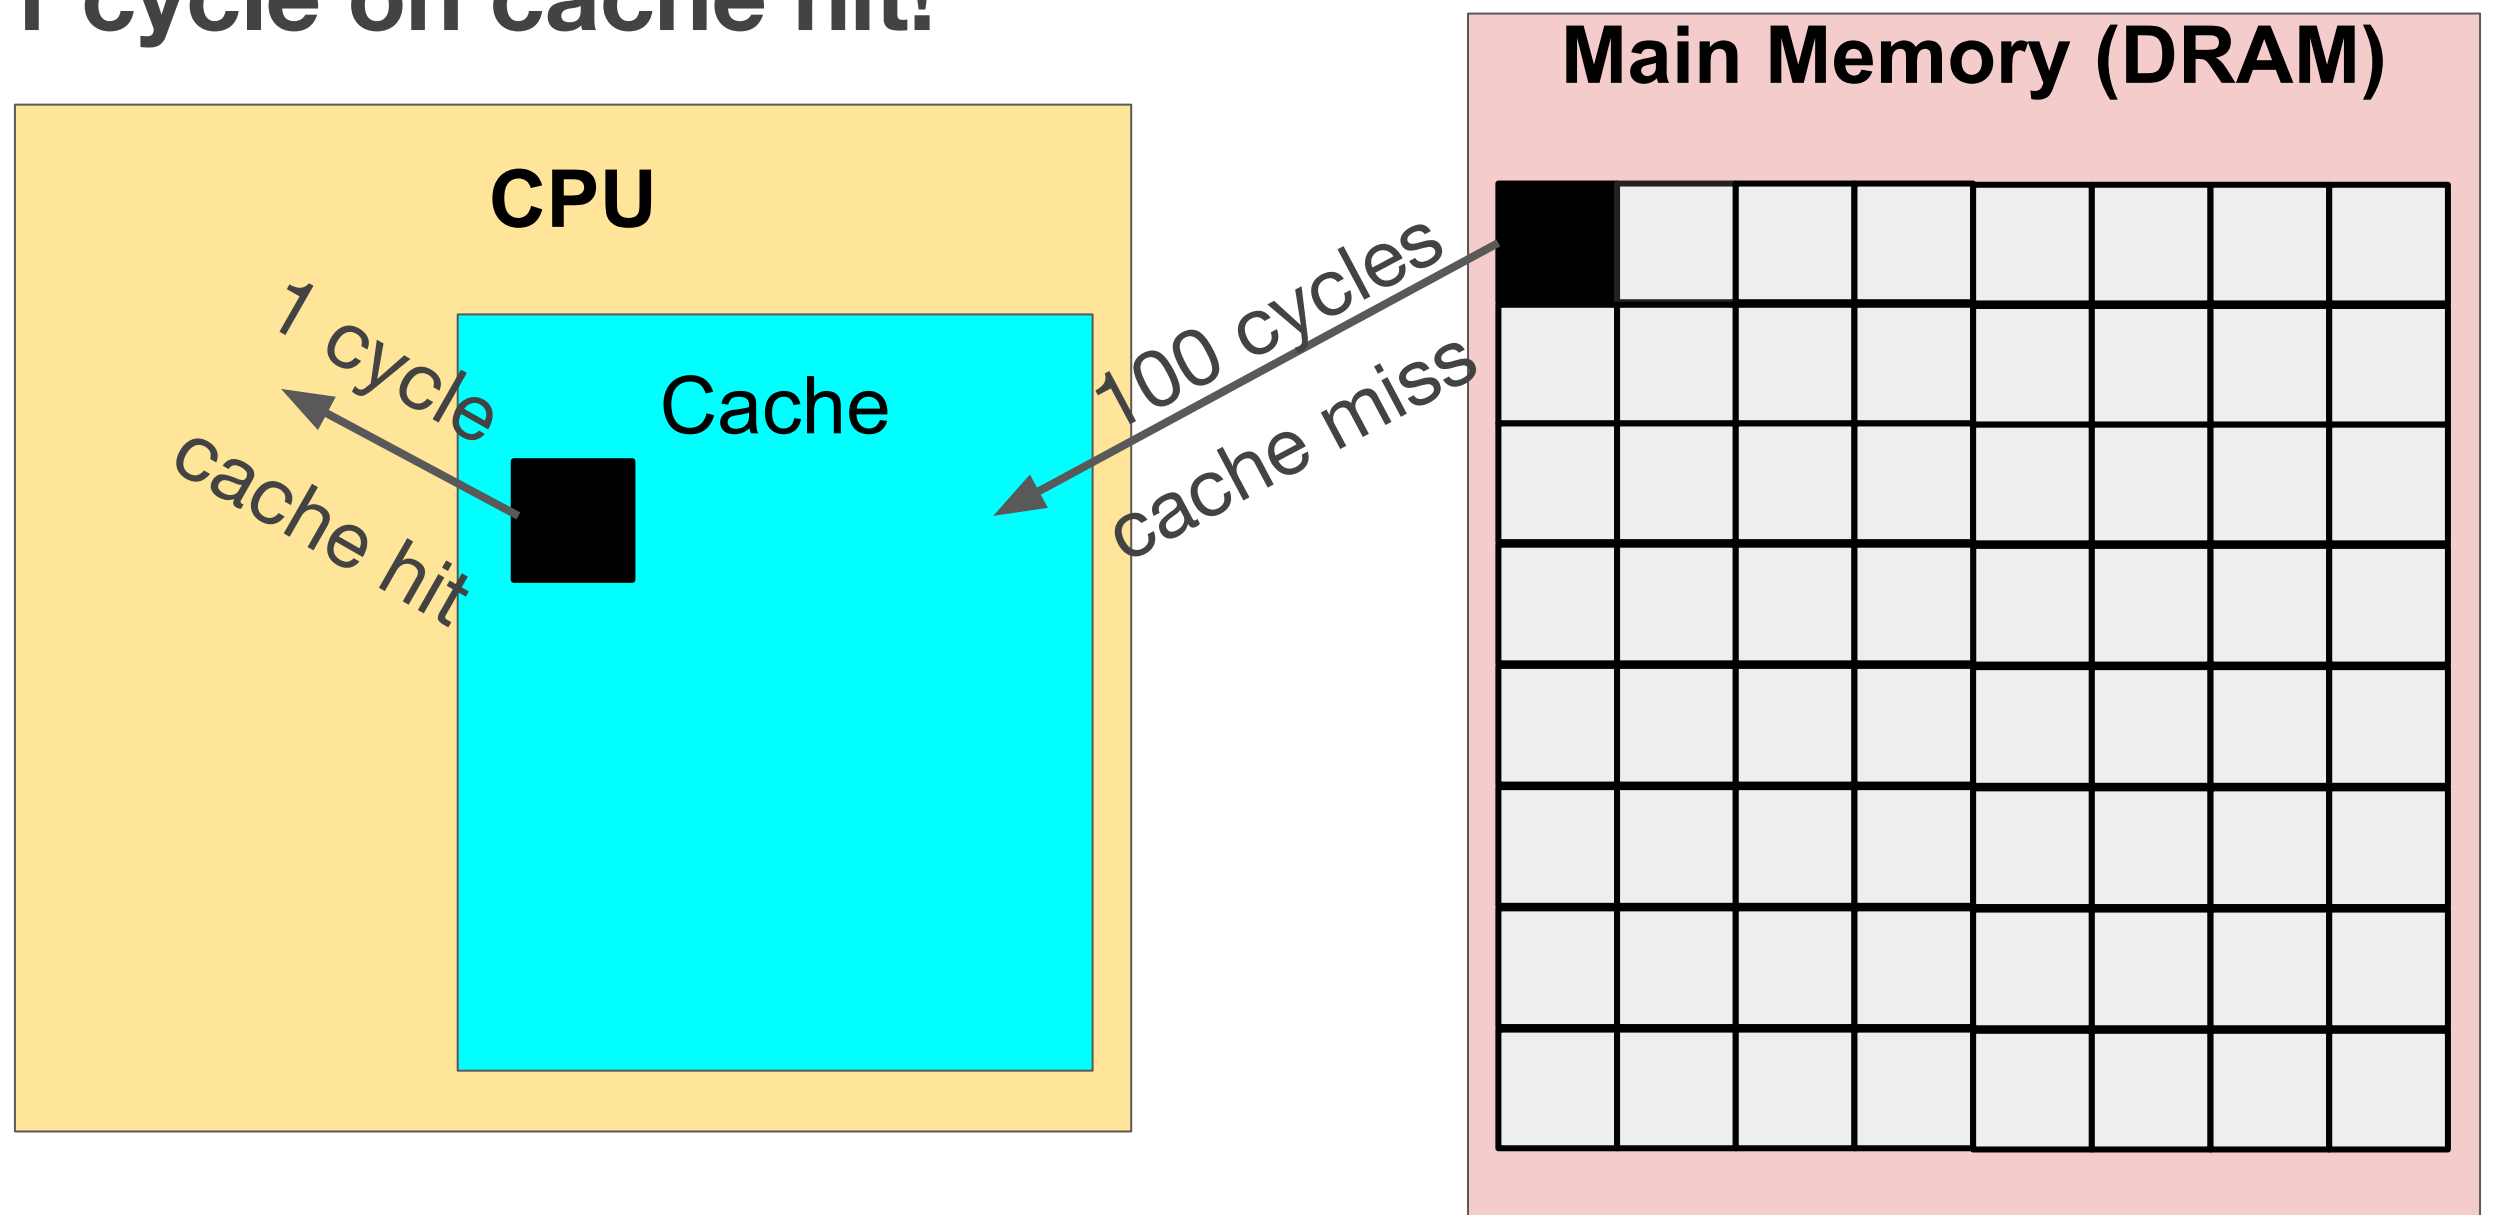
Take 100 cycles every access



Caches

100 cycles on cache miss (main memory access).

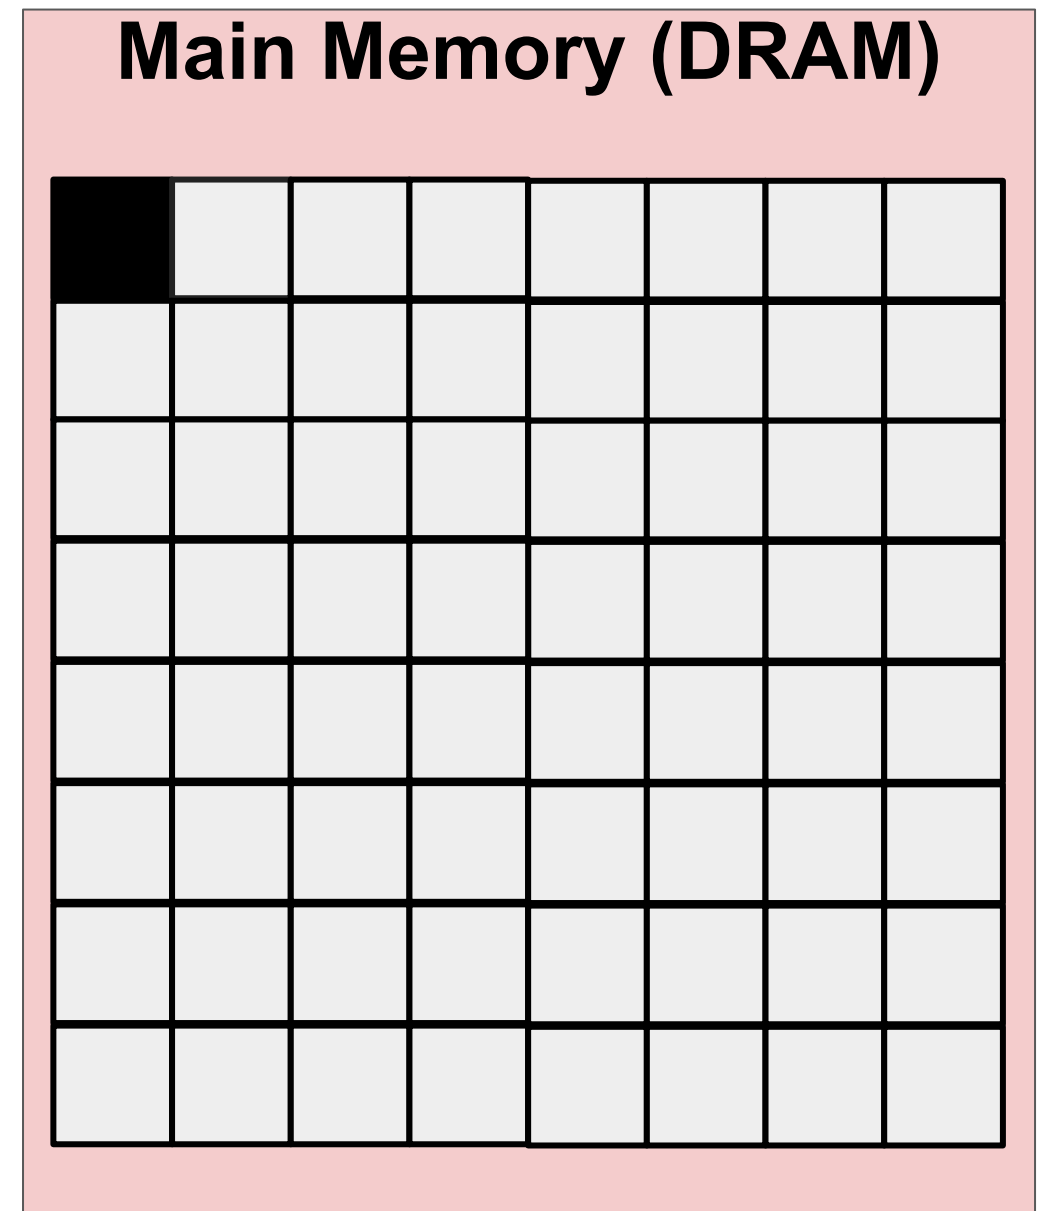
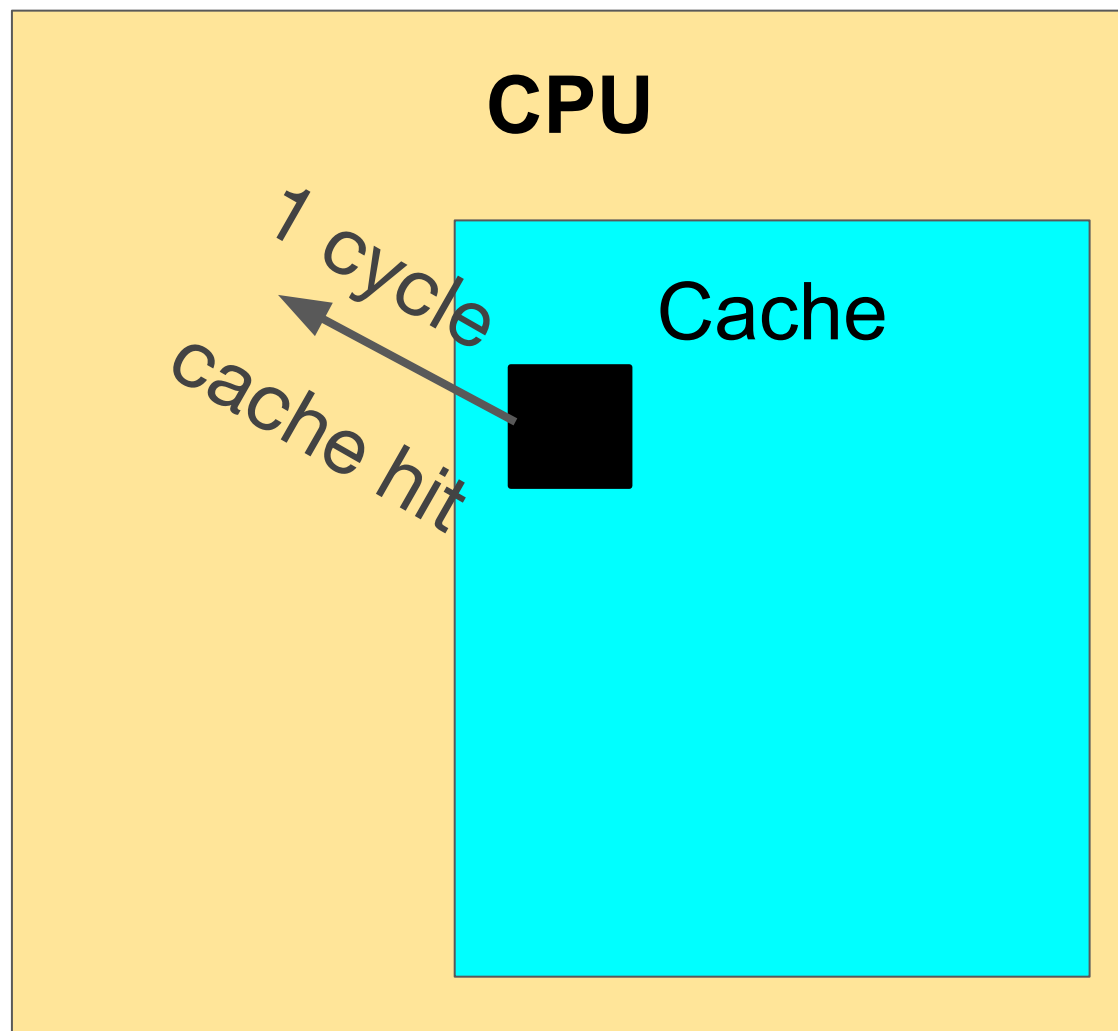
1 cycle on cache hit!



Caches

Concept 1: *Temporal Locality*

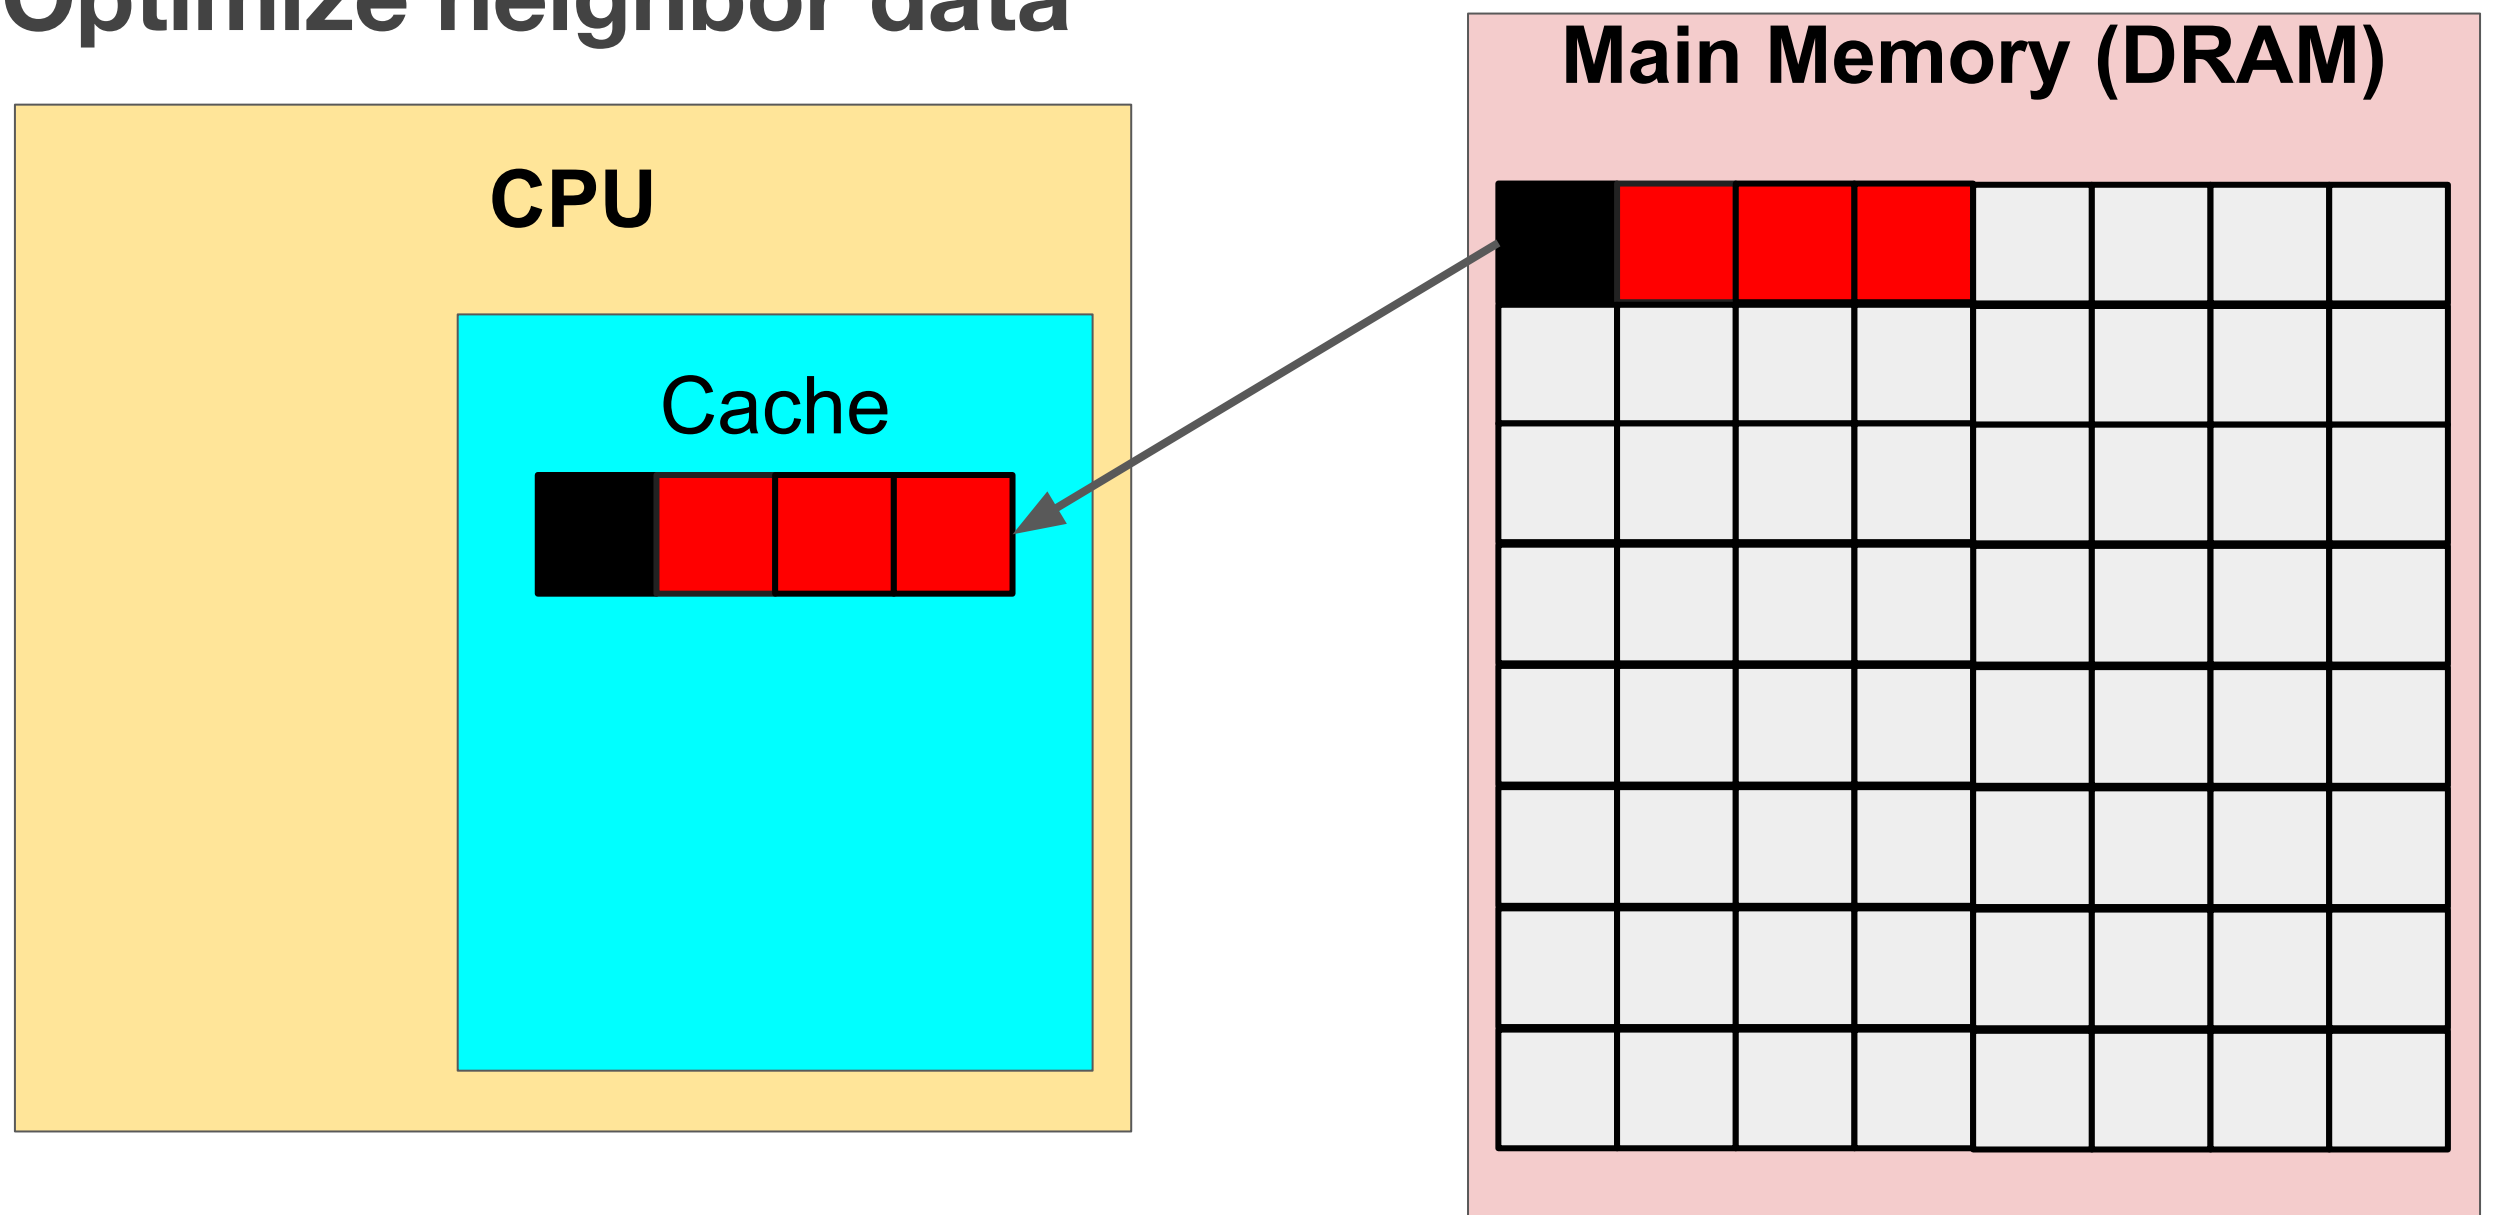
Optimize frequently used data



Caches

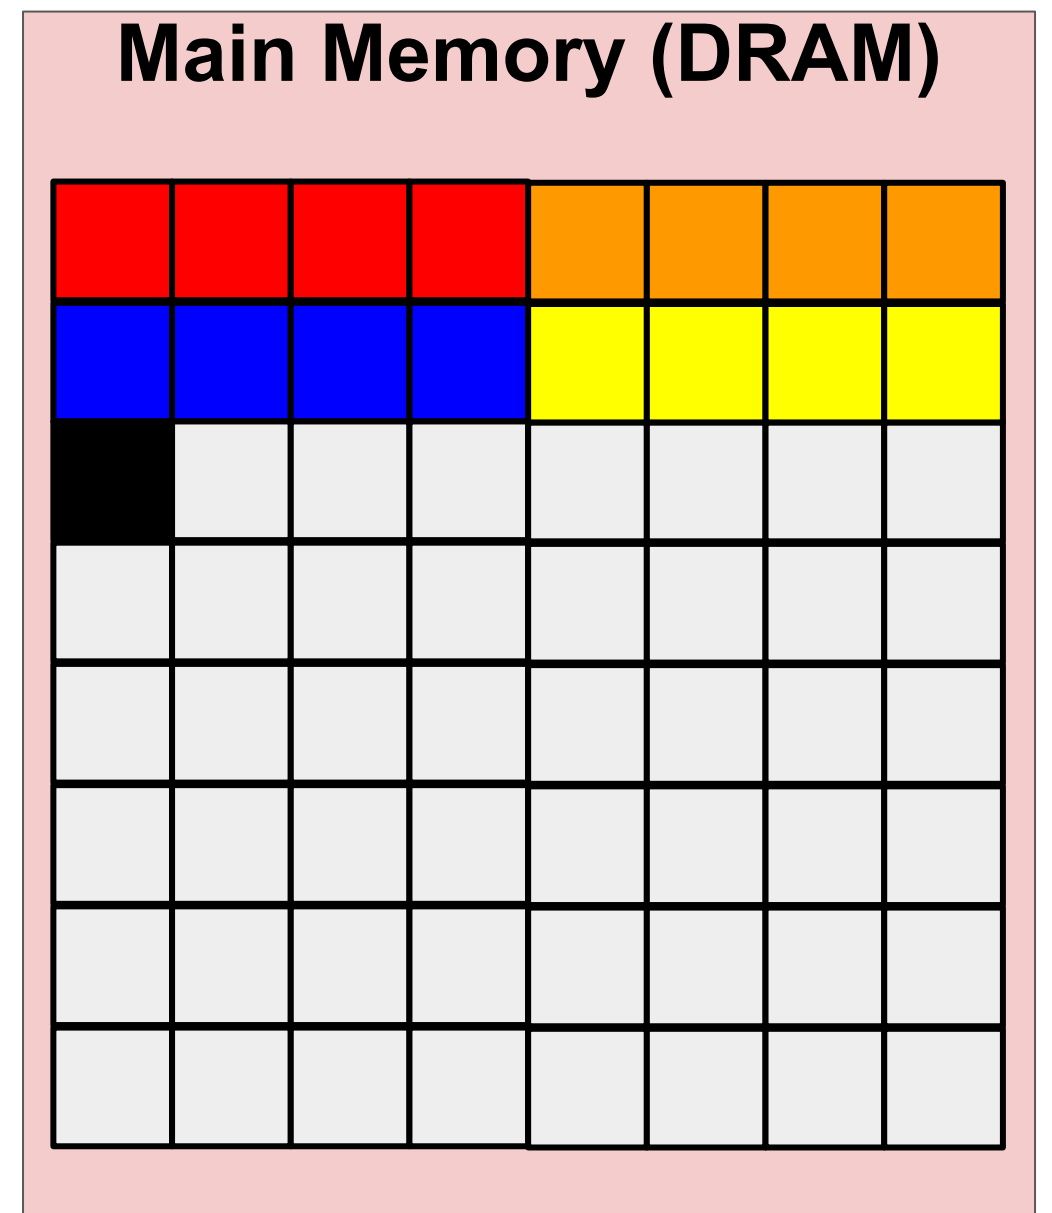
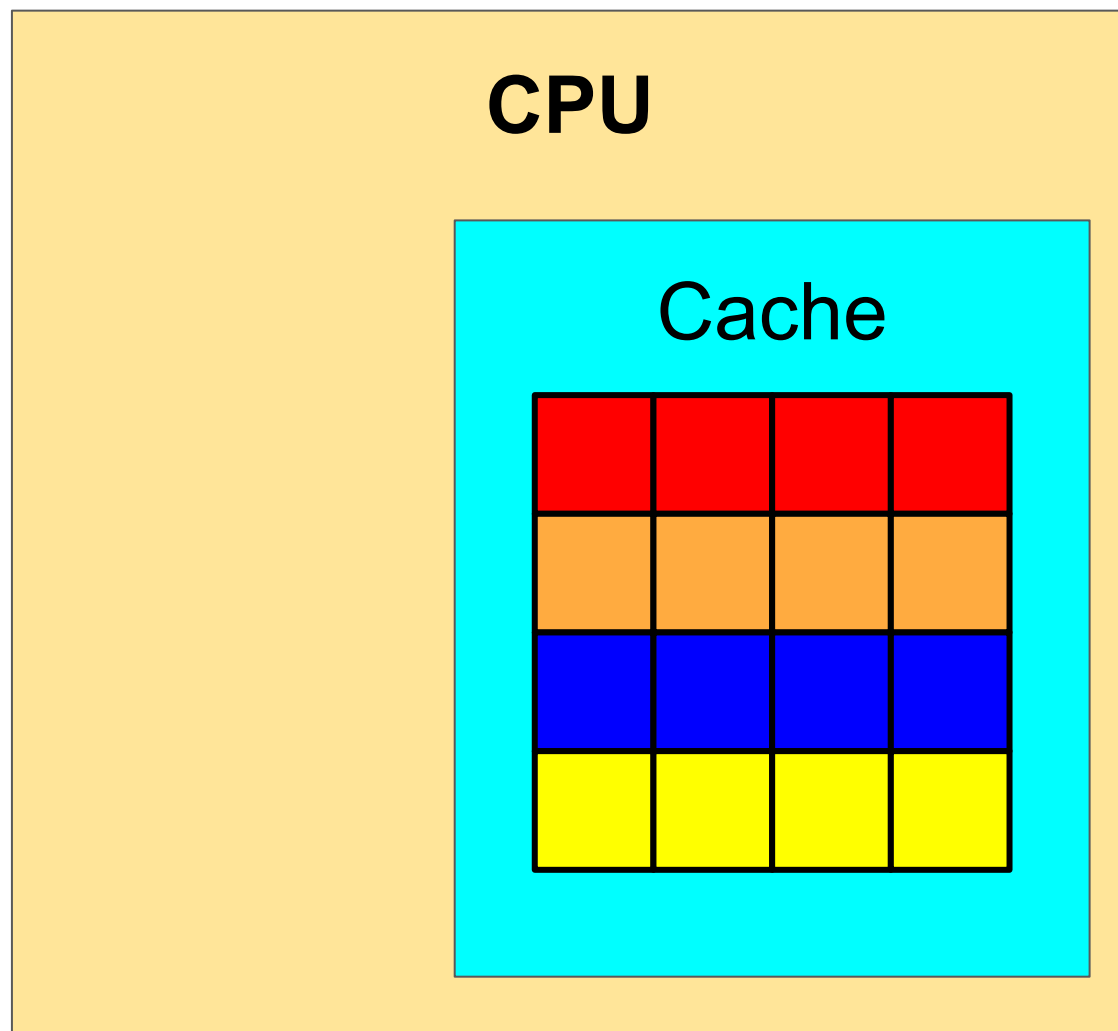
Concept 2: *Spatial Locality*

Optimize neighbor data



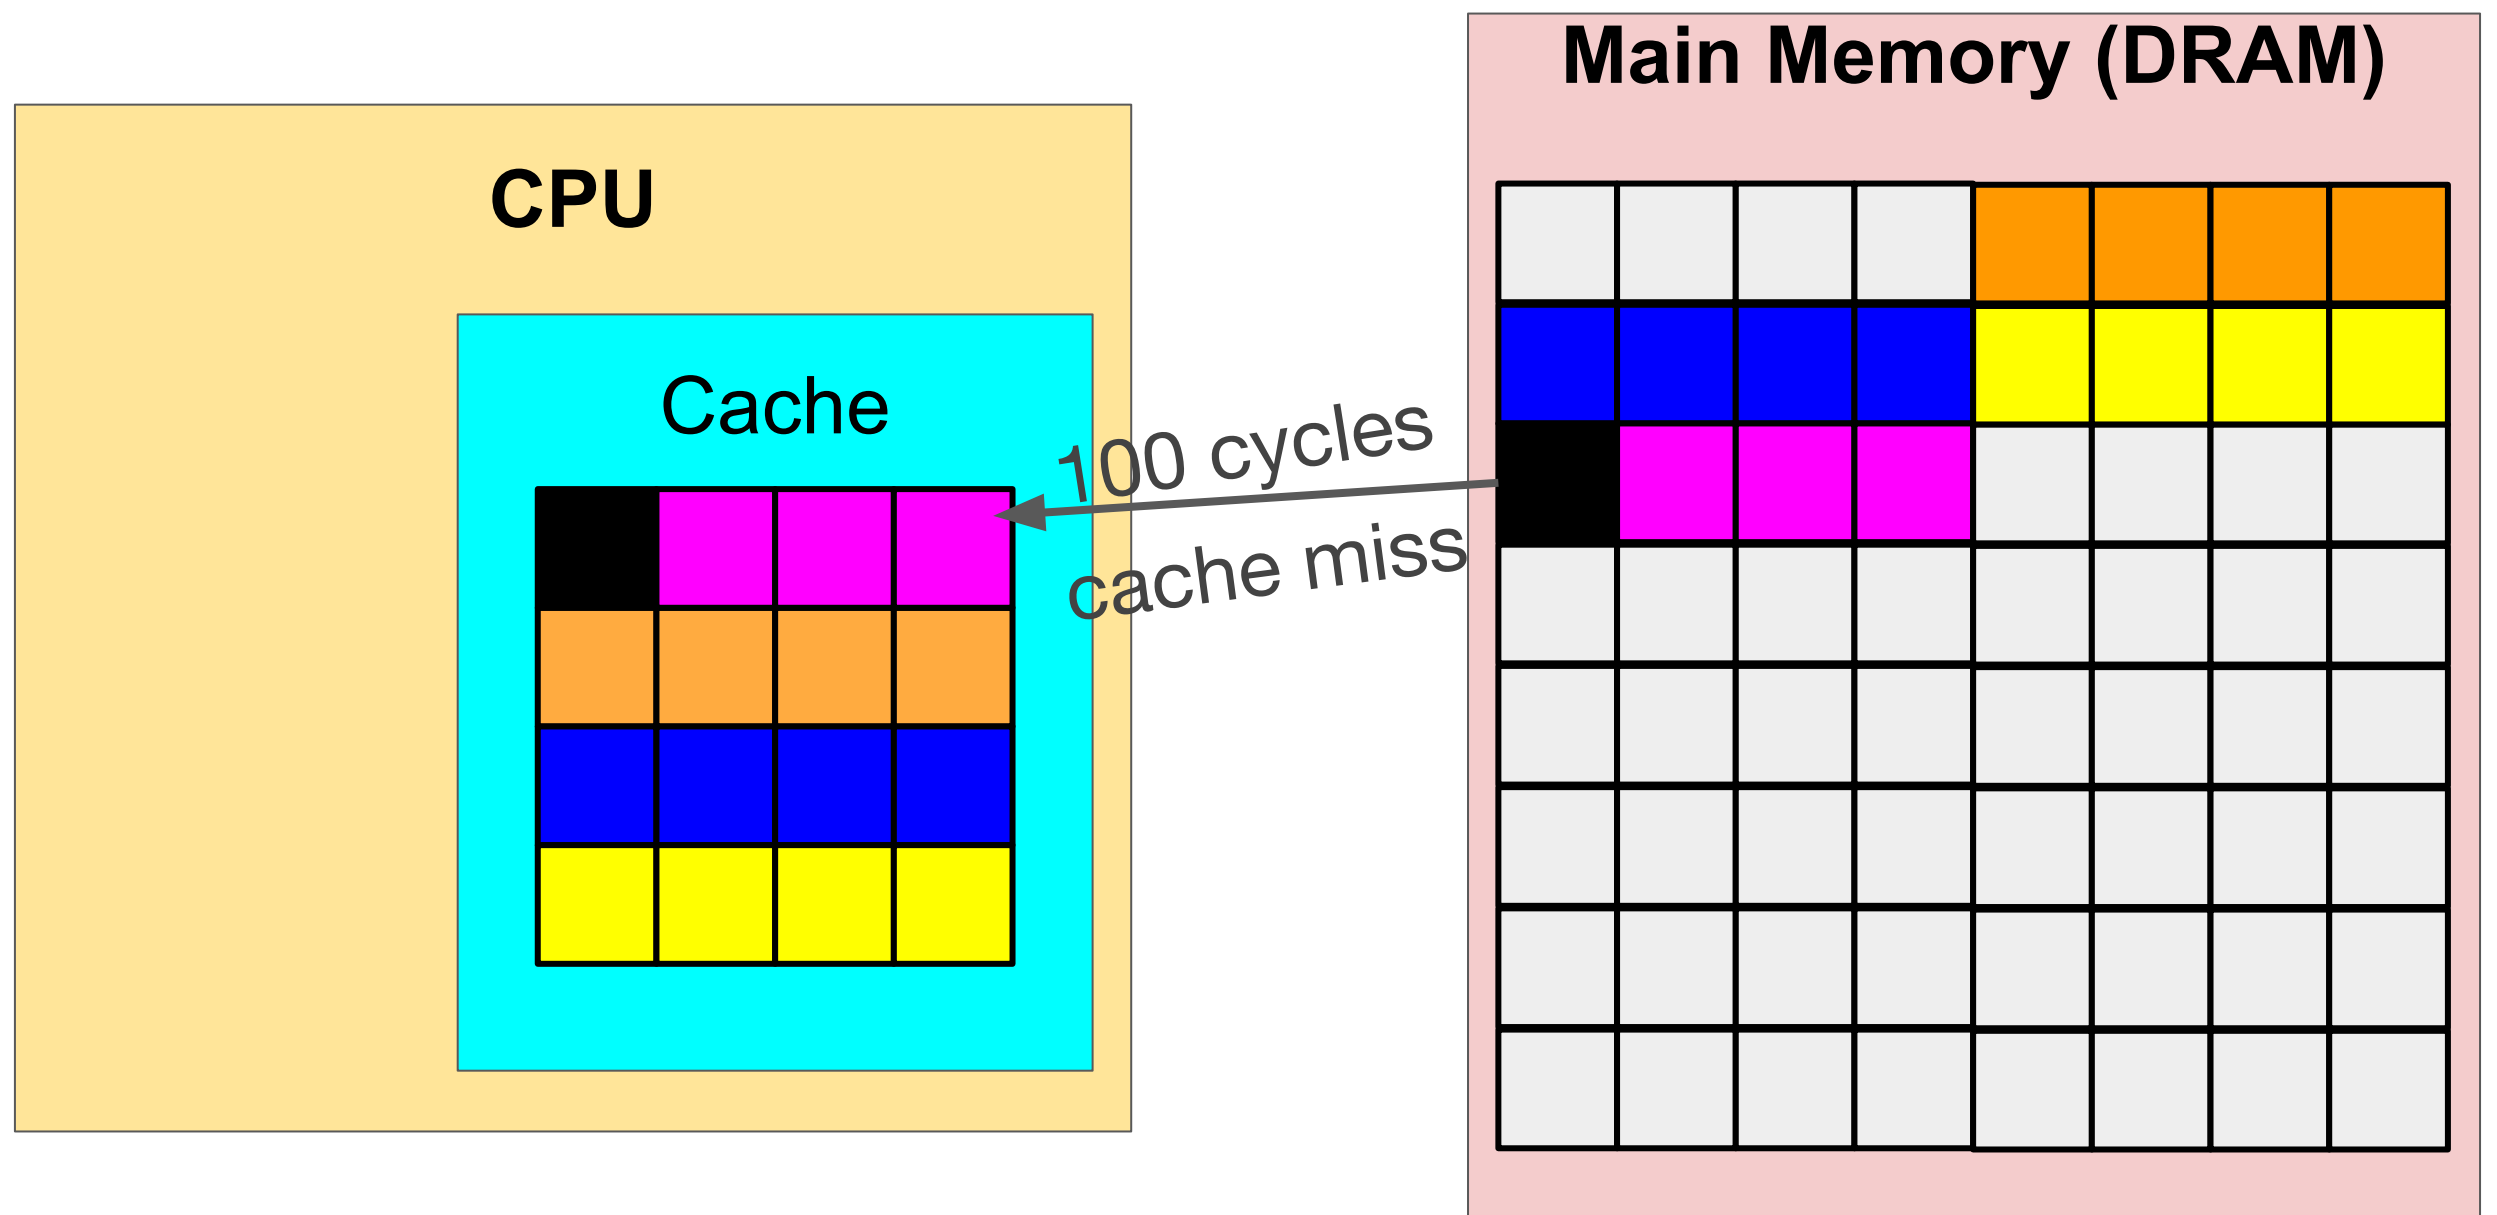
Caches

Limited cache size



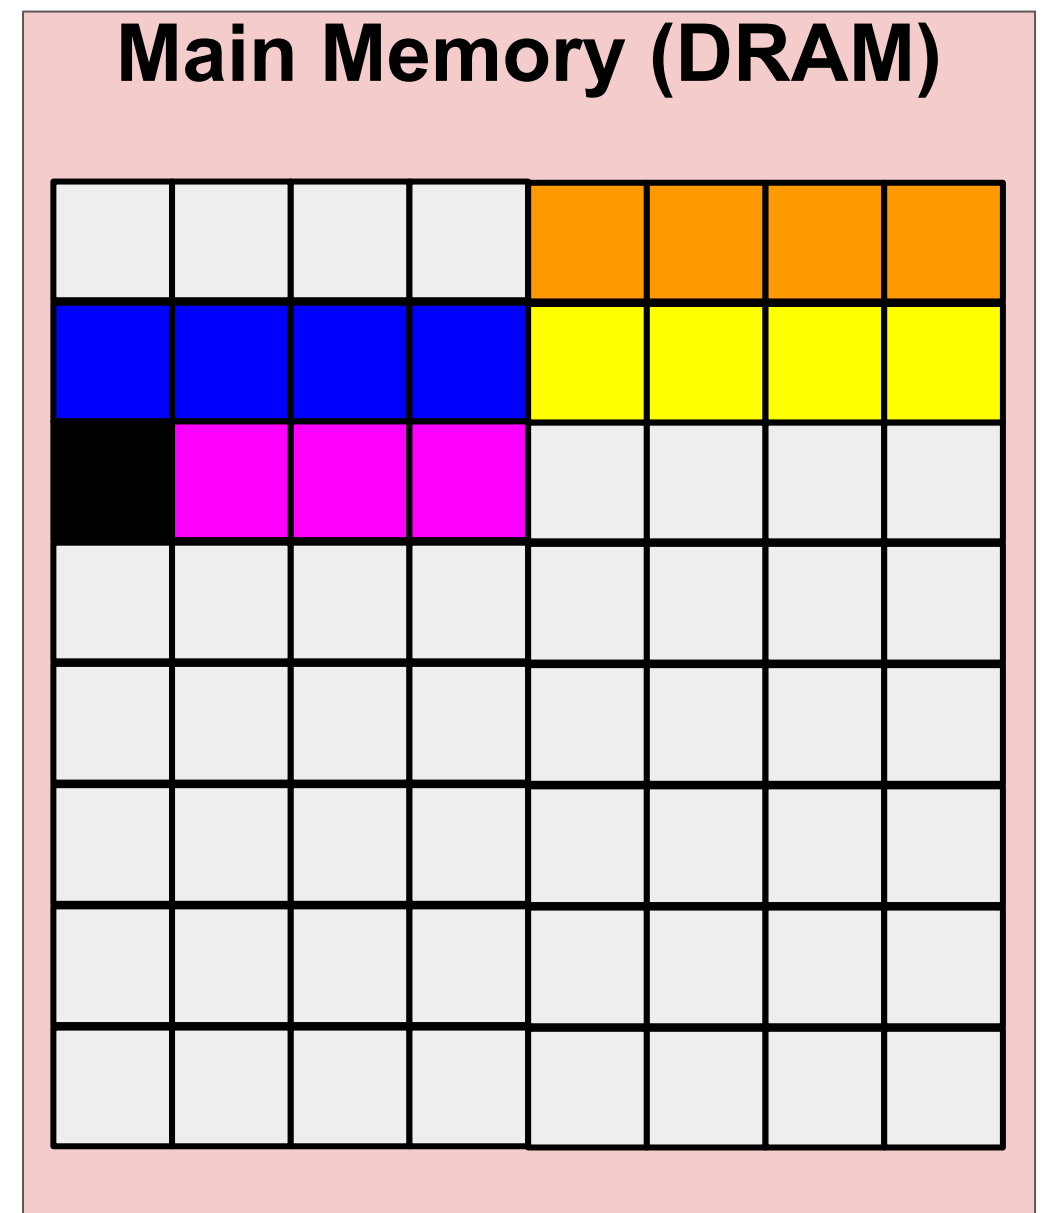
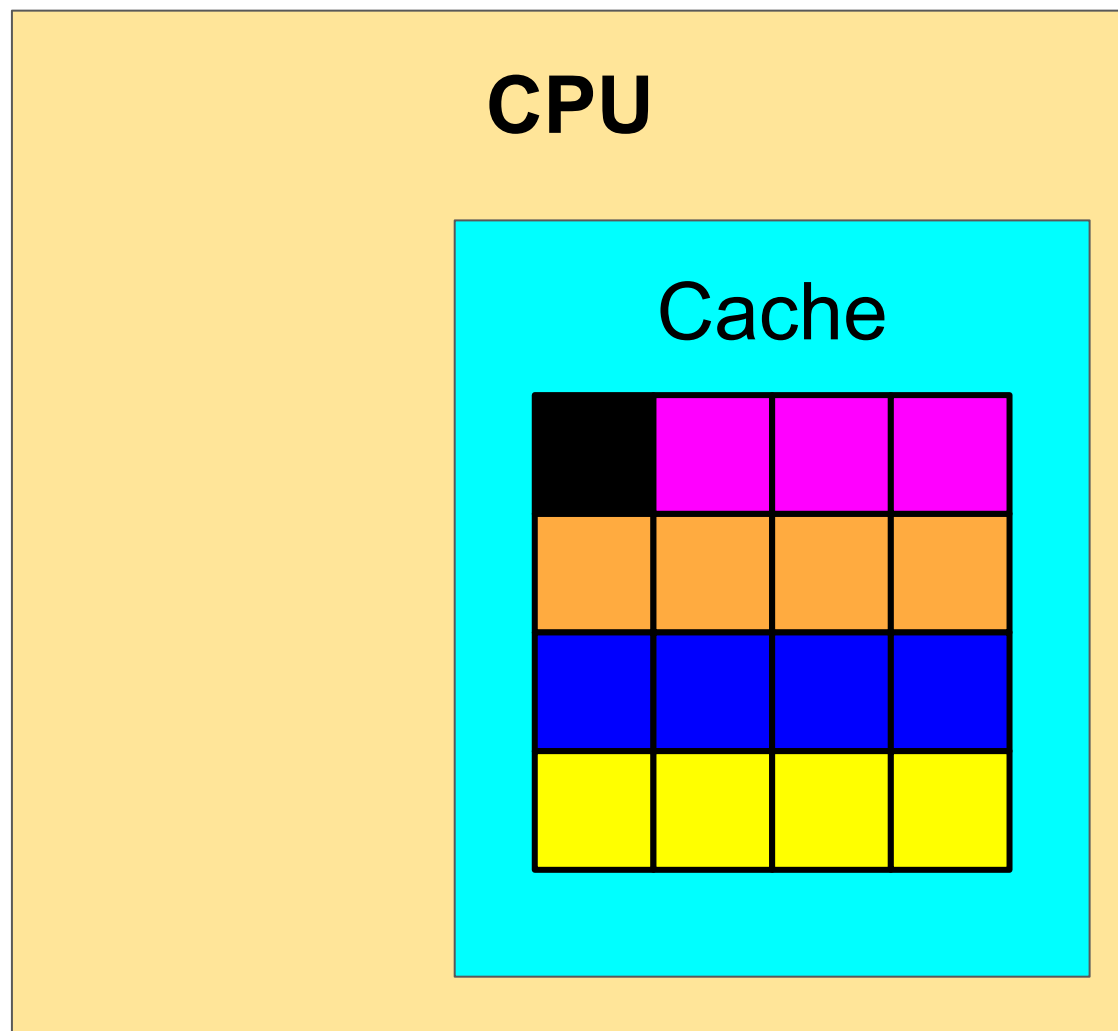
Caches

Eviction



Caches

Eviction

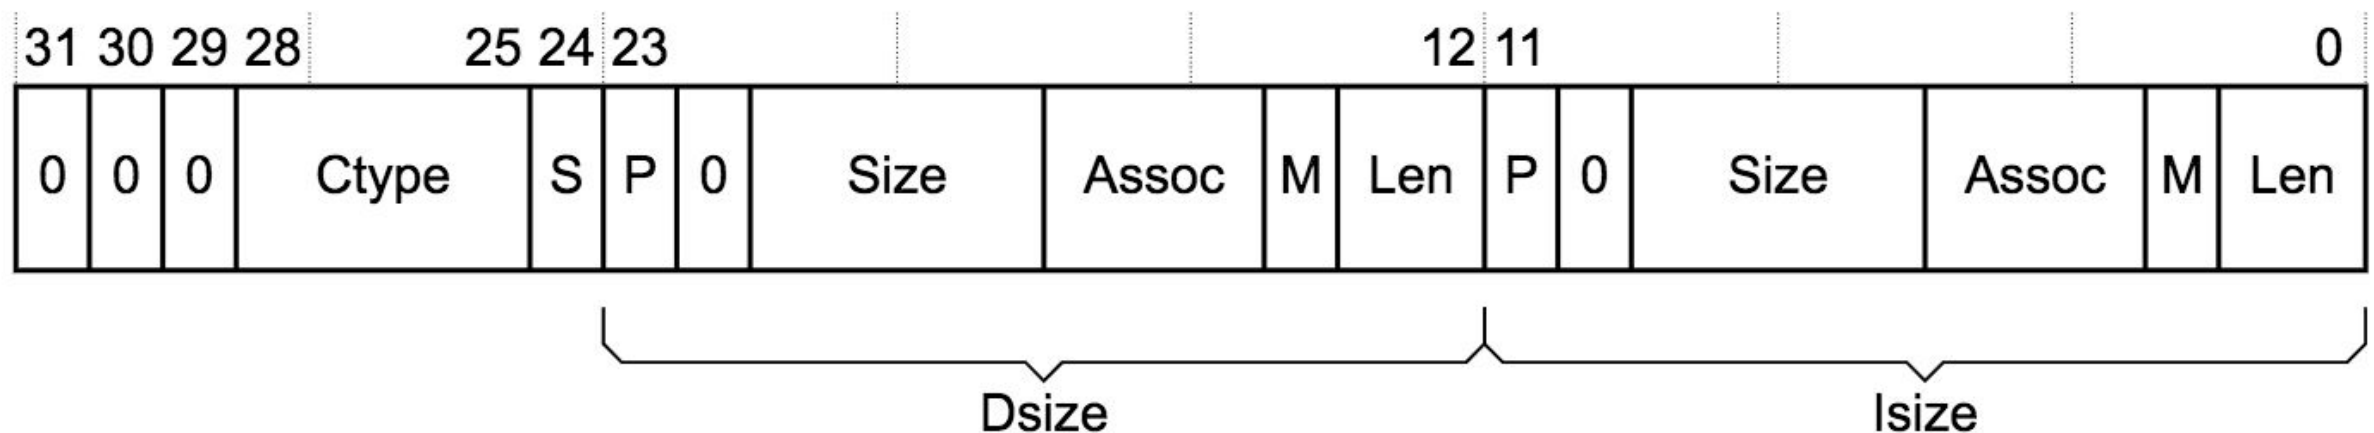


BCM 2835 Data Cache

L1 Total Size: 16 KB

L1 Block Size: 32 Bytes

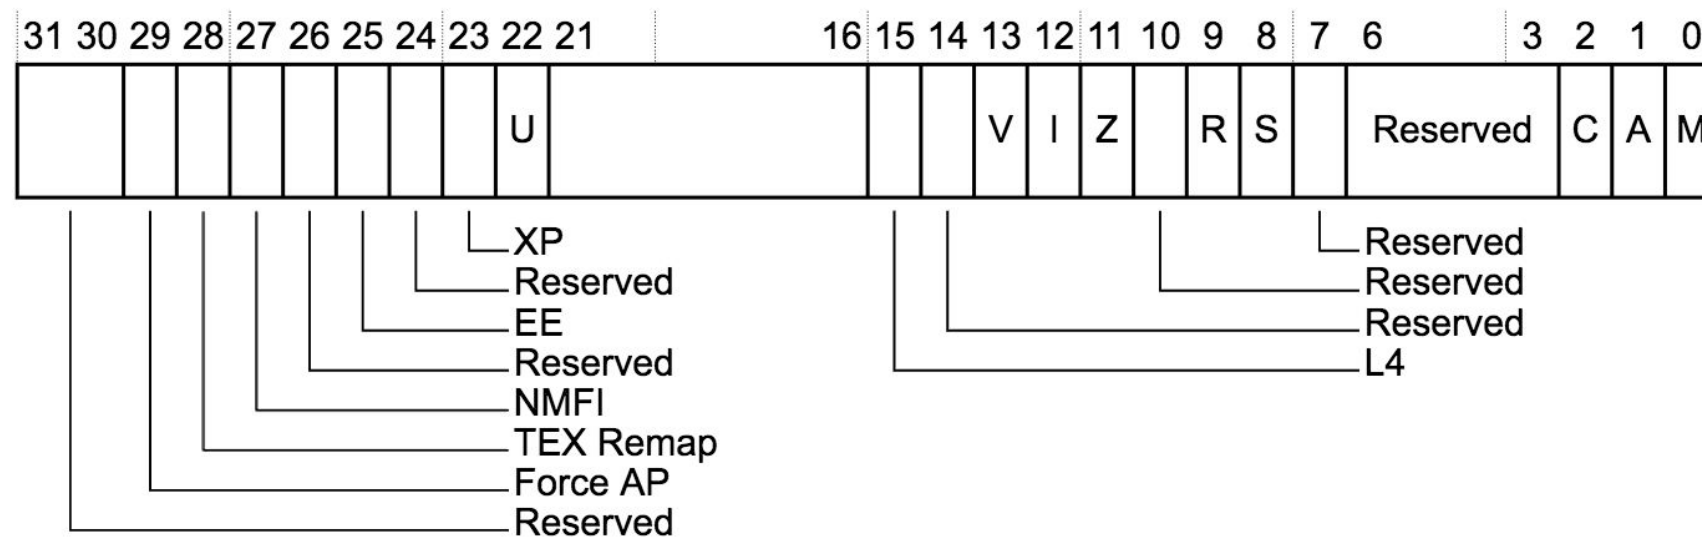
Figure 3.3. Cache Type Register format



```
// system.c
unsigned system_get_cache_type(void) {
    // See section 3.3.2 in arm1176
    unsigned reg;
    __asm__ volatile("mrc    p15, 0, %0, c0, c0, 1" : "=r"(reg));
    return reg;
}
```

Enabling Cache

Figure 3.18. Control Register format



```
// system.c
```

```
system_enable_dcache(void) {
    // See section 3.2.7 in arm1176
    unsigned reg;
    __asm__ volatile("mrc    p15, 0, %0, c1, c0, 0" : "=r"(reg));
    reg |= SYSTEM_DCACHE_ENABLE; //(1<<2)
    __asm__ volatile("mcr    p15, 0, %0, c1, c0, 0" : : "r"(reg));
}
```

Demo: clear

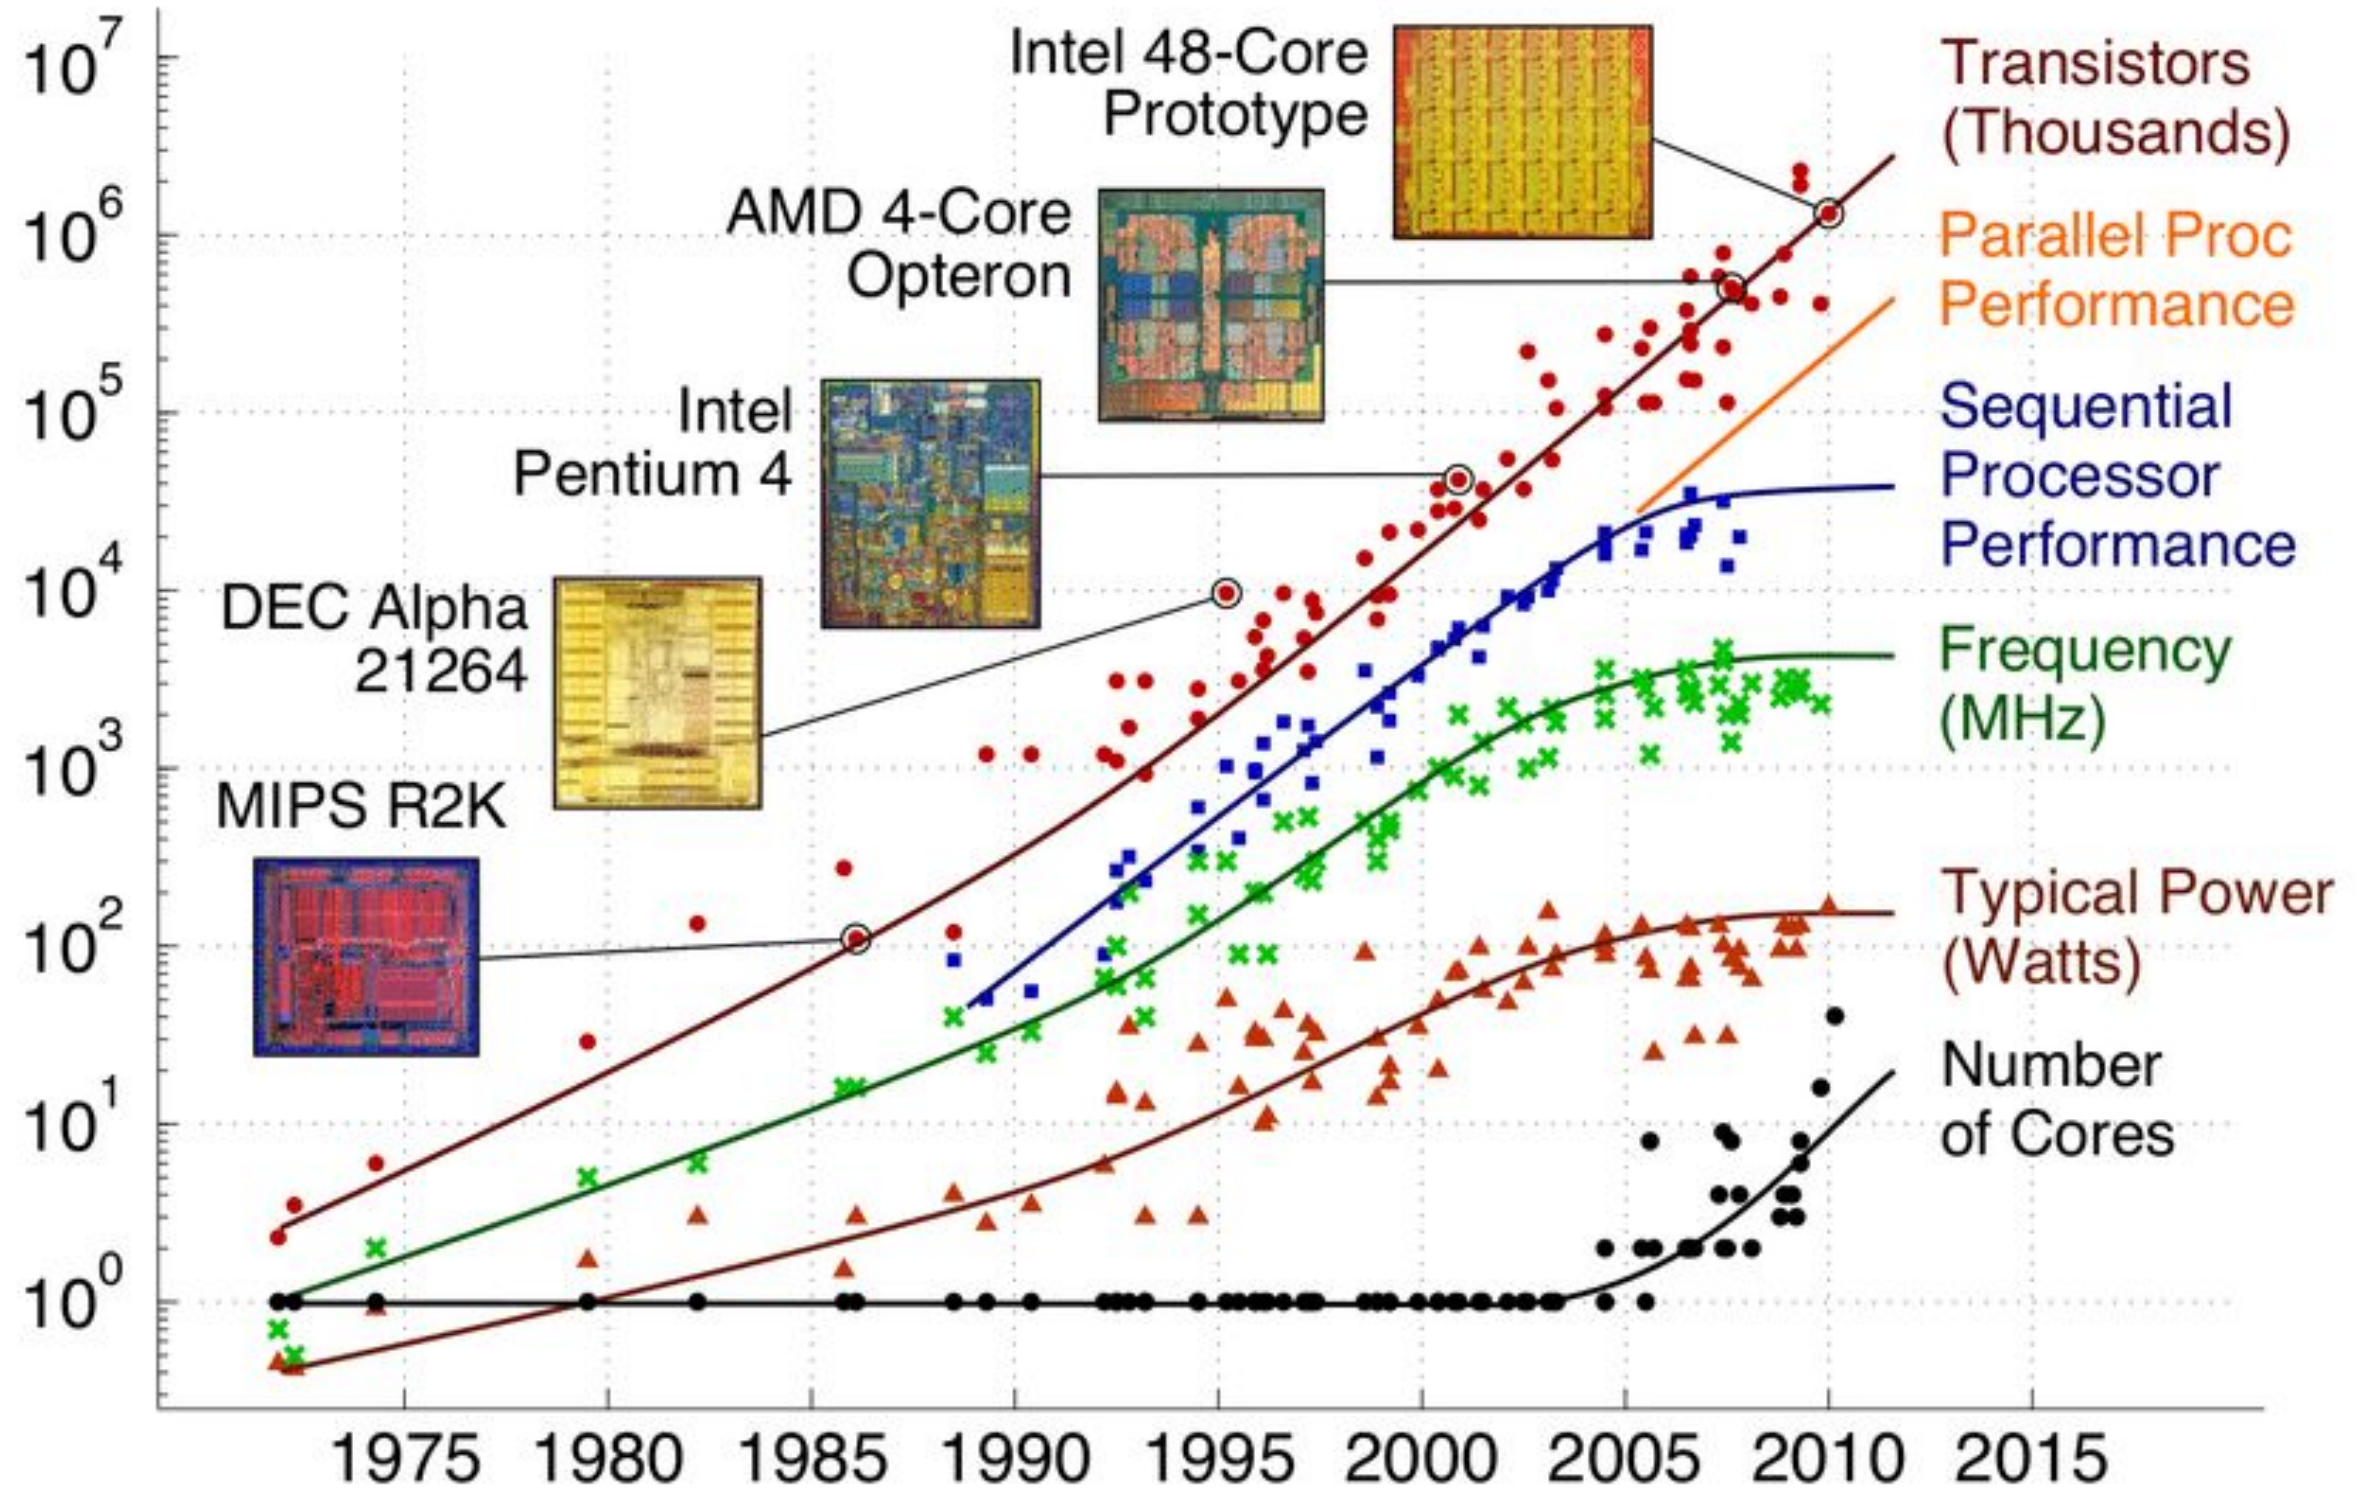
w/ cache enabled

Know your hardware!

YOU DON'T HAVE TO OPTIMIZE CODE^x

IF NEXT YEAR'S PROCESSOR RUNS 2X

Evolution of Processors



Memory System Performance

Processor	Alpha 21164	
Machine	AlphaServer 8200	
Clock Rate	300 MHz	
Memory Performance	Latency	Bandwidth
I Cache (8KB on chip)	6.7 ns (2 clocks)	4800 MB/sec
D Cache (8KB on chip)	6.7 ns (2 clocks)	4800 MB/sec
L2 Cache (96KB on chip)	20 ns (6 clocks)	4800 MB/sec
L3 Cache (4MB off chip)	26 ns (8 clocks)	960 MB/sec
Main Memory Subsystem	253 ns (76 clocks)	1200 MB/sec
Single DRAM component	≈60ns (18 clocks)	≈30–100 MB/sec

[Patterson, David, et al. "A case for intelligent RAM."]

**Moving data between
the cpu and memory is
the bottleneck**

strcpy

```
for (int i = 0; i <= strlen(src); i++) {  
    dst[i] = src[i]  
}
```

All we're doing is loading data from memory into the CPU
and storing it back into memory

Avoiding the Memory Bottleneck

- Raspberry Pi has a **DMA Controller** that allows us to read and write memory without having to go through the processor (avoiding the load/store latency)
- Section 4 of BCM2835-ARM-Peripherals.pdf

Demo: dma

Measuring Performance

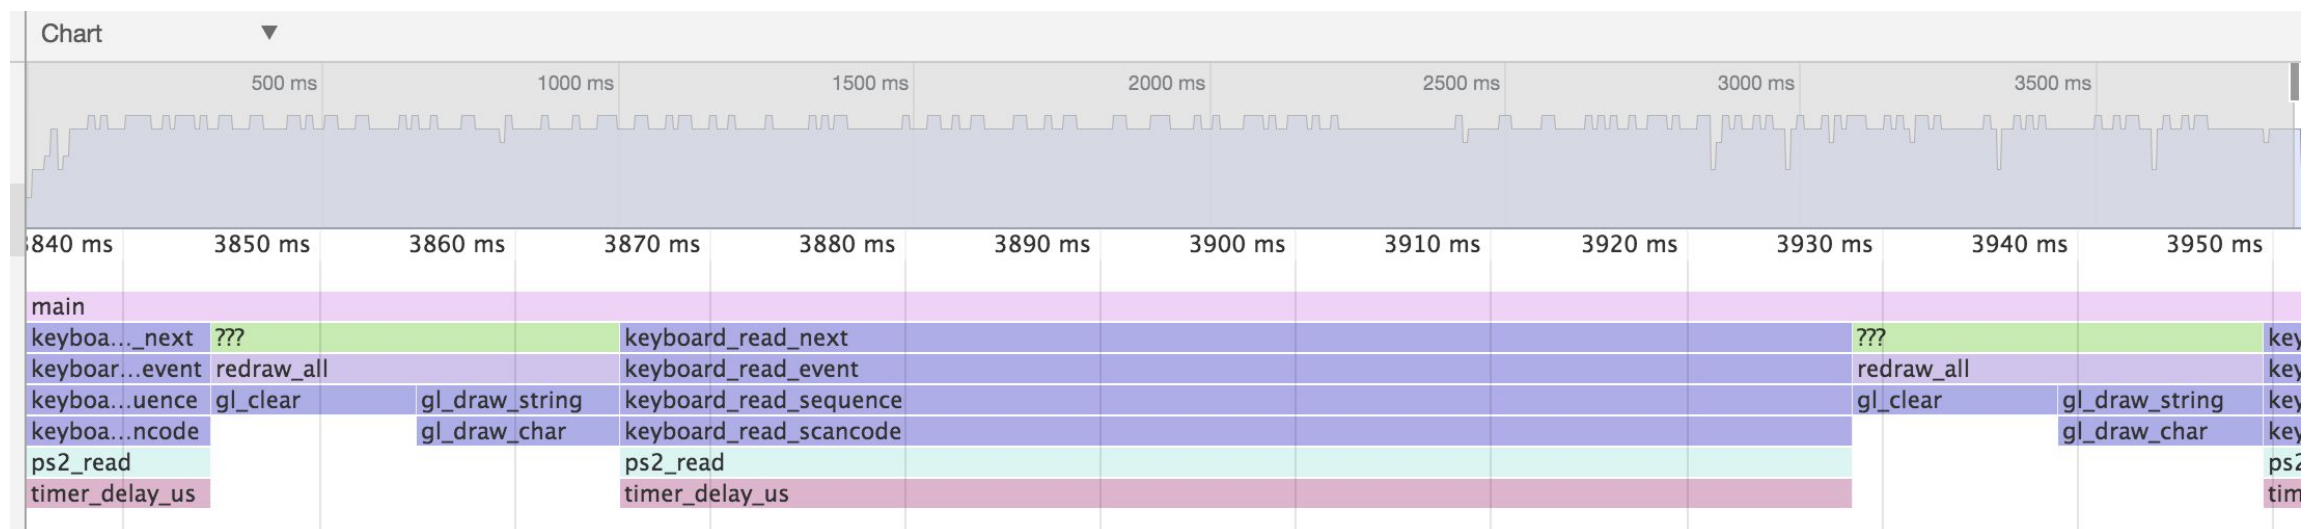
Don't optimize blind

Taking Measurements

- **Hardware interrupts** - gprof.c
- **Code Instrumentation** - timer
- Also: instruction set simulation, OS hooks, performance counters
- Many techniques rely on sampling (statistical profilers) to trade off accuracy for speed

Visualizing Measurements

Chrome Developer Tools



- Output profiling information in standard format (linux *perf*)
- Use **thlorenz/cpuprofilify** to convert into *.cpuprofile* format

Demo: *stackprof*

Performance Metrics

Latency: time to complete an action

- How fast to load data from memory?
- How fast can instruction X complete?

Throughput: actions over a period of time

- How many instructions/second?
- How many compute operations per second? (GOPs, GFLOPS)

Stages in Instruction Execution

How does an instruction execute on a processor?

```
add r0, r1, r2
```

[IF] Instruction Fetch: load instruction from memory based on PC

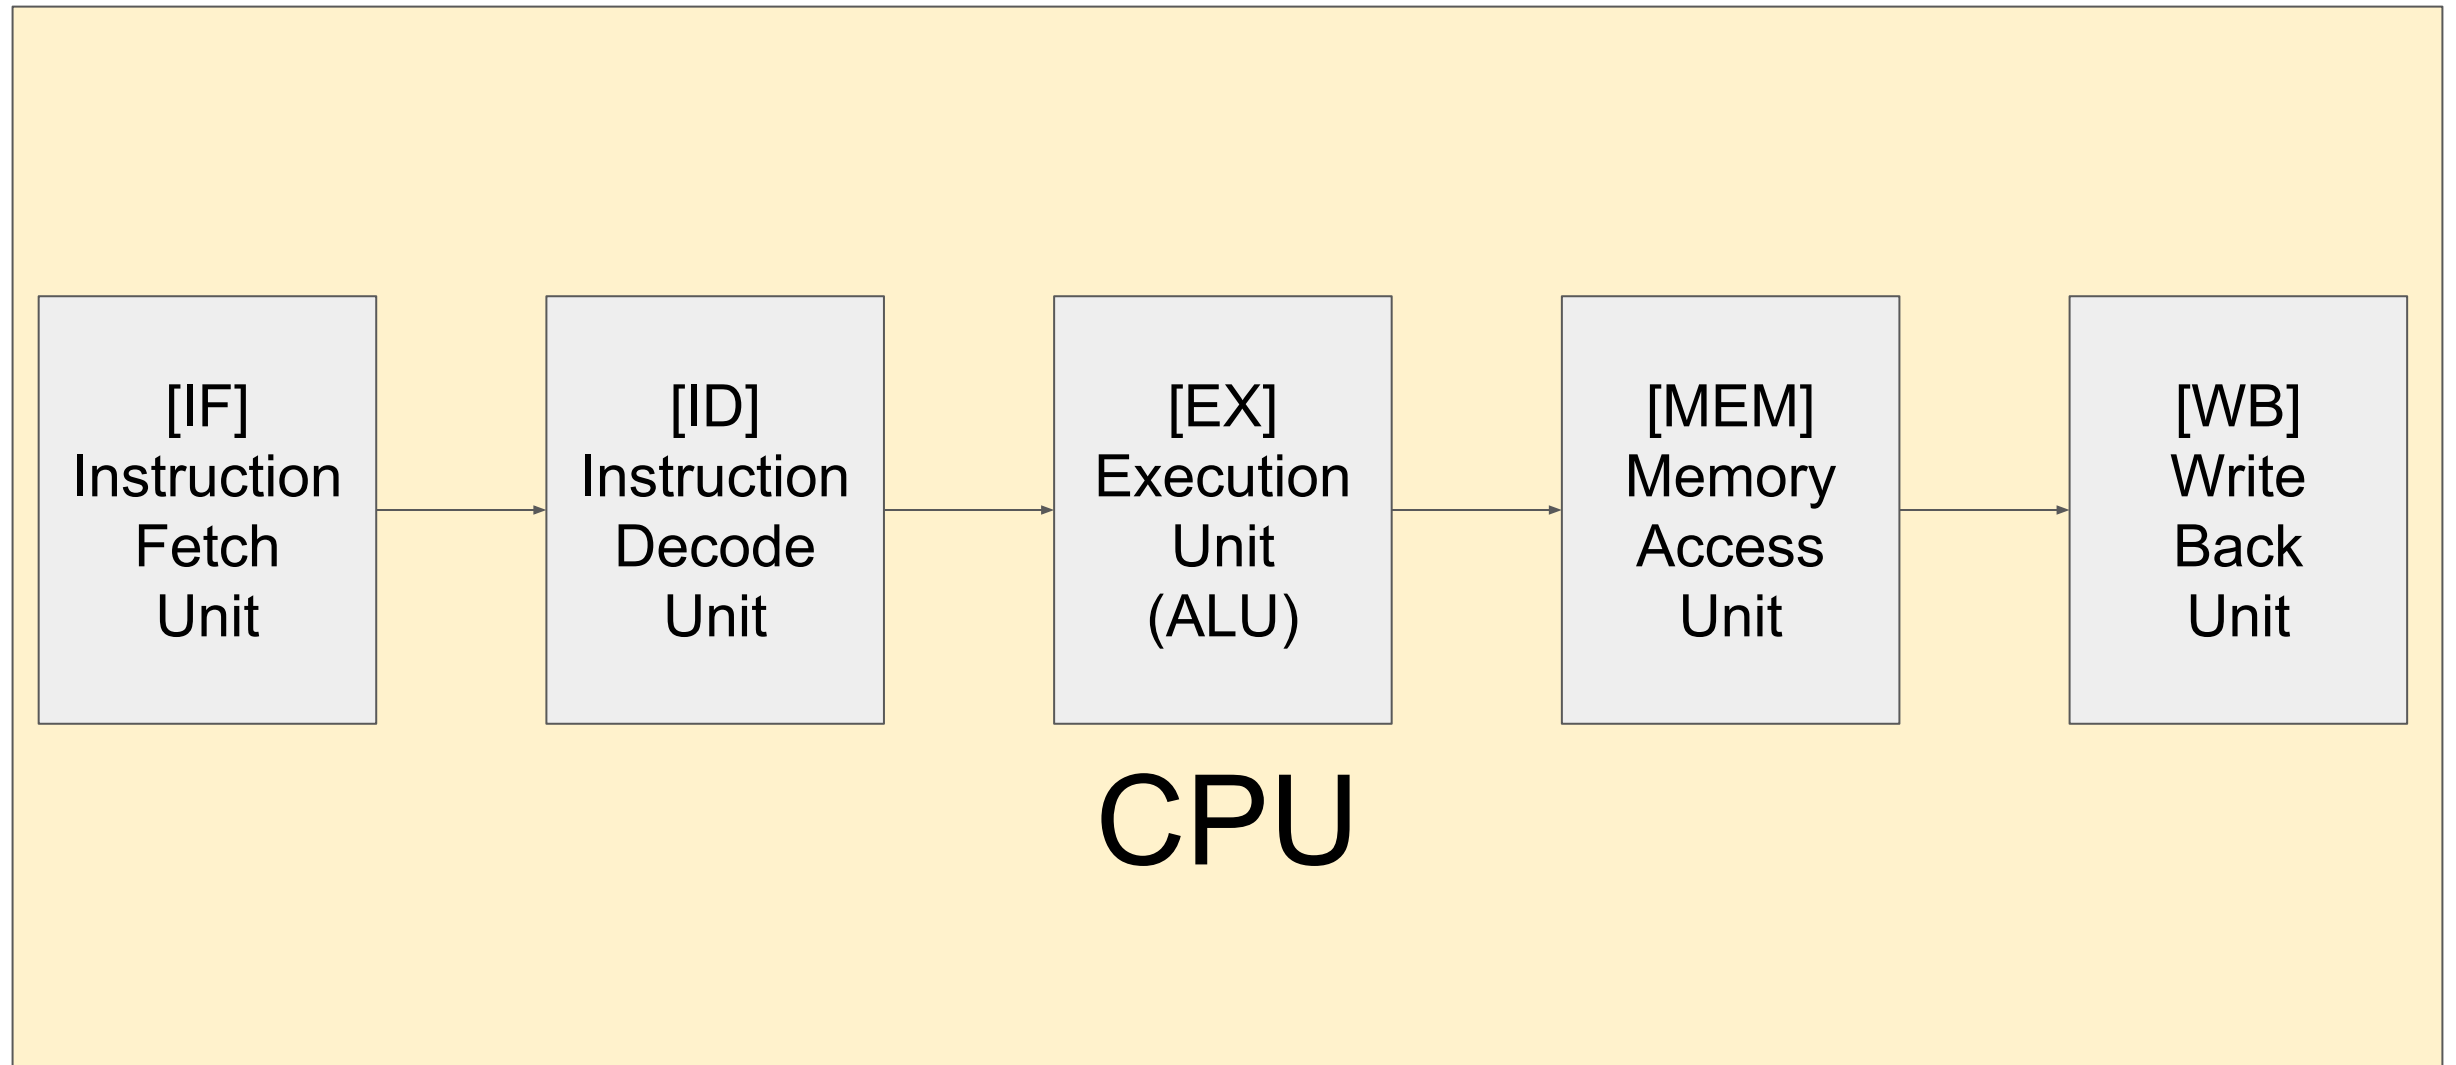
[ID] Instruction Decode: decode the instruction operation, dst, src

[EX] Execute (ALU): execute arithmetic

[MEM] Memory Access: load or store data value to memory

[WB] Write Back: write value back to register

Processor Units



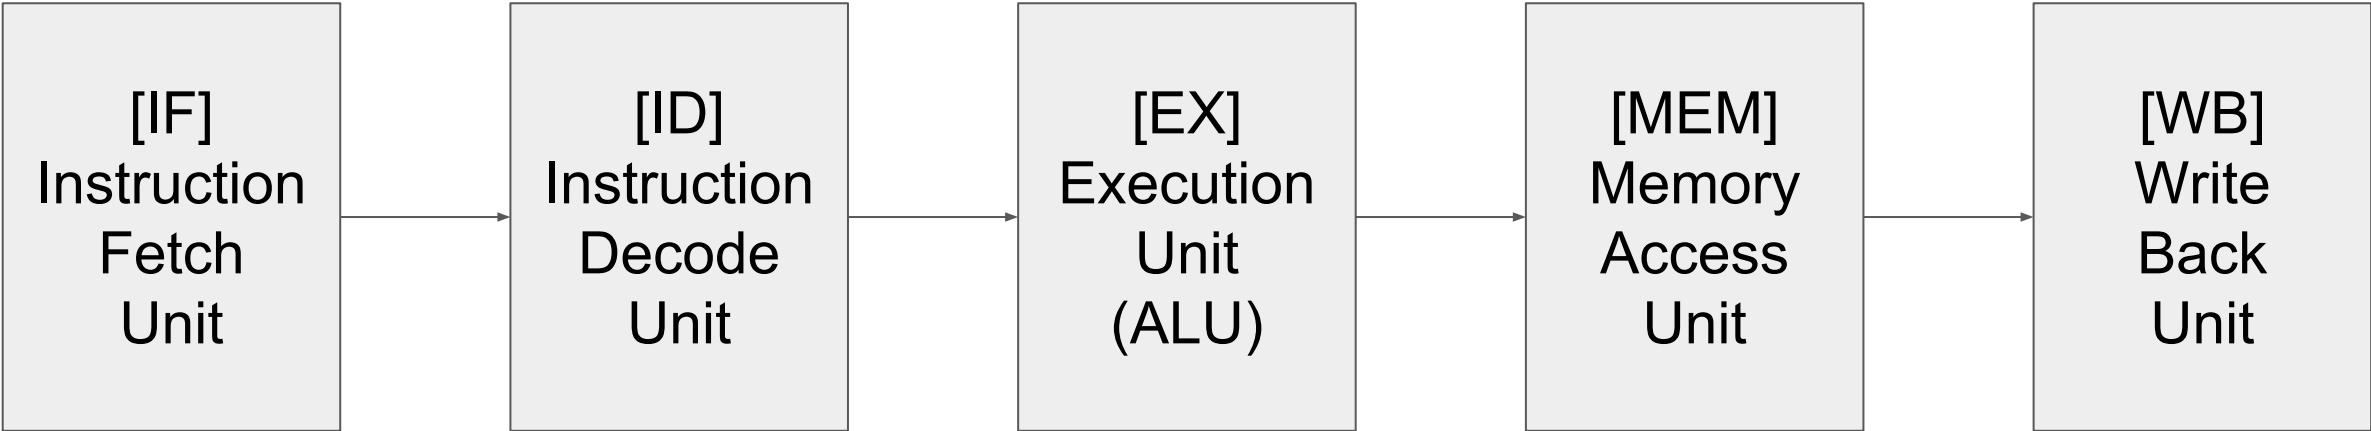
Single-Cycle Processor

```
ldr r0, [r1, r2]
add r0, r0, r2
str r1, [r1, r2]
```

PC = 0

	REG FILE
r0	0x0000
r1	0x8000
r2	0x0004

	MEMORY
0x8000	0xFFFF
0x8004	0x107E



```
ldr r0, [r1, r2]
```

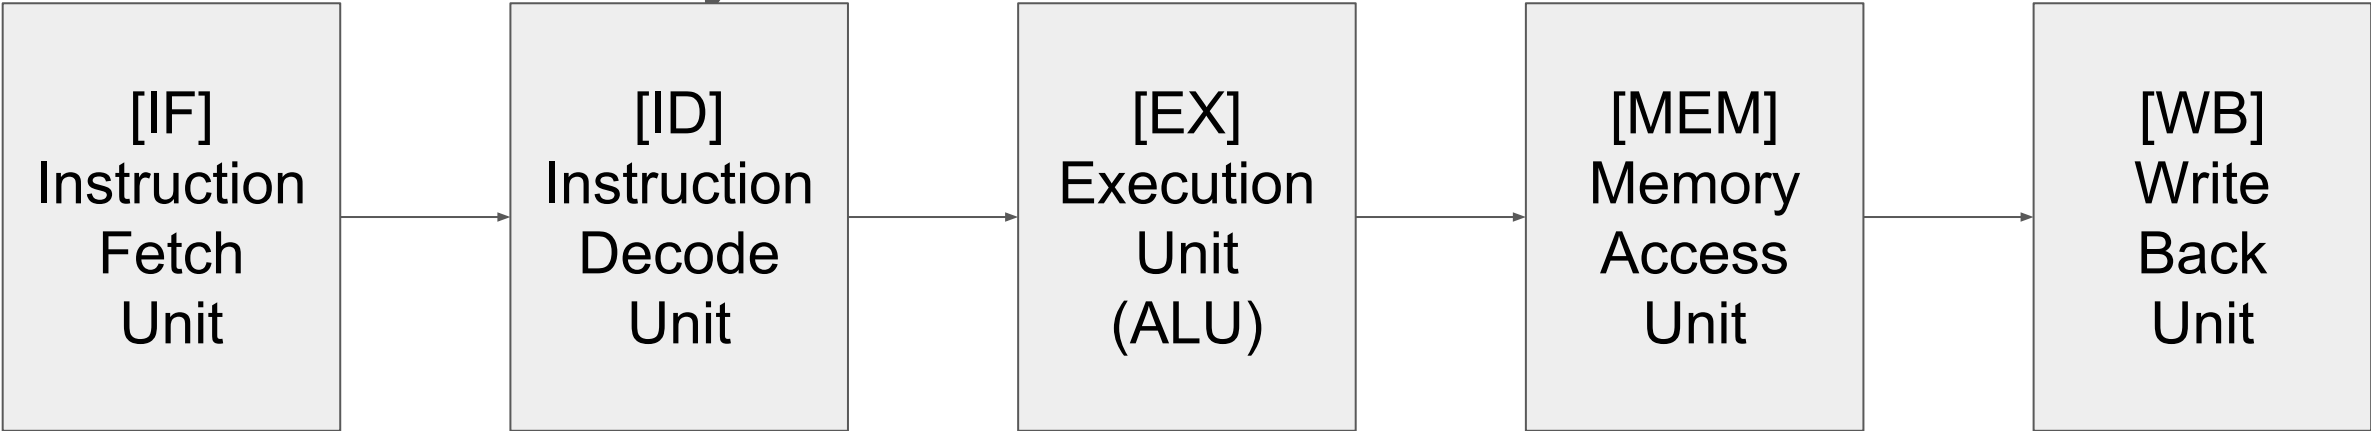


Single-Cycle Processor

```
ldr r0, [r1, r2]
add r0, r0, r2
str r1, [r1, r2]
```

	REG FILE
r0	0x0000
r1	0x8000
r2	0x0004

	MEMORY
0x8000	0xFFFF
0x8004	0x107E



```
OP: ldr
r1 = 0x8000
r2 = 0x0004
```

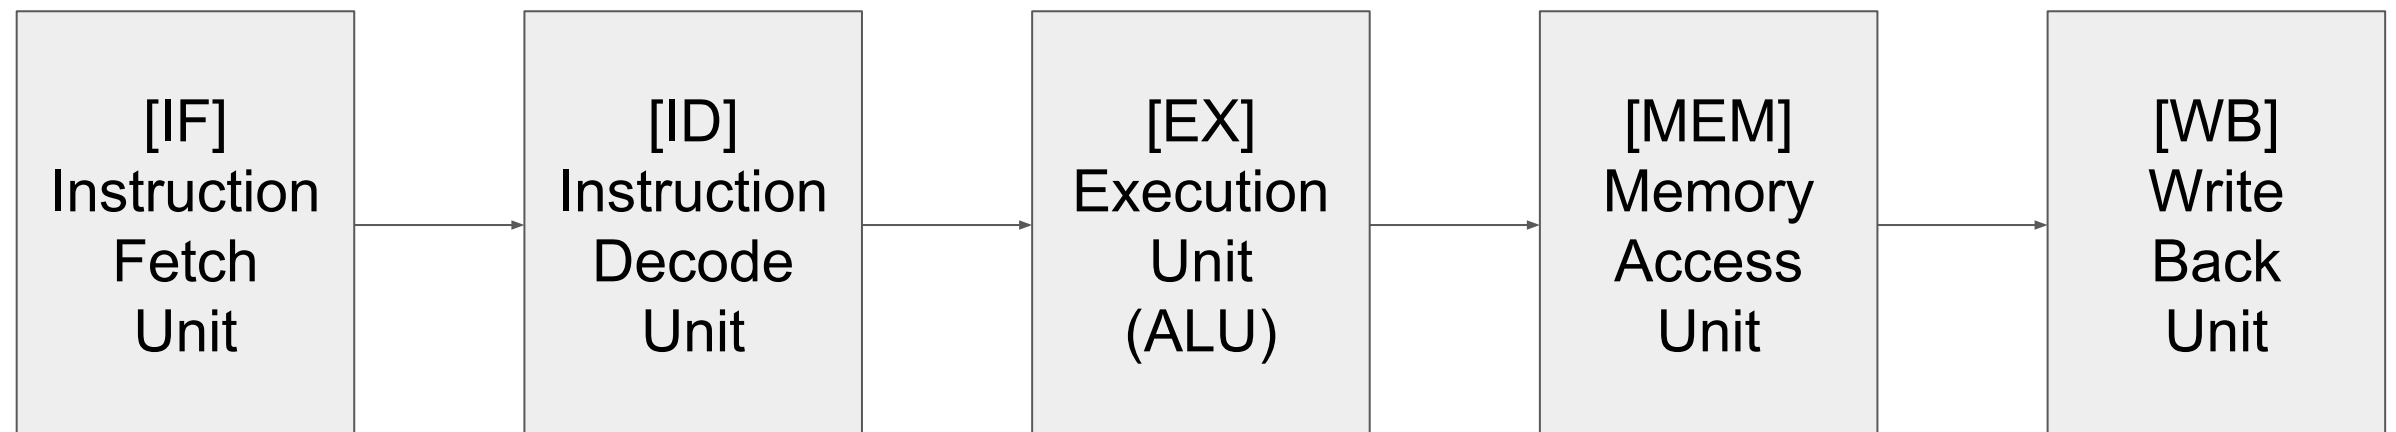


Single-Cycle Processor

```
ldr r0, [r1, r2]
add r0, r0, r2
str r1, [r1, r2]
```

	REG FILE
r0	0x0000
r1	0x8000
r2	0x0004

	MEMORY
0x8000	0xFFFF
0x8004	0x107E



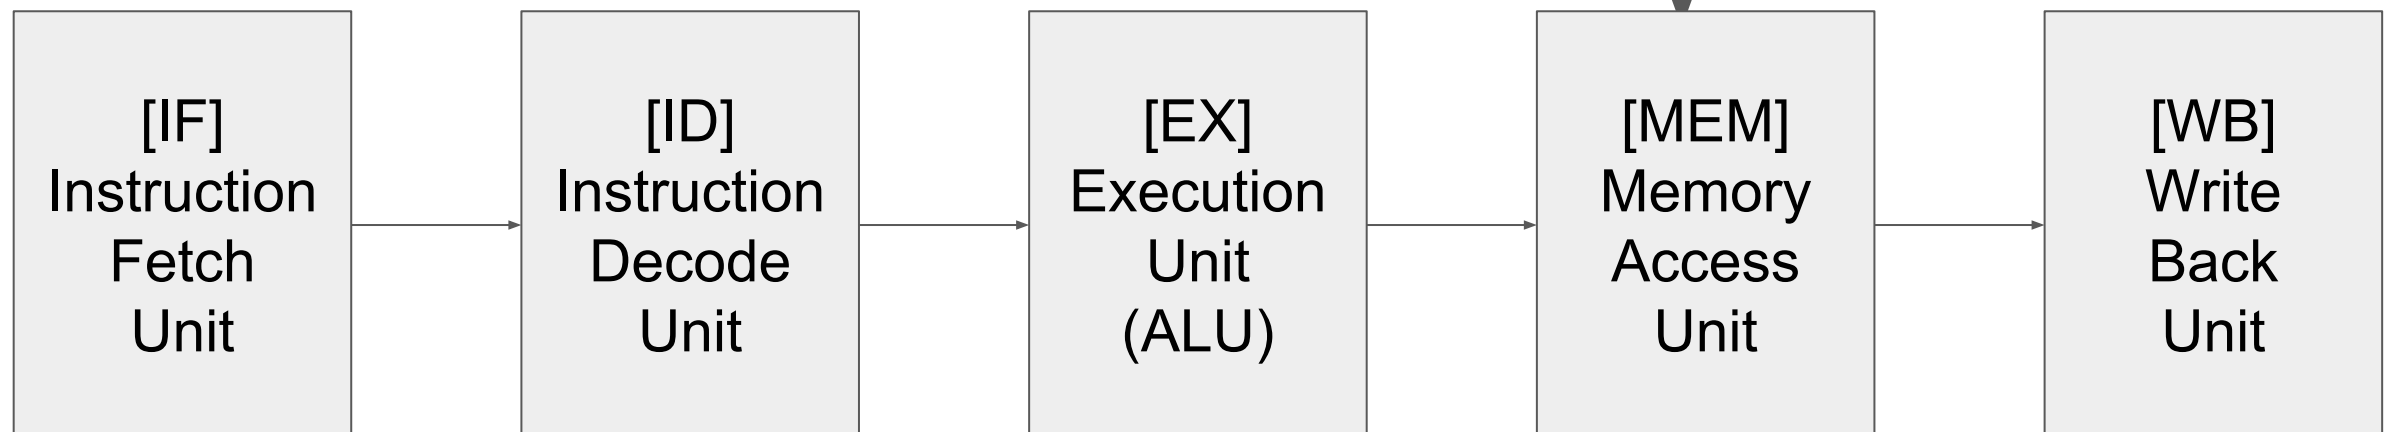
OP: ldr
ALU_val
= r1 + r2
= 0x8000 + 0x4
= 0x8004

Single-Cycle Processor

```
ldr r0, [r1, r2]
add r0, r0, r2
str r1, [r1, r2]
```

	REG FILE
r0	0x0000
r1	0x8000
r2	0x0004

	MEMORY
0x8000	0xFFFF
0x8004	0x107E



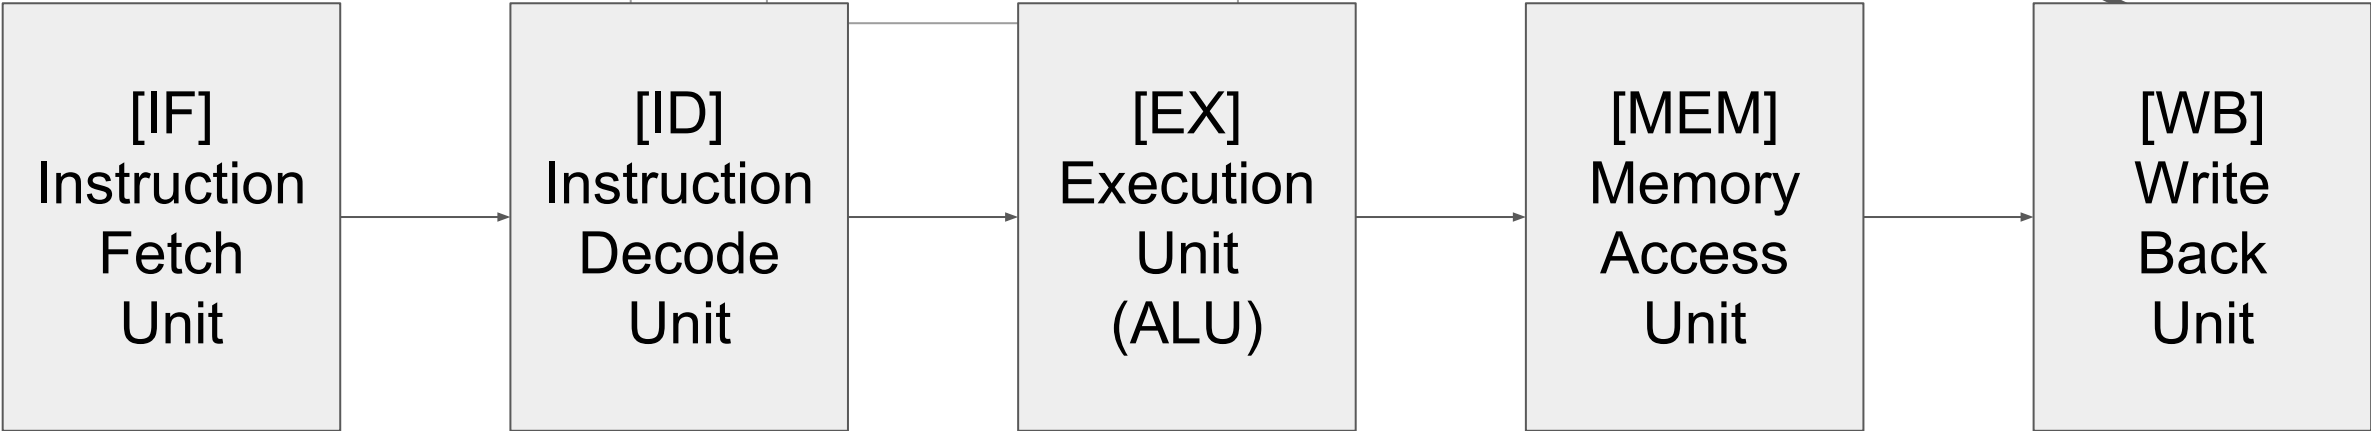
OP: ldr
MEM_val
= MEM[ALU_val]
= MEM[0x8004]
= 0x107E

Single-Cycle Processor

```
ldr r0, [r1, r2]
add r0, r0, r2
str r1, [r1, r2]
```

	REG FILE
r0	0x0000 0x107E
r1	0x8000
r2	0x0004

	MEMORY
0x8000	0xFFFF
0x8004	0x107E



```
r0 <- MEM_val
r0 <- 0x107E
```



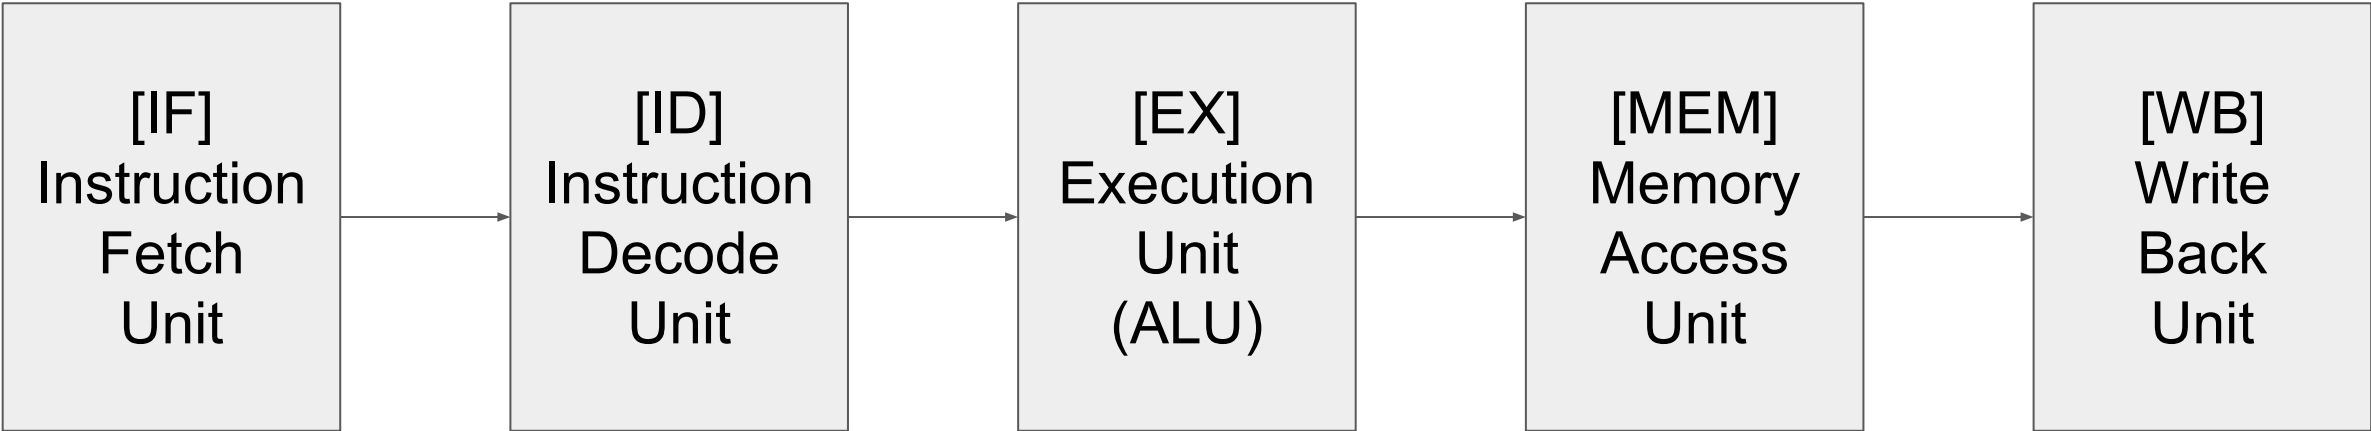
Single-Cycle Processor

```
ldr r0, [r1, r2]
add r0, r0, r2
str r1, [r1, r2]
```

PC = 4

	REG FILE
r0	0x107E
r1	0x8000
r2	0x0004

	MEMORY
0x8000	0xFFFF
0x8004	0x107E



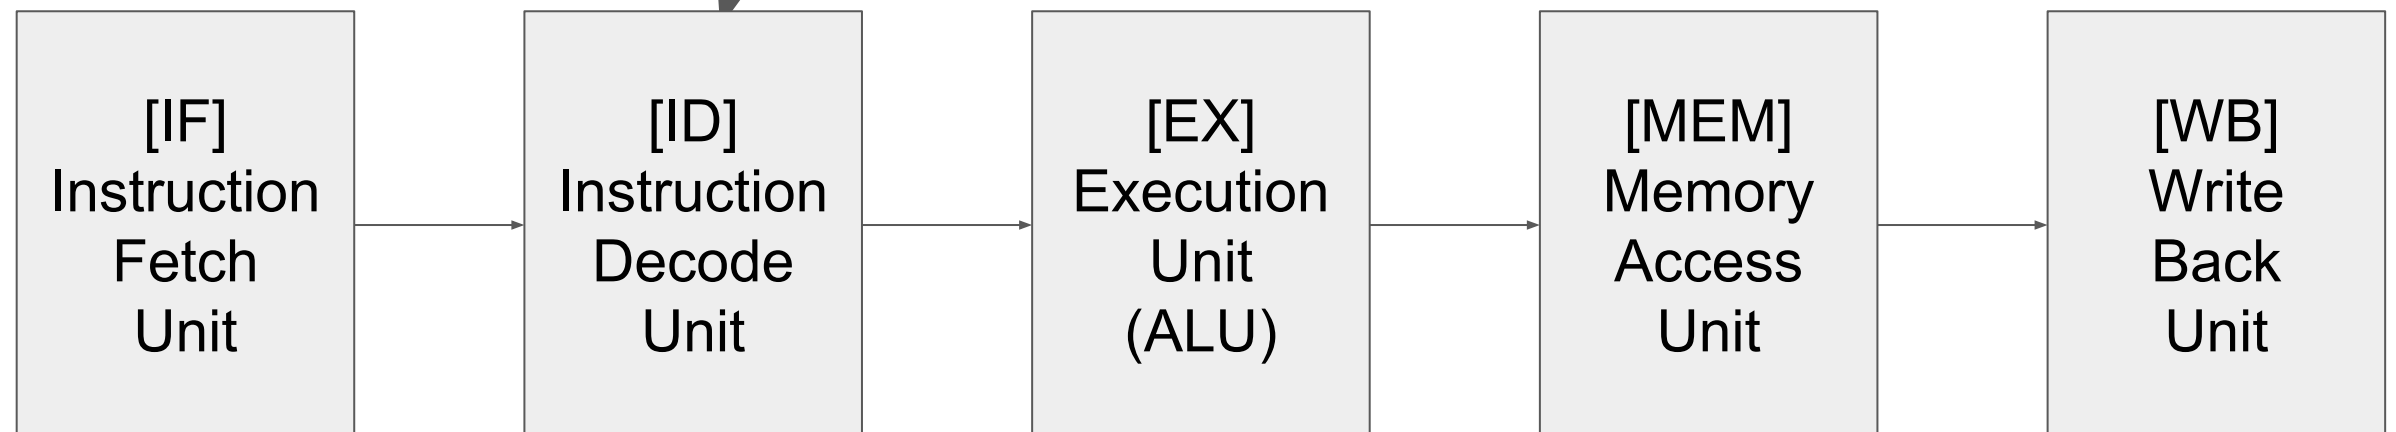
```
add r0, r0, r2
```

Single-Cycle Processor

```
ldr r0, [r1, r2]  
add r0, r0, r2  
str r1, [r1, r2]
```

	REG FILE
r0	0x107E
r1	0x8000
r2	0x0004

	MEMORY
0x8000	0xFFFF
0x8004	0x107E



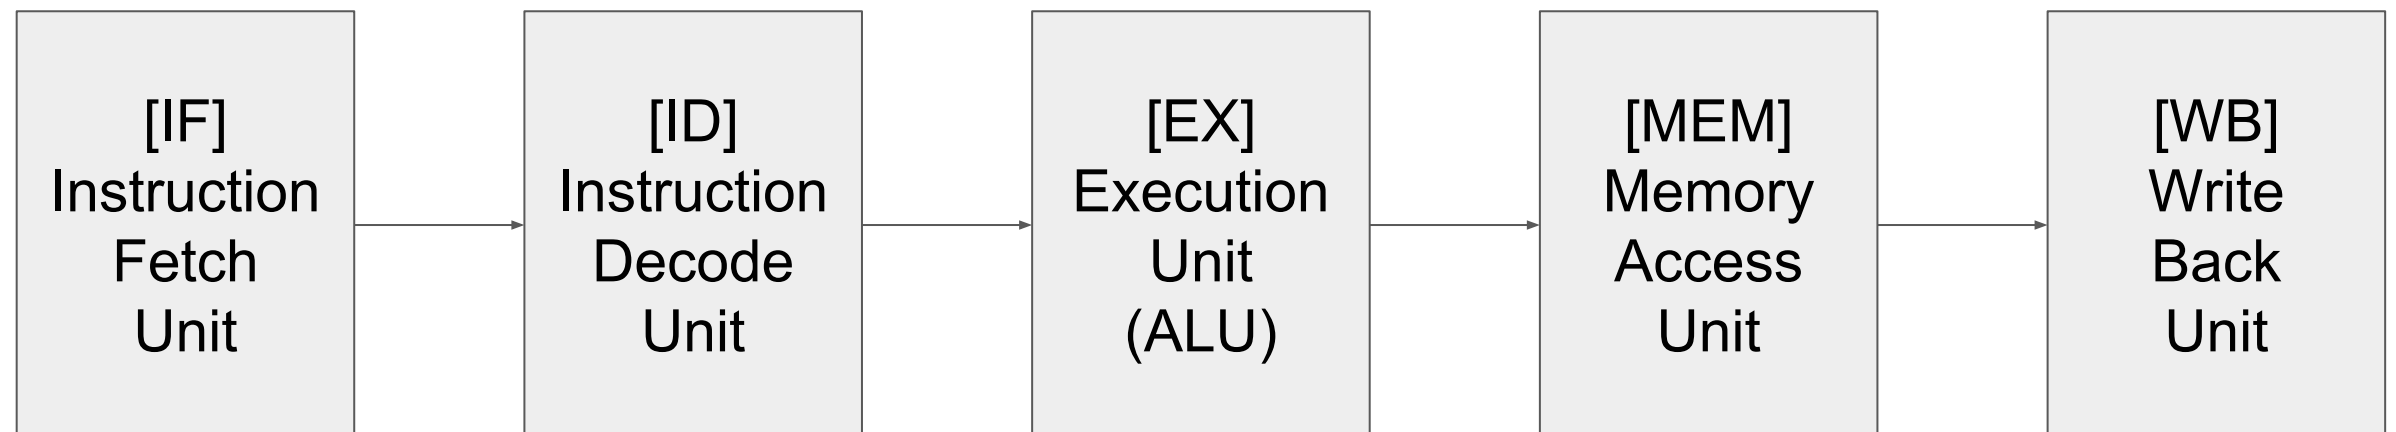
OP: **add**
r0 = 0x107E
r2 = 0x0004

Single-Cycle Processor

```
ldr r0, [r1, r2]  
add r0, r0, r2  
str r1, [r1, r2]
```

	REG FILE
r0	0x107E
r1	0x8000
r2	0x0004

	MEMORY
0x8000	0xFFFF
0x8004	0x107E



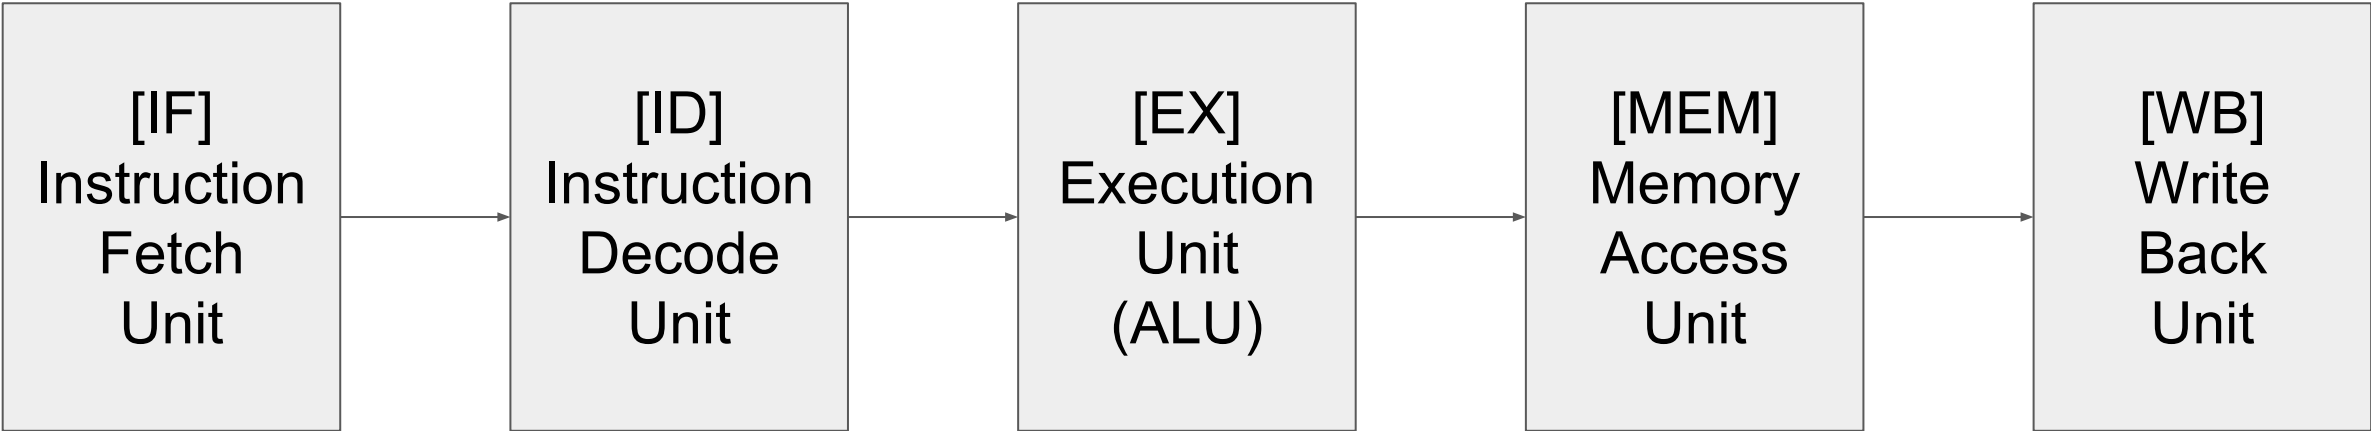
OP: ldr
ALU_val
= r0 + r2
= 0x107E + 0x4
= 0x1083

Single-Cycle Processor

```
ldr r0, [r1, r2]
add r0, r0, r2
str r1, [r1, r2]
```

	REG FILE
r0	0x107E
r1	0x8000
r2	0x0004

	MEMORY
0x8000	0xFFFF
0x8004	0x107E



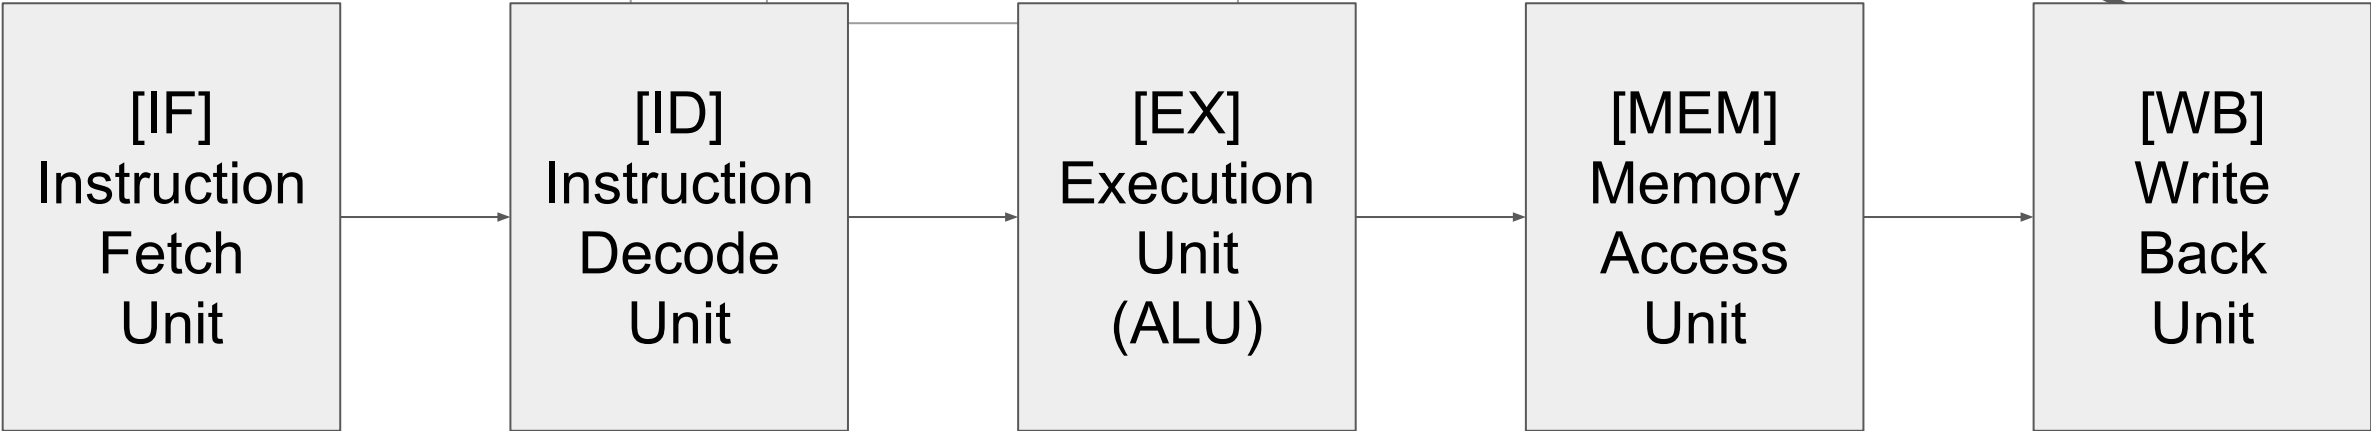
OP: add

Single-Cycle Processor

```
ldr r0, [r1, r2]
add r0, r0, r2
str r1, [r1, r2]
```

	REG FILE
r0	0x107E 0x1083
r1	0x8000
r2	0x0004

	MEMORY
0x8000	0xFFFF
0x8004	0x107E



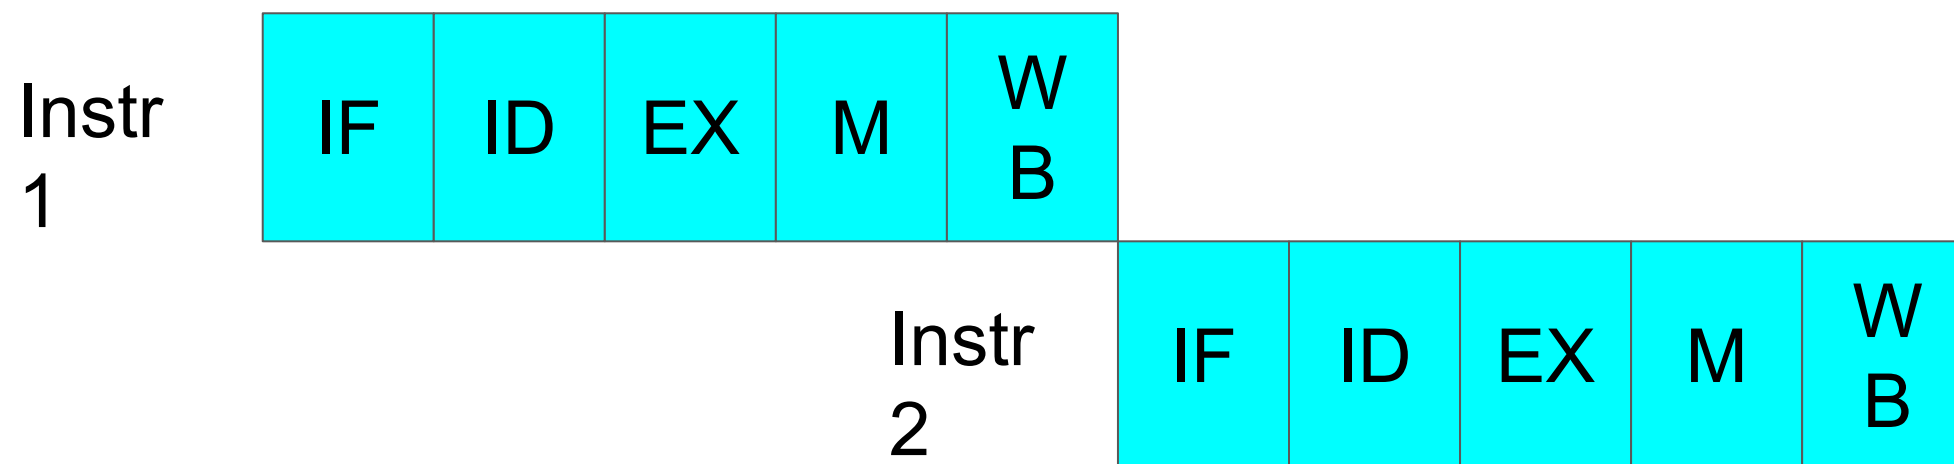
```
r0 <- ALU_val
r0 <- 0x1083
```

Disadvantages of Single Cycle CPU

CPU resources not fully utilized

5 resource units: IF, ID, EX, MEM, WB

Only 1 unit used at a time → 20% utilization



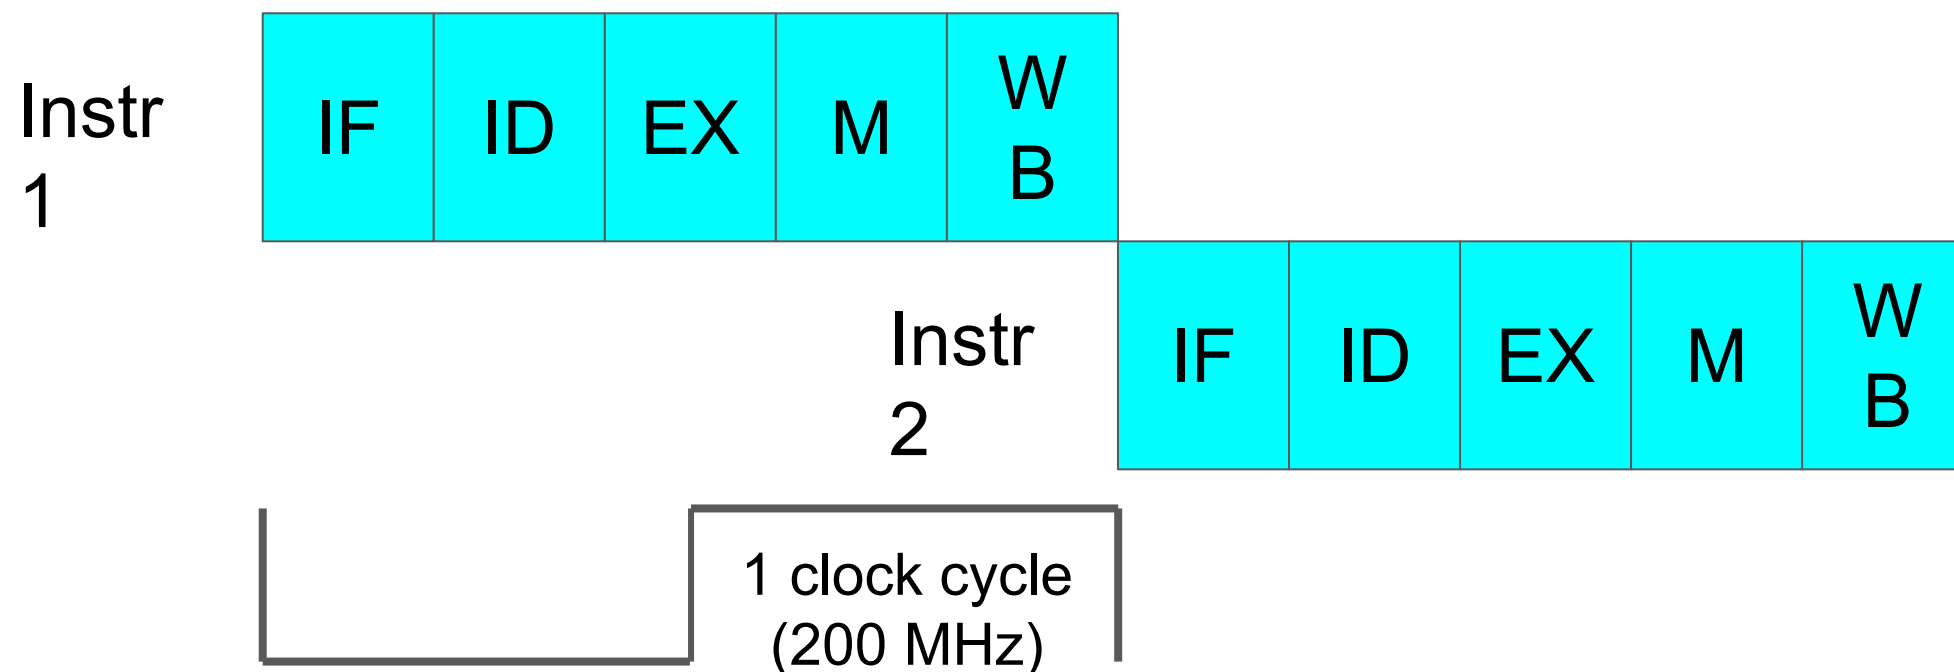
Disadvantages of Single Cycle CPU

CPU clock cycle is *slow* & 1 instruction per clock cycle

@ 200 MHz clock \rightarrow 200,000,000 instructions per second

Latency: 5 ns per instruction

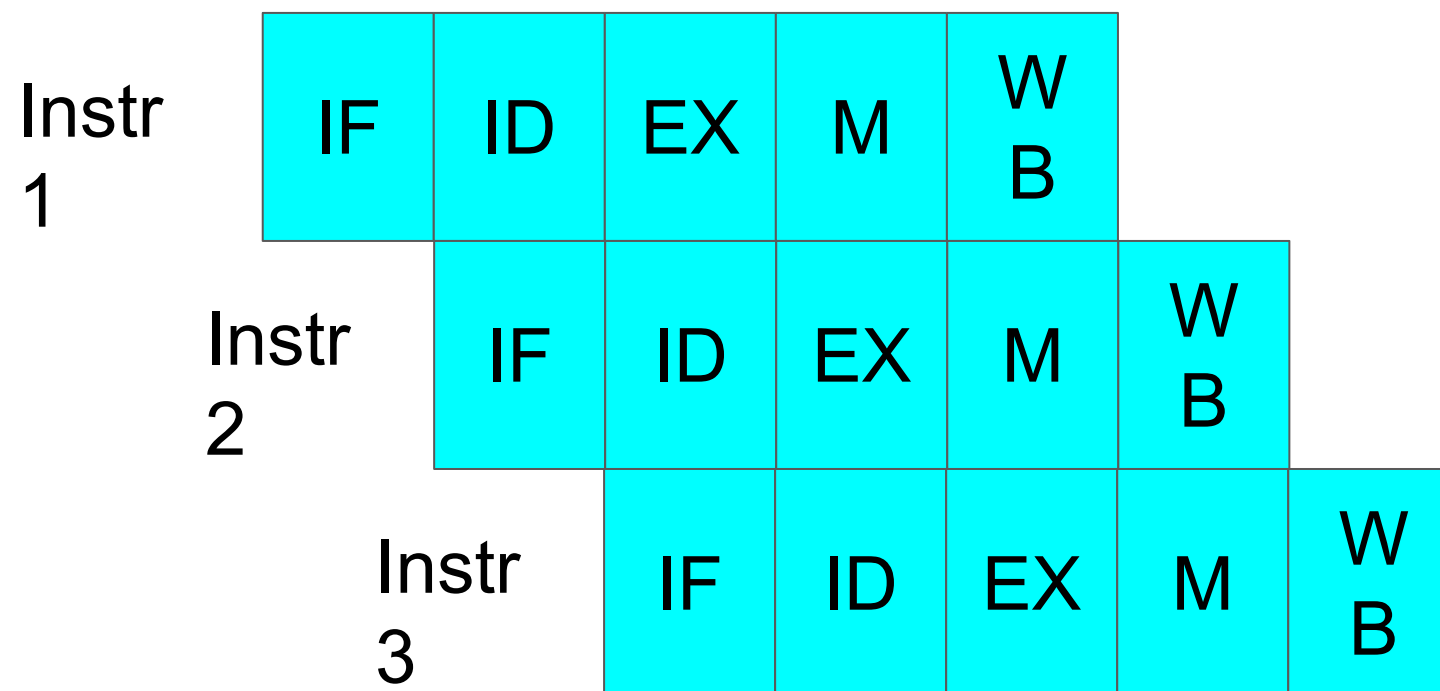
Throughput: 200 M-instr/sec



Fixing the Single Cycle CPU

Pipelining

- Each instruction can start using a free unit before the previous instruction is completed
- Increases *throughput* and CPU *utilization* (100%)

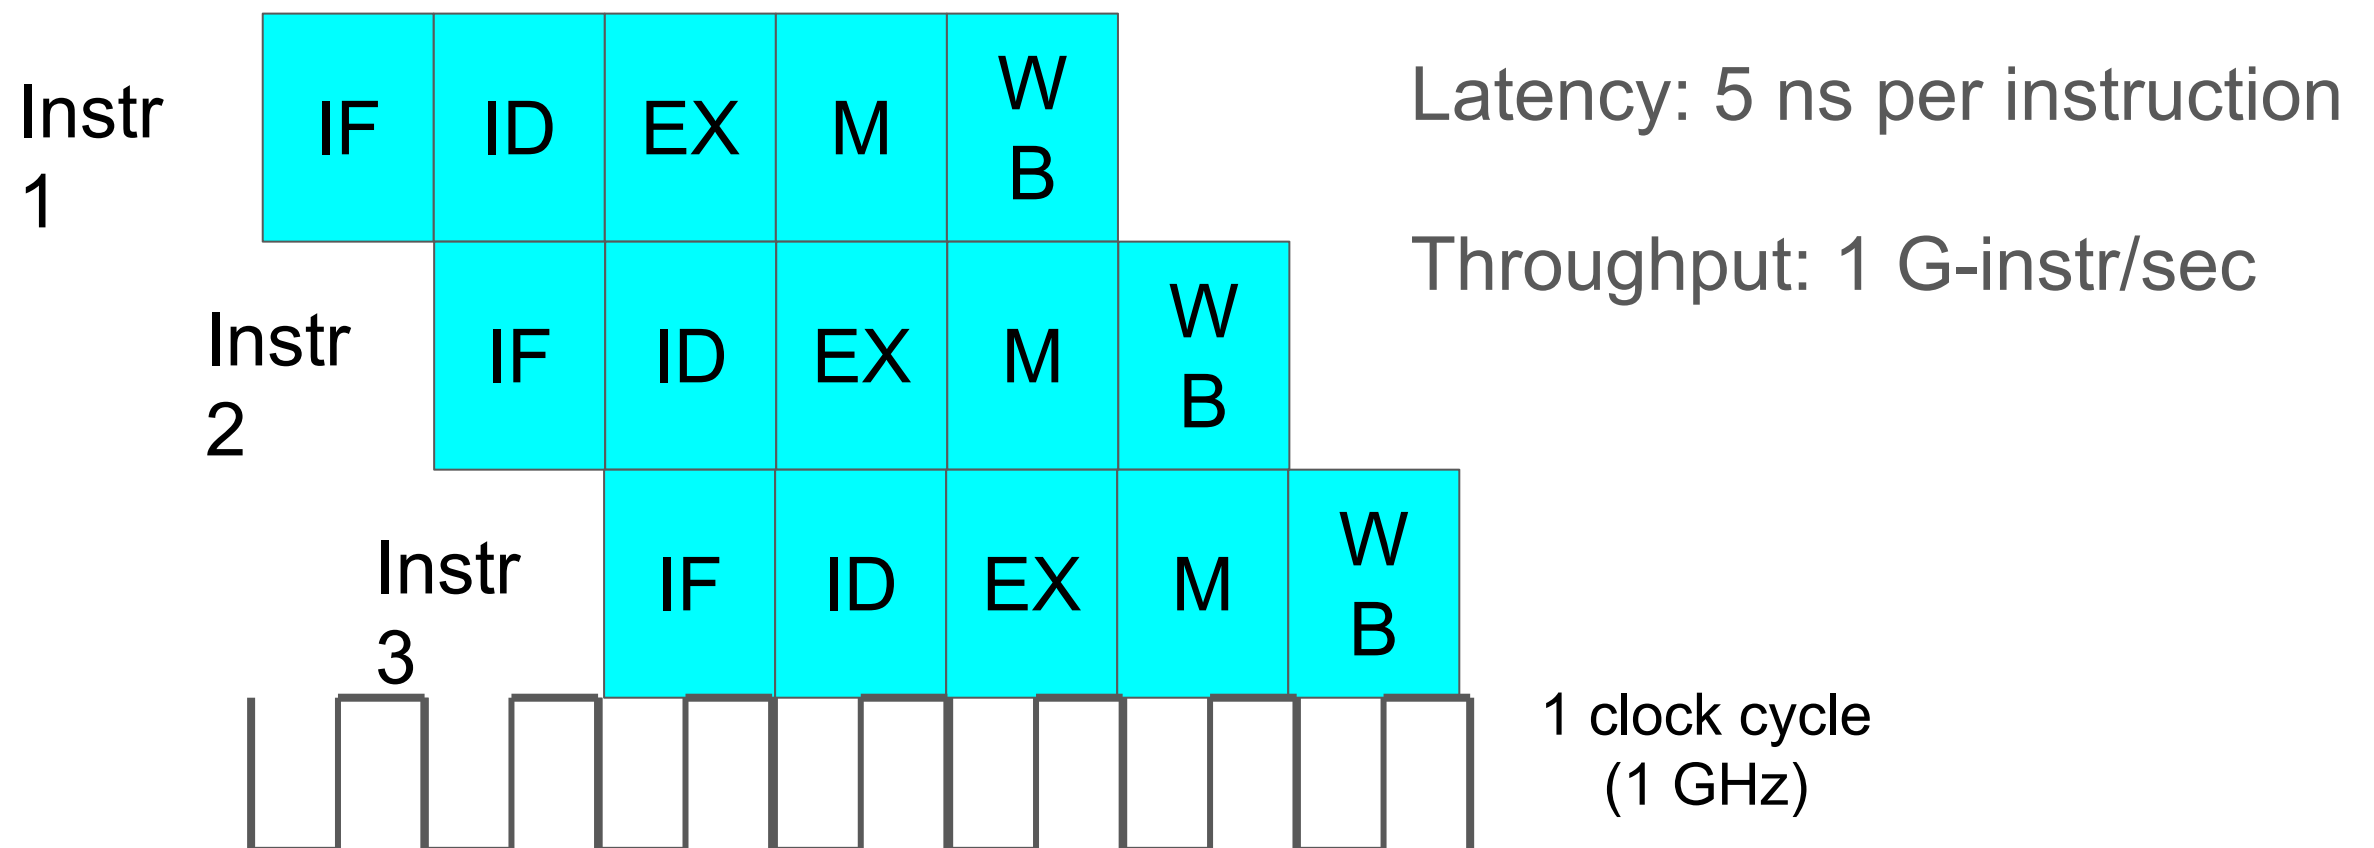


Fixing the Single Cycle CPU

Multi-Cycle

- Each instruction takes more clock cycles, but CPU can be clocked faster
- 1 clock cycle per functional unit ($200 \text{ Mhz} \times 5 = 1 \text{ Ghz}$)

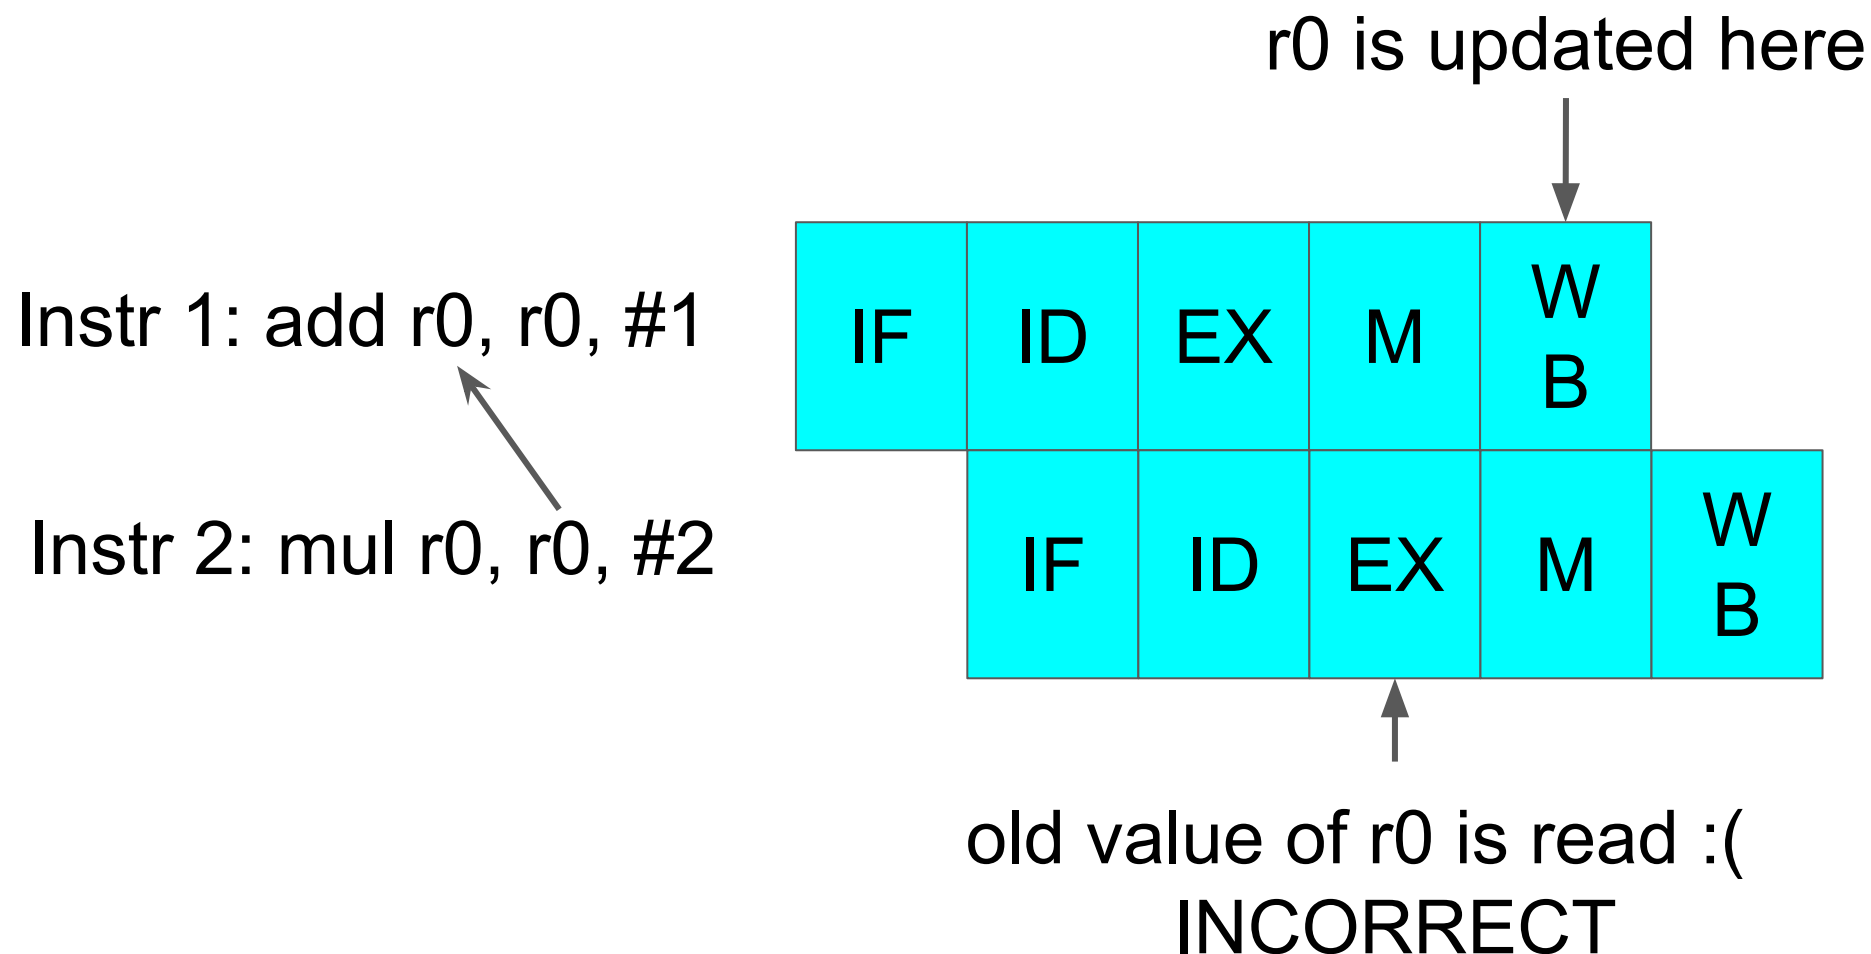
@ 1 GHz clock \rightarrow 1,000,000,000 instructions per second



Data Hazards

Instruction 2 in the pipeline depends on result of Instruction 1.

If Instruction 2 reads the register file (EX stage) before Instruction 1 has updated it (WB stage).

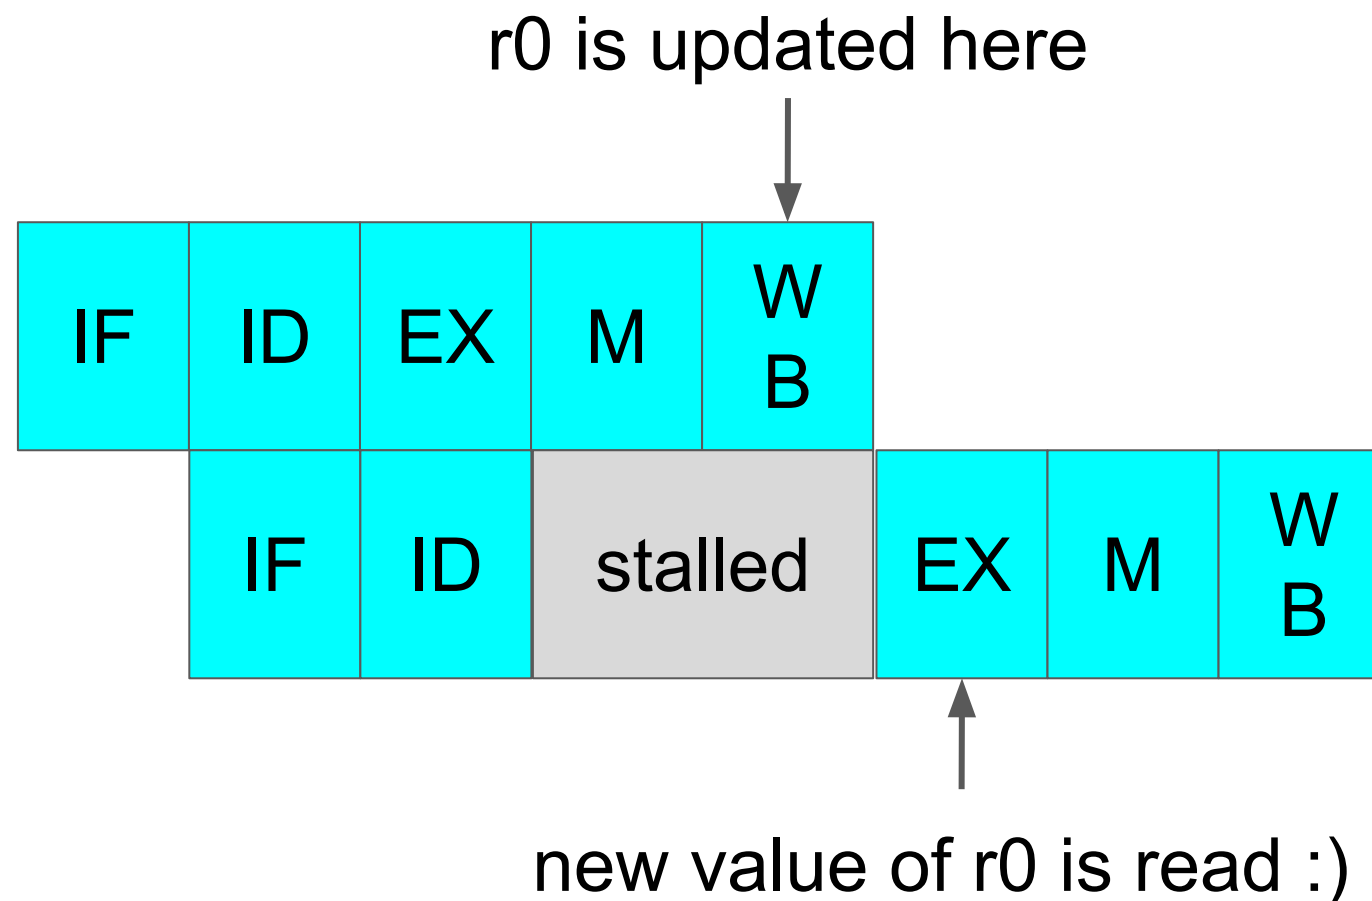


Data Hazards

Solution 1: *Detect* the dependency & *stall* until previous WB is complete.

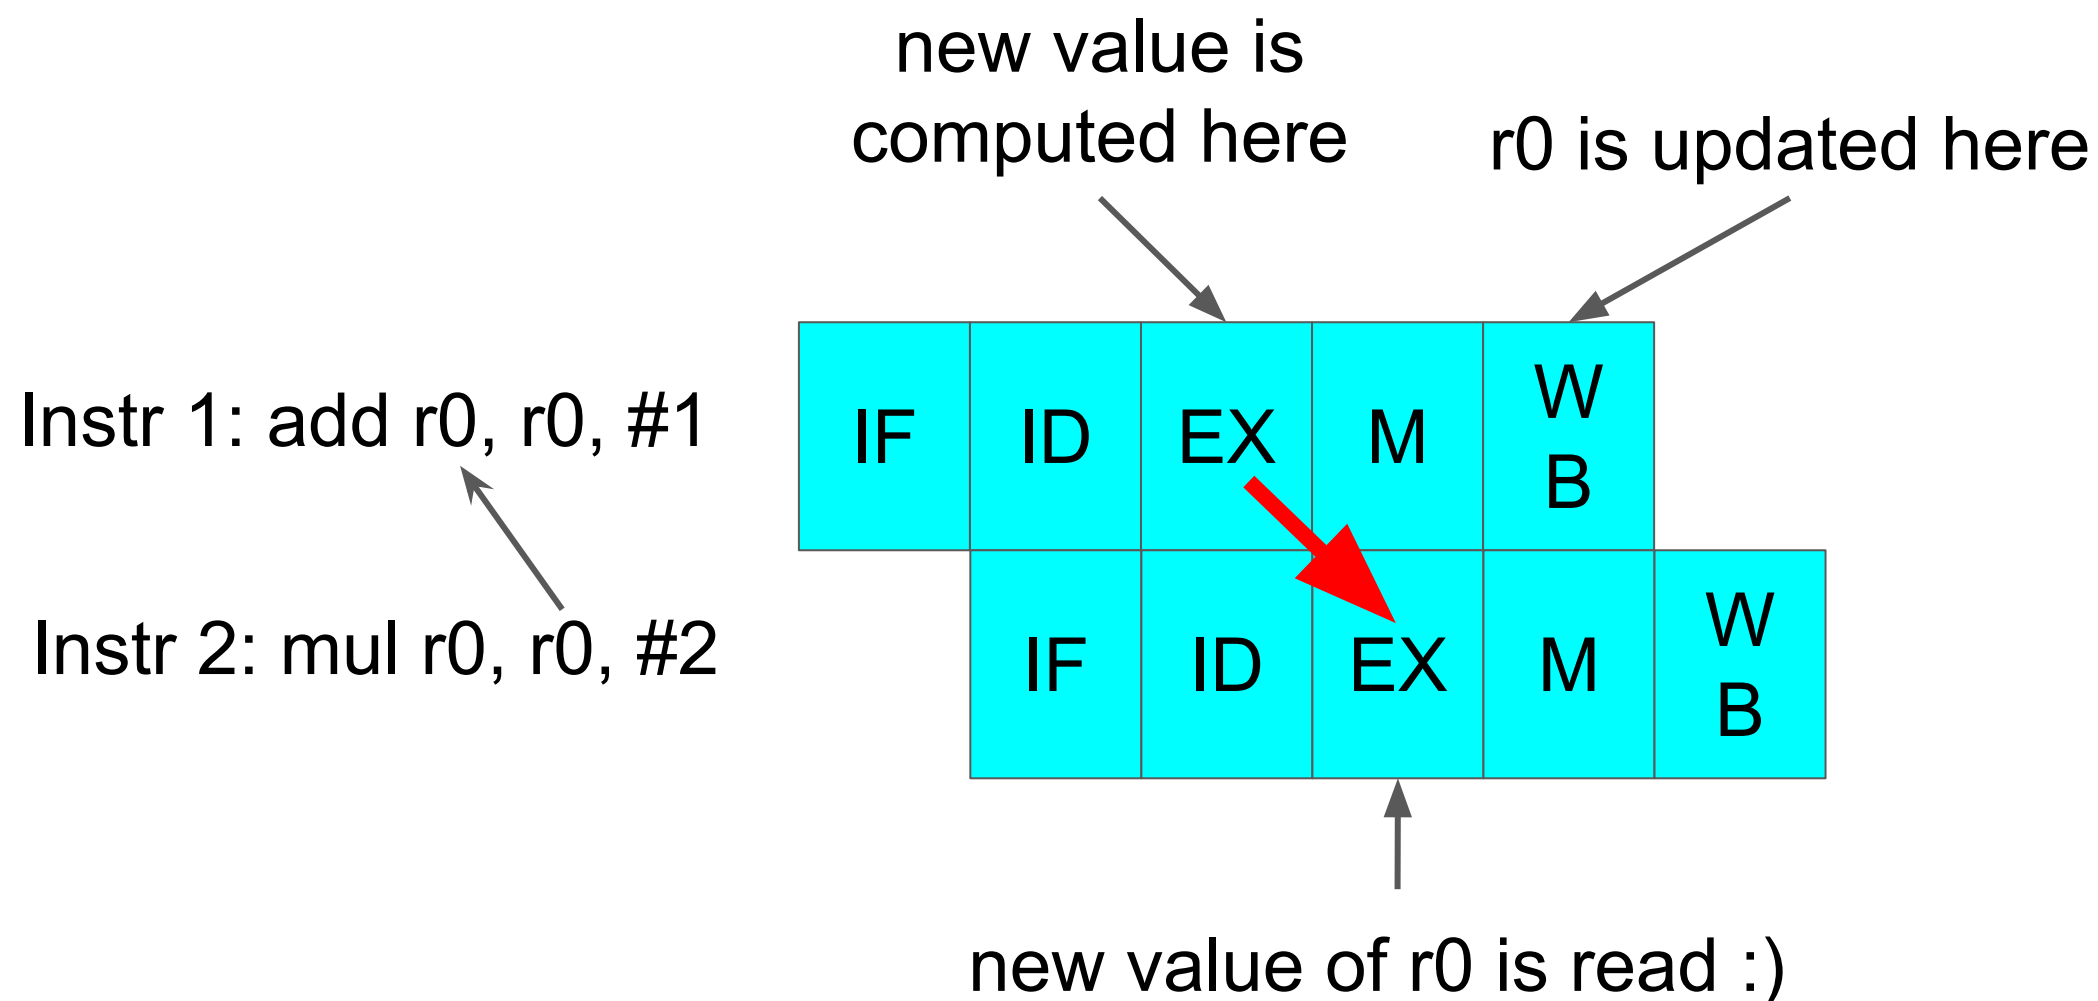
Instr 1: add r0, r0, #1

Instr 2: mul r0, r0, #2



Data Hazards

Solution 2: *Detect* the dependency & *forward* the result to next instruction.



Control Hazards

Which instruction to fetch after a branch instruction?

```
TAKE_LOOP:
    add r0, r0, #1
    cmp r0, r1      // r0=1, r1=10
    bne TAKE_LOOP
ldr r0, r3          //not taken
```

What stage of the pipeline do we compute where to branch?

Hint: PC + offset

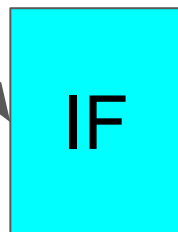
Control Hazards

Which instruction to fetch after a branch instruction?

```
TAKE_LOOP:
    add r0, r0, #1
    cmp r0, r1      // r0=1, r1=10
    bne TAKE_LOOP
    ldr r0, r3      //not taken
```

Naively increment
next instruction fetch
(PC = PC + 4)

Instr 1: bne TAKE_LOOP



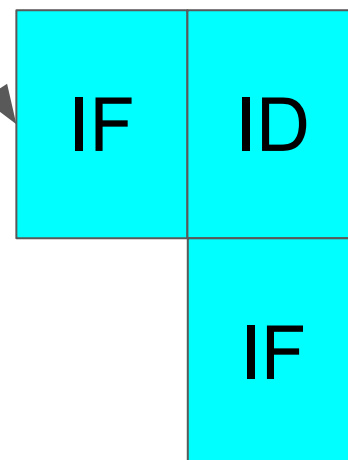
Control Hazards

Which instruction to fetch after a branch instruction?

```
TAKE_LOOP:  
    add r0, r0, #1  
    cmp r0, r1      // r0=1, r1=10  
    bne TAKE_LOOP  
    ldr r0, r3      //not taken
```

Naively increment
next instruction fetch
(PC = PC + 4)

Instr 1: bne TAKE_LOOP



Instr 2: ldr r0, r3

Wrong instruction fetched :(
INCORRECT

Control Hazards

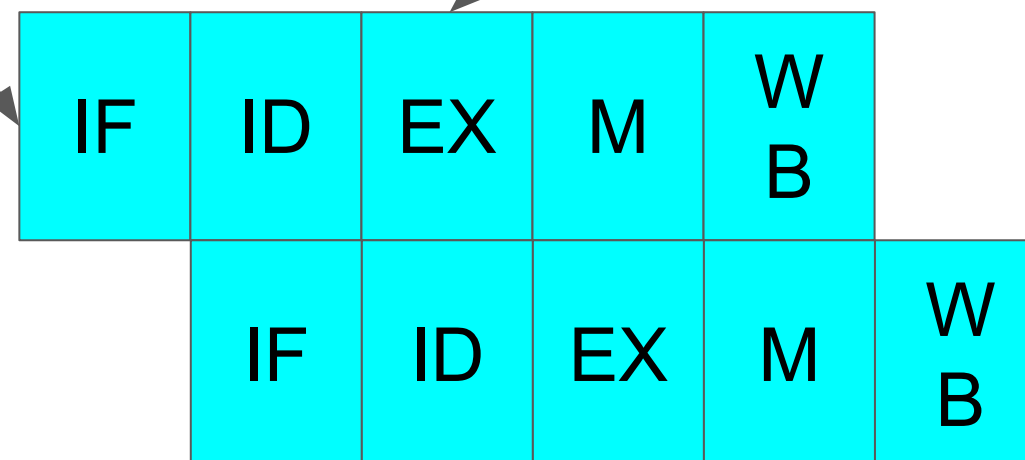
Which instruction to fetch after a branch instruction?

```
TAKE_LOOP:
    add r0, r0, #1
    cmp r0, r1      // r0=1, r1=10
    bne TAKE_LOOP
    ldr r0, r3      //not taken
```

Naively increment
next instruction fetch
($PC = PC + 4$)

Instruction to branch to:
 $PC = PC - 12$

Instr 1: bne TAKE_LOOP



Instr 2: ldr r0, r3

Wrong instruction fetched :(
INCORRECT

Control Hazards

Solution 1: *Stall* until branch address is resolved.

TAKE_LOOP:

```
add r0, r0, #1
```

```
cmp r0, r1      // r0=1, r1=10
```

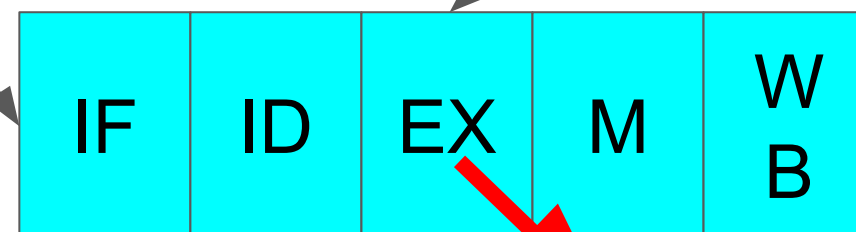
```
bne TAKE_LOOP
```

```
ldr r0, r3      //not taken
```

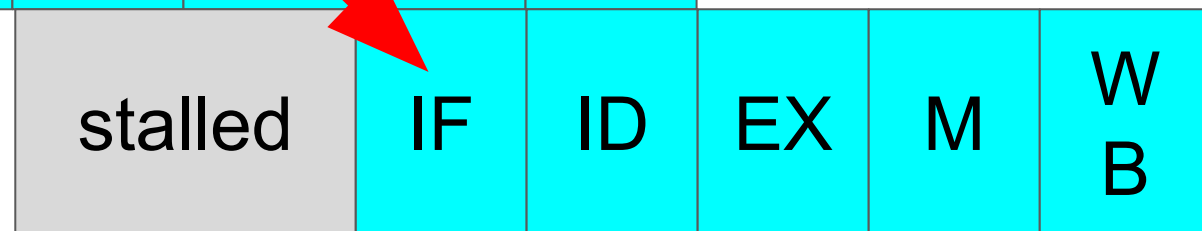
Naively increment
next instruction fetch
($PC = PC + 4$)

Instruction to branch to:
 $PC = PC - 12$

Instr 1: bne TAKE_LOOP



Instr 2: add r0, r0, #1



Correct instruction fetched :)

Control Hazards

Solution 2: *Speculate* - Try to predict branch taken or not taken.

TAKE_LOOP:

```
    add r0, r0, #1
```

```
    cmp r0, r1      // r0=1, r1=10
```

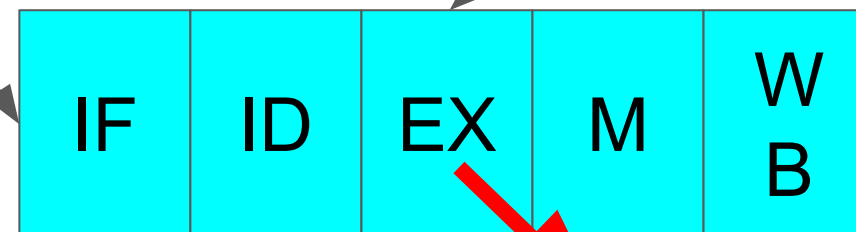
```
    bne TAKE_LOOP
```

```
ldr r0, r3          //not taken
```

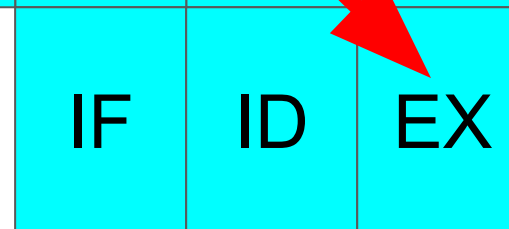
Naively increment
next instruction fetch
($PC = PC + 4$)

Instruction to branch to:
 $PC = PC - 12$

Instr 1: bne TAKE_LOOP



Instr 2: add r0, r0, #1



Predict taken:
 $PC = PC - 12$

IMPORTANT:
CHECK PREDICTION
 $PC \neq PC$

Correct instruction fetched :)

Control Hazards

Solution 2: *Speculate* - Try to predict branch taken or not taken.

TAKE_LOOP:

```
    add r0, r0, #1
```

```
    cmp r0, r1      // r0=1, r1=10
```

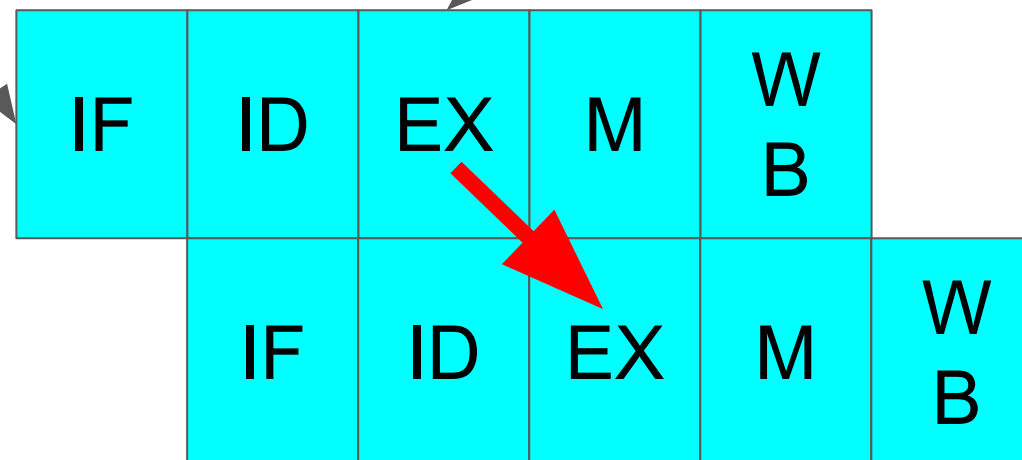
```
    bne TAKE_LOOP
```

```
ldr r0, r3          //not taken
```

Naively increment
next instruction fetch
($PC = PC + 4$)

Instruction to branch to:
 $PC = PC - 12$

Instr 1: bne TAKE_LOOP



Instr 2: add r0, r0, #1

Predict taken:
 $PC = PC - 12$

IMPORTANT:
CHECK PREDICTION
 $PC \neq PC$

Correct instruction fetched :)

Control Hazards

Speculation: What if branch is predicted incorrectly?

TAKE_LOOP:

```
    add r0, r0, #1
```

```
    cmp r0, r1      // r0=10, r1=10
```

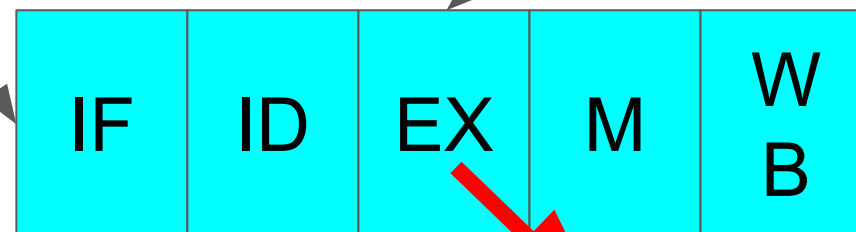
```
    bne TAKE_LOOP
```

```
ldr r0, r3          //not taken
```

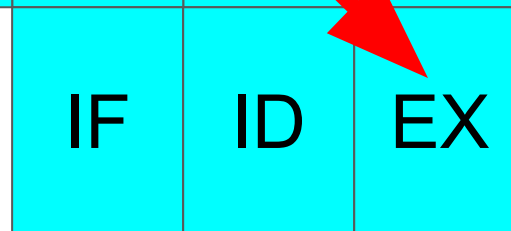
Naively increment
next instruction fetch
($PC = PC + 4$)

Instruction to branch to:
 $PC = PC + 0$

Instr 1: bne TAKE_LOOP



Instr 2: add r0, r0, #1



Predict taken:
 $PC = PC - 12$

IMPORTANT:
CHECK PREDICTION
 $PC \neq PC$

INCORRECT Instruction

Control Hazards

Speculation: What if branch is predicted incorrectly?

TAKE_LOOP:

add r0, r0, #1

cmp r0, r1 // r0=10, r1=10

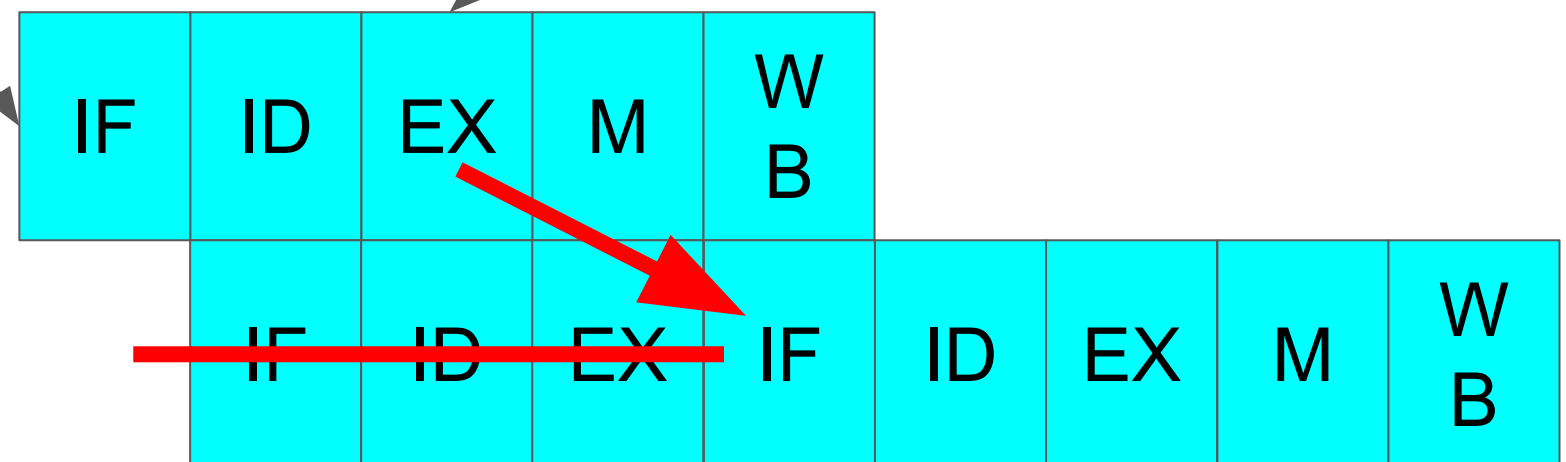
bne TAKE_LOOP

ldr r0, r3 //not taken

Naively increment
next instruction fetch
(PC = PC + 4)

Instruction to branch to:
PC = PC + 0

Instr 1: bne TAKE_LOOP



SQUASH

~~Instr 2: add r0, r0, #1~~

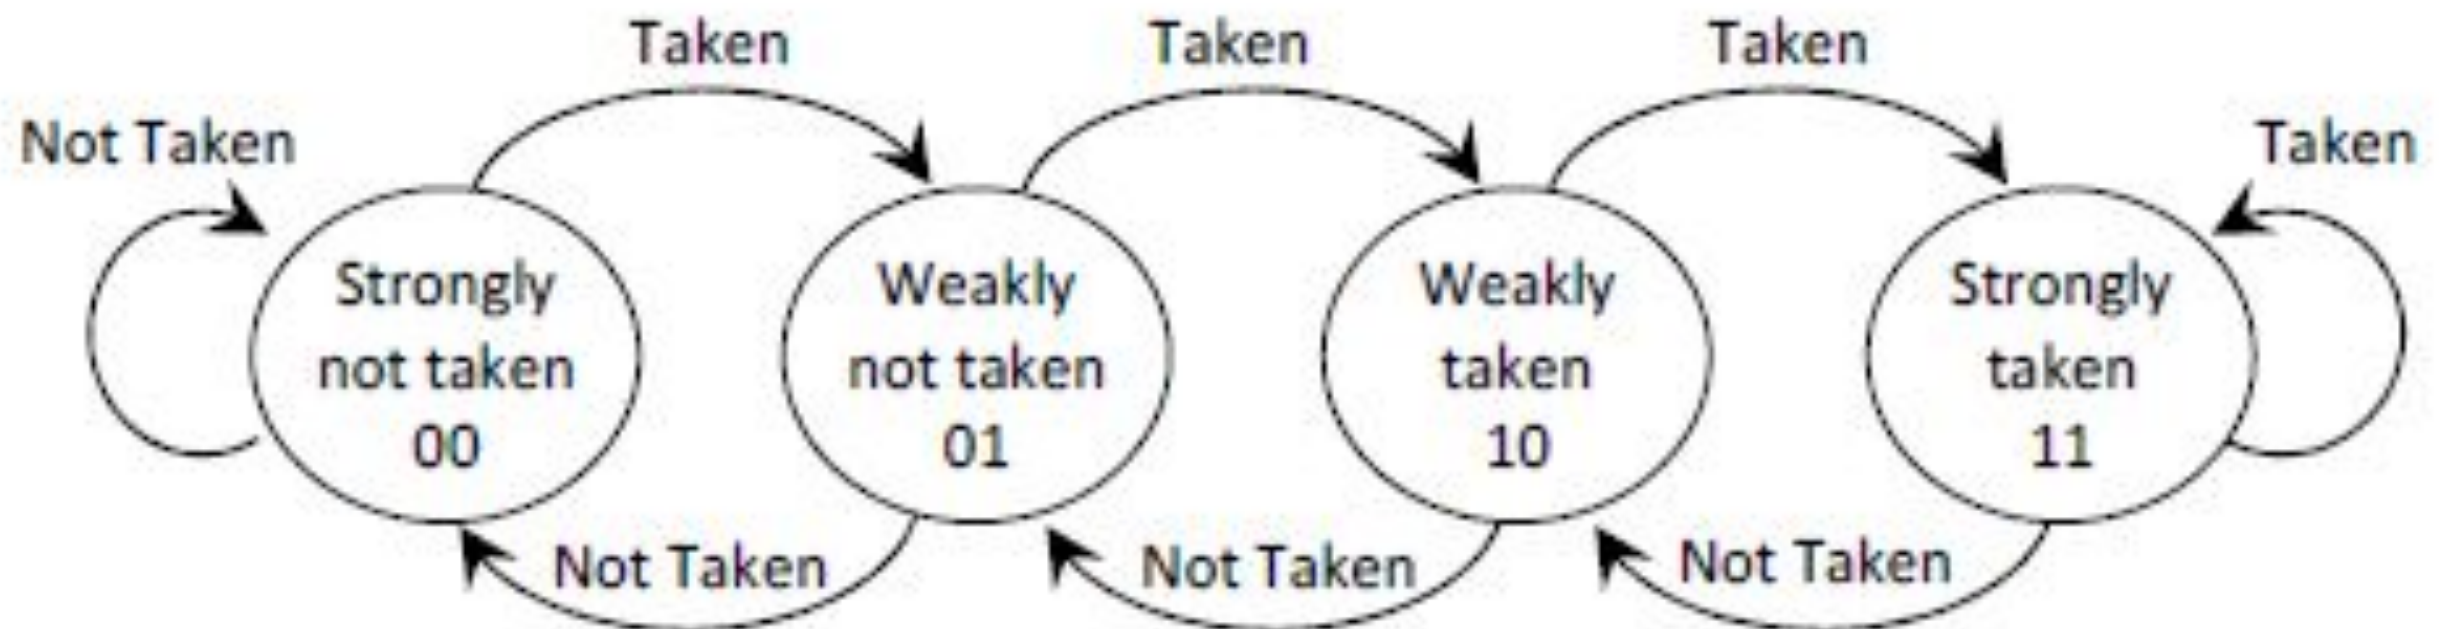
Instr 2: ldr r0, r3

Branch Prediction

Piece of *hardware* to predict if branch will be *taken* or *not taken*.

Simplest: Always predict TAKEN

2-bit Predictor:



Demo: clear

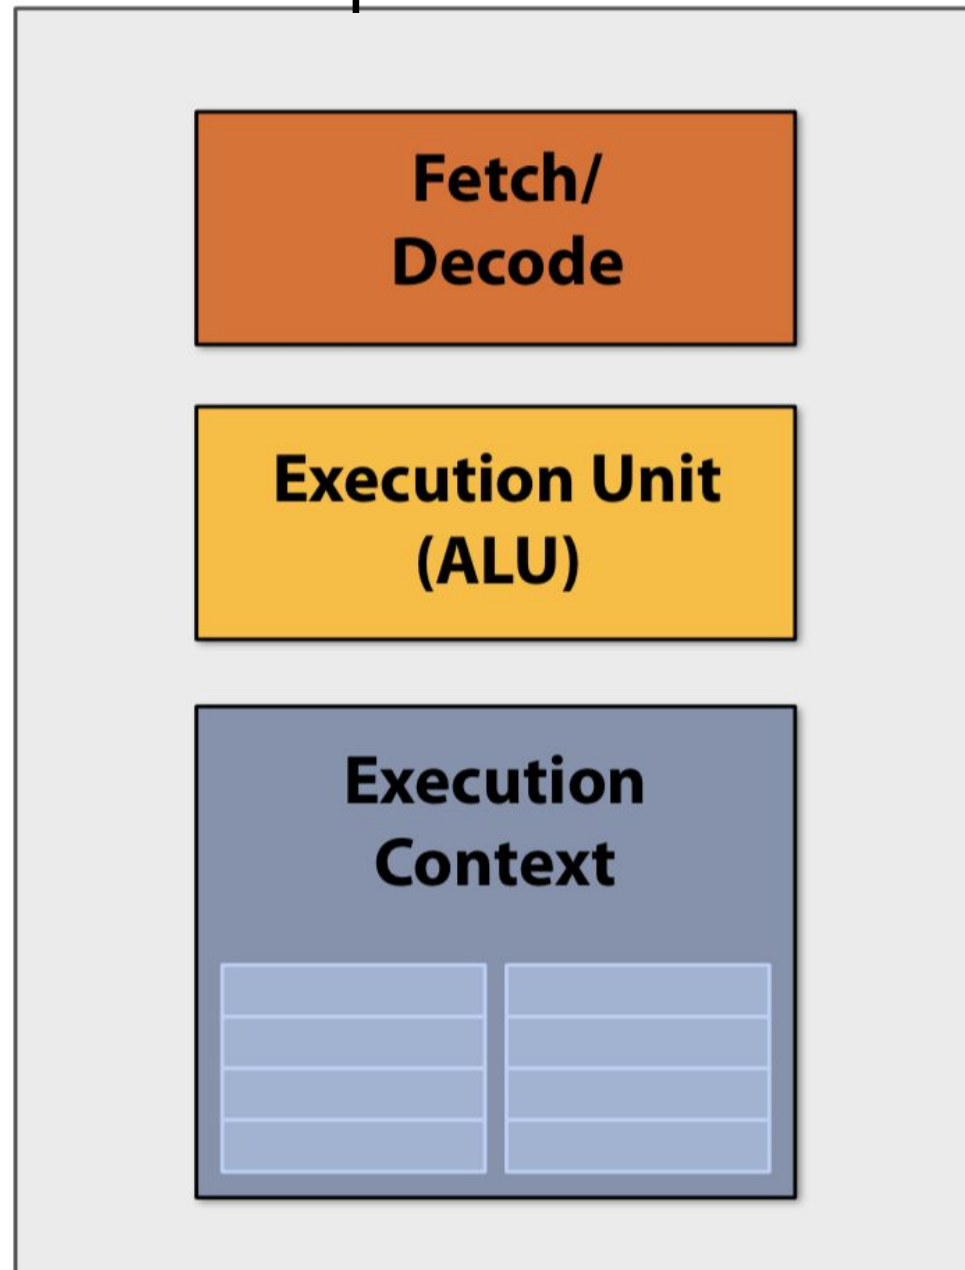
w/ branch prediction

Parallel Architecture

Images courtesy of CS 149 (cs149.stanford.edu)

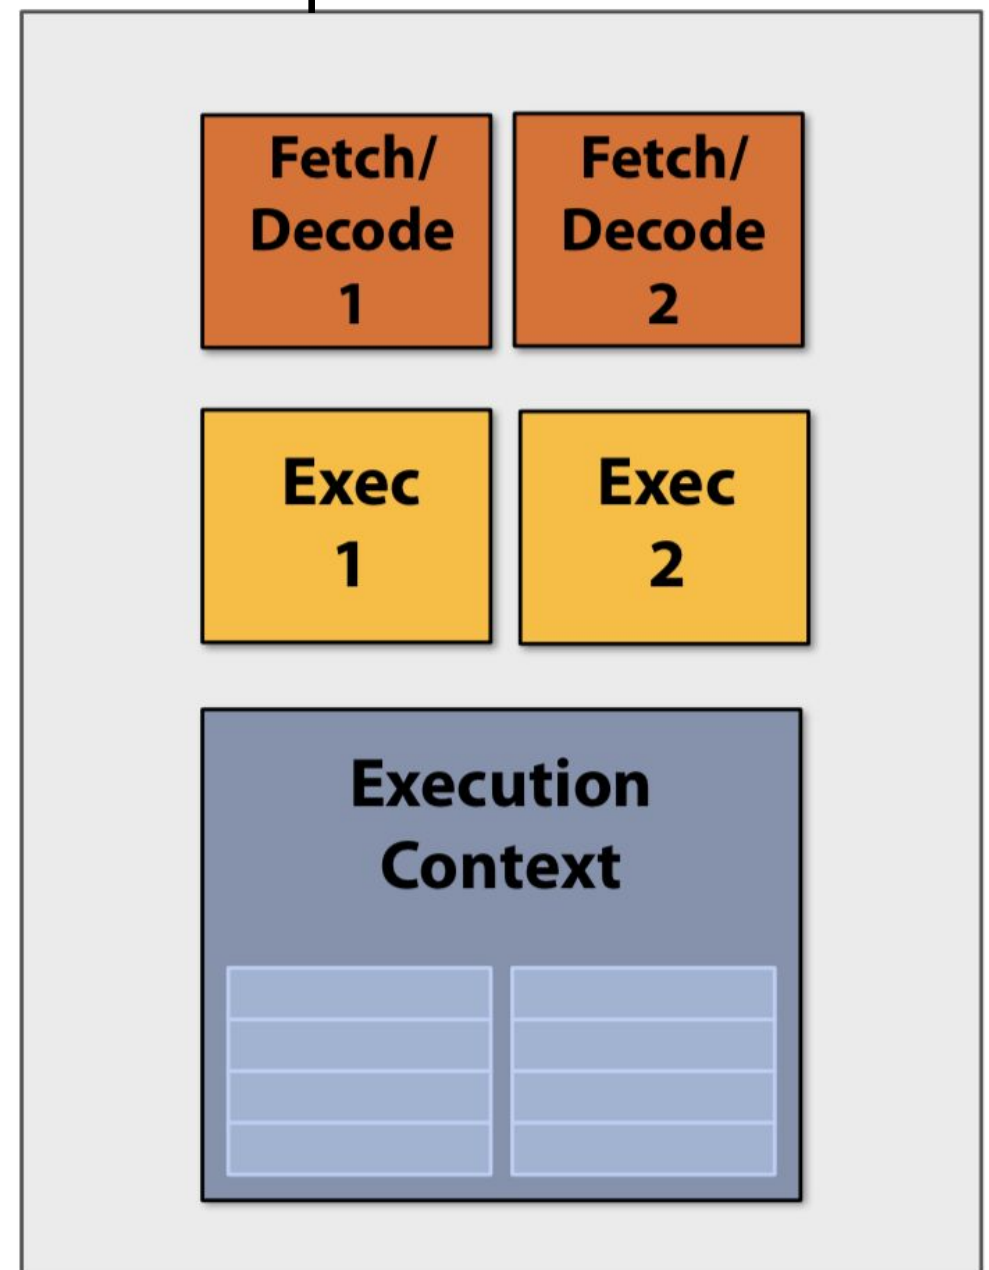
Instruction Level Parallelism (ILP)

Pipelined CPU



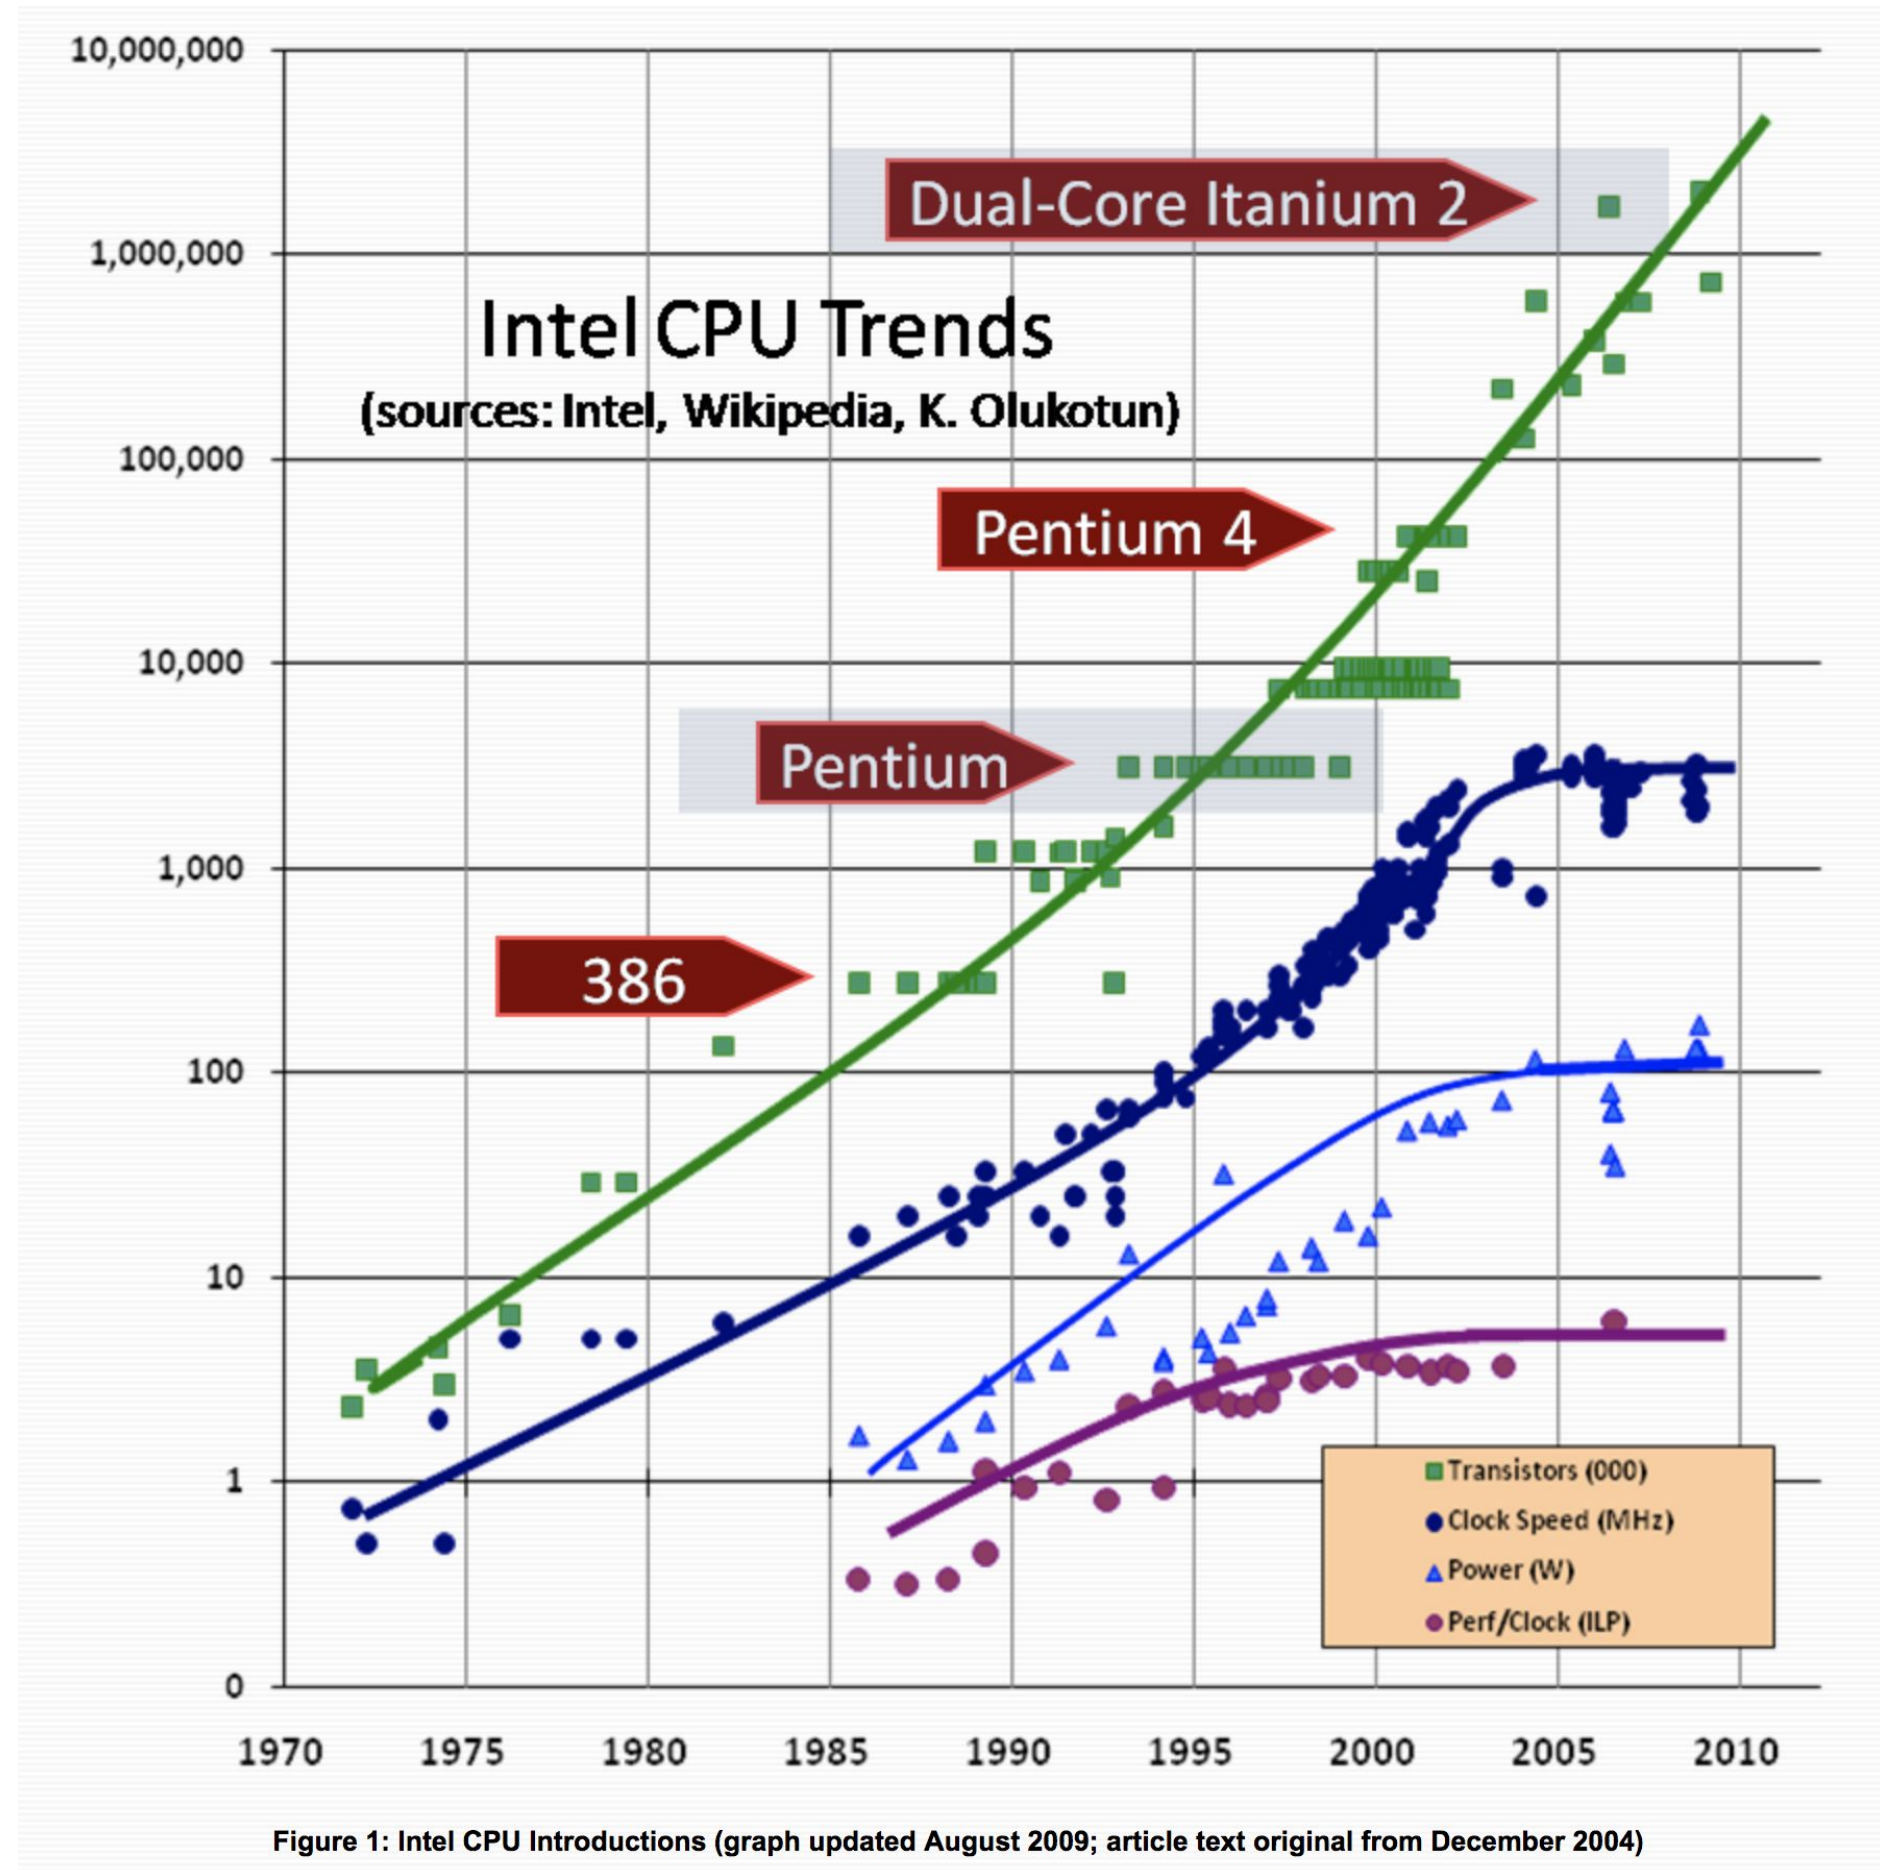
1 instr/cycle

Superscalar CPU

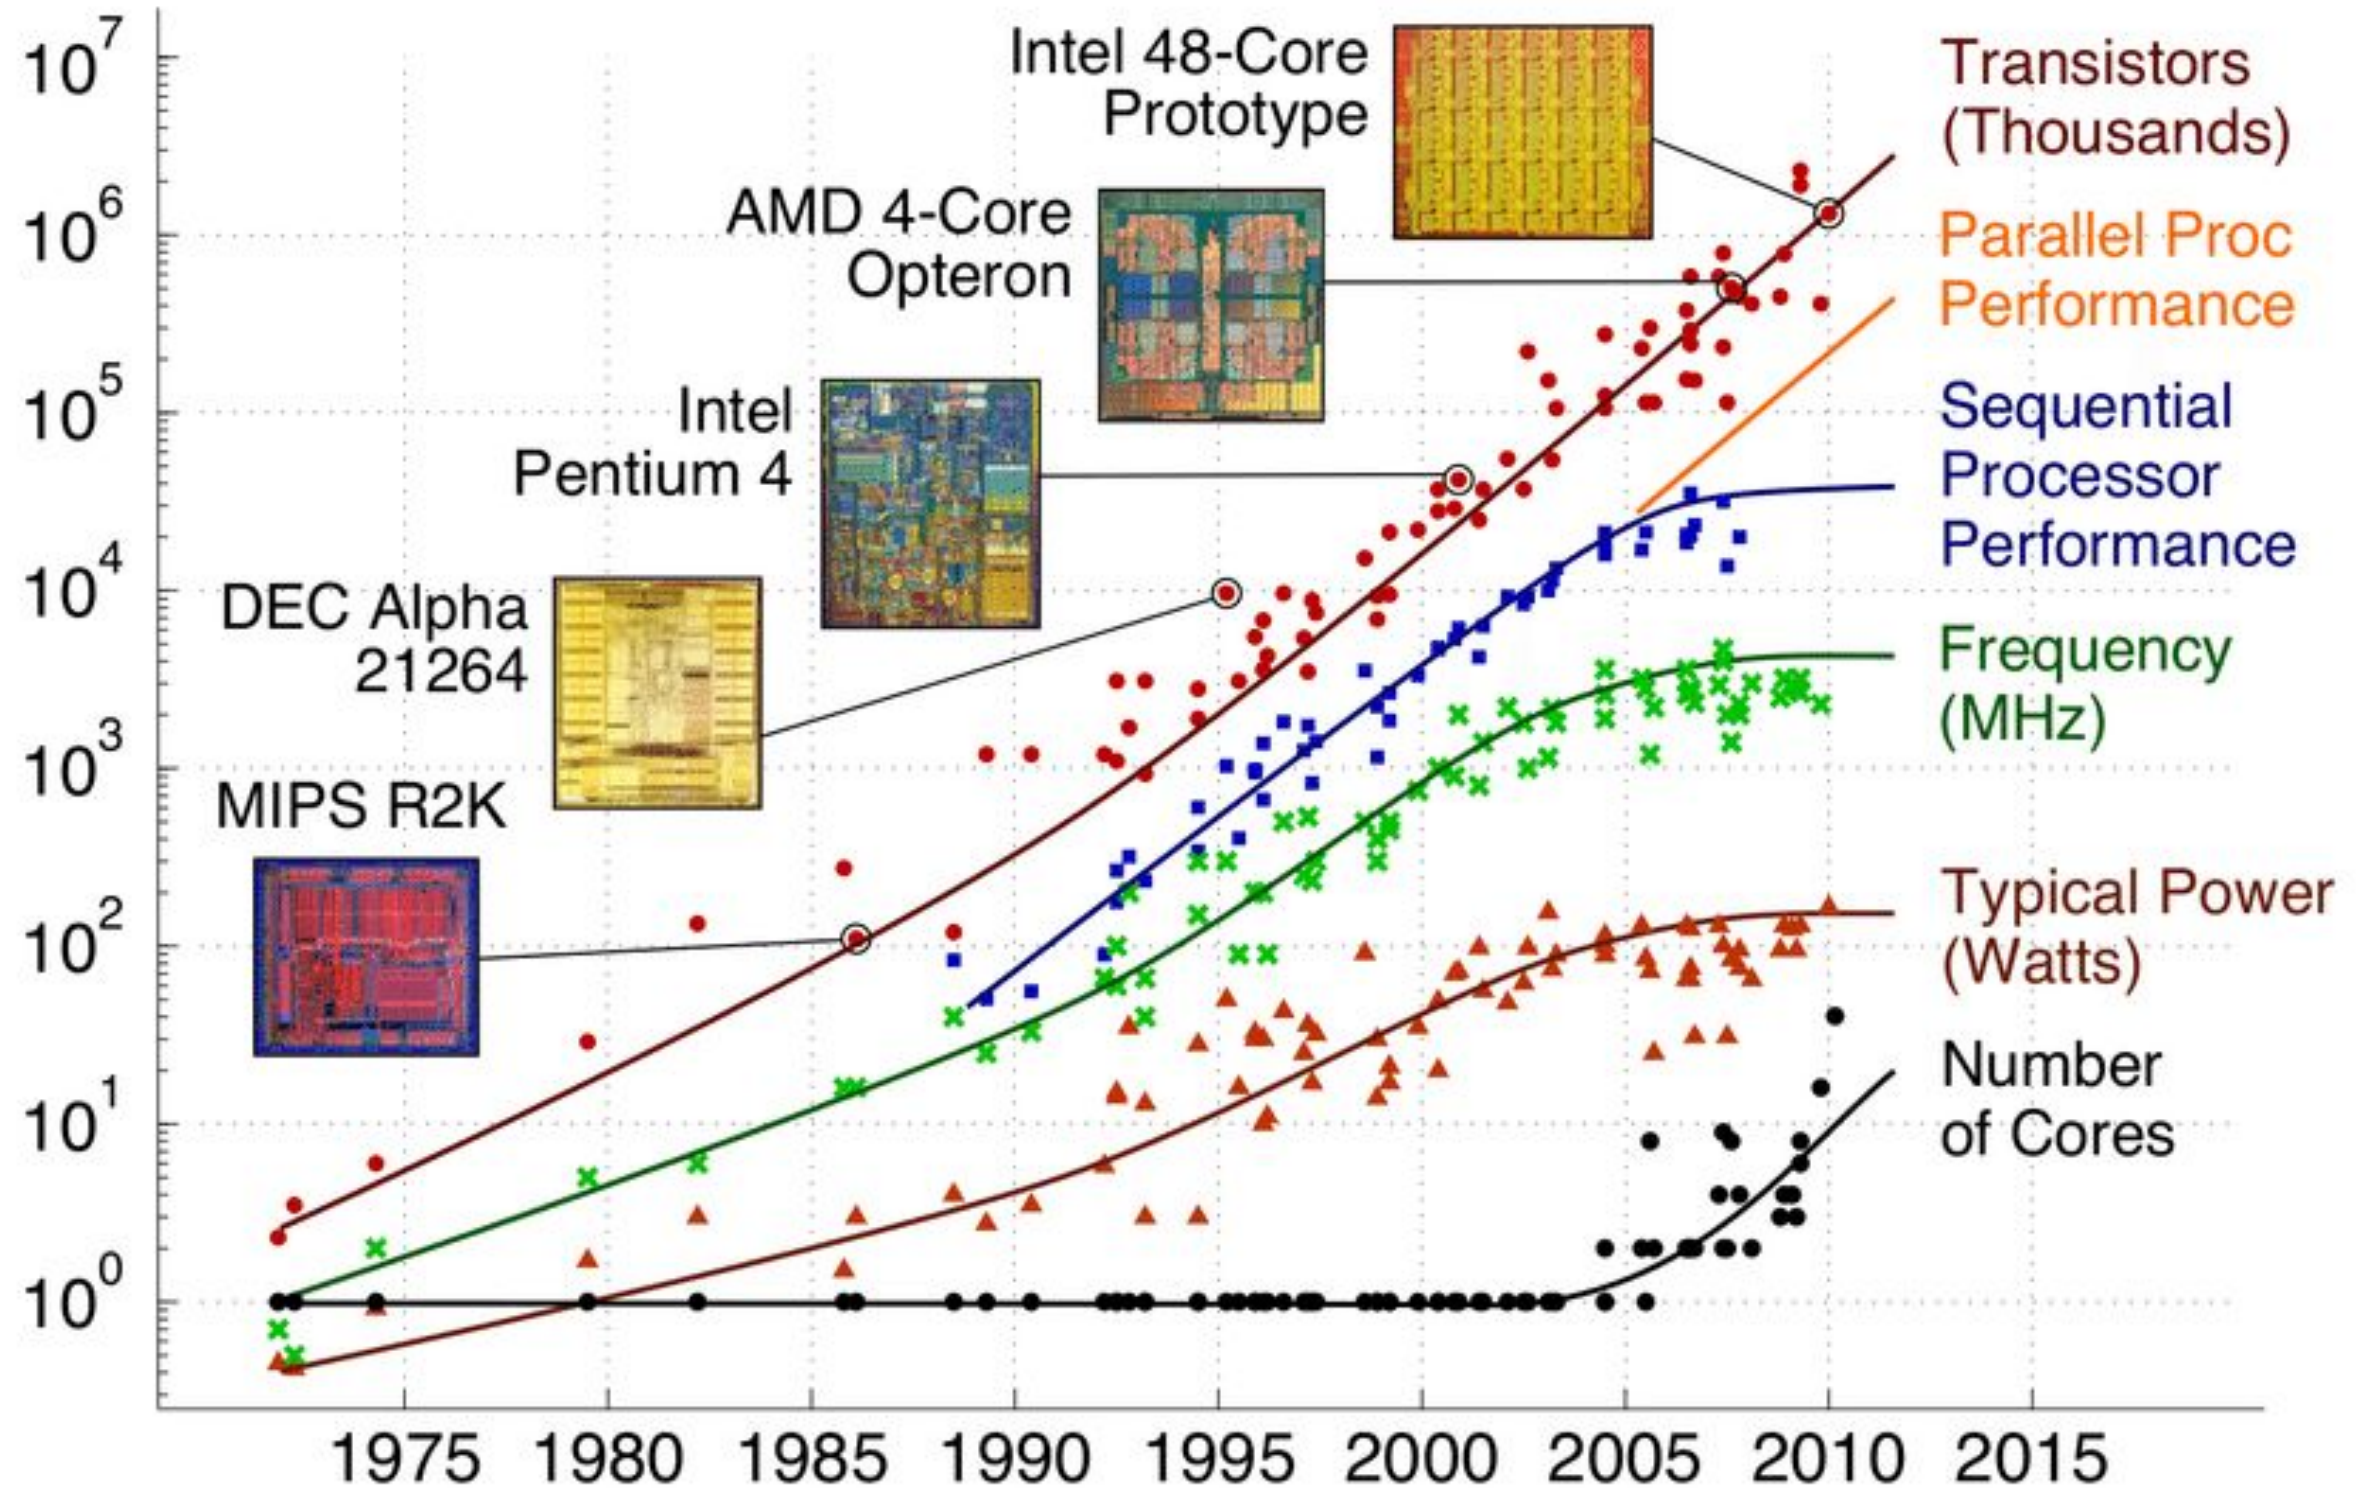


2 instr/cycle

Limits of ILP



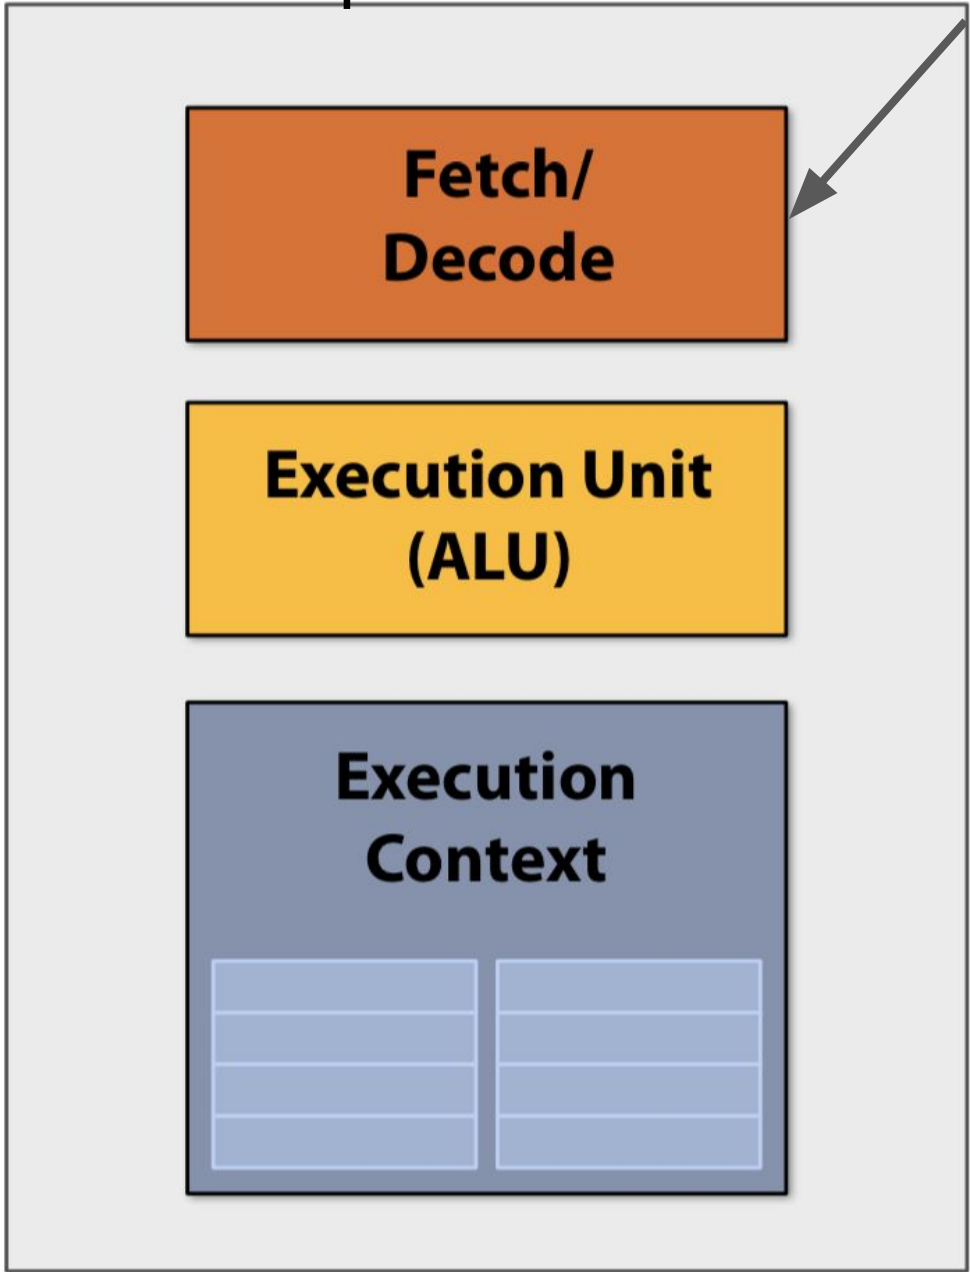
Why Parallelism?



Data Level Parallelism (DLP)

Single Instruction, Multiple Data
SIMD CPU

Pipelined CPU

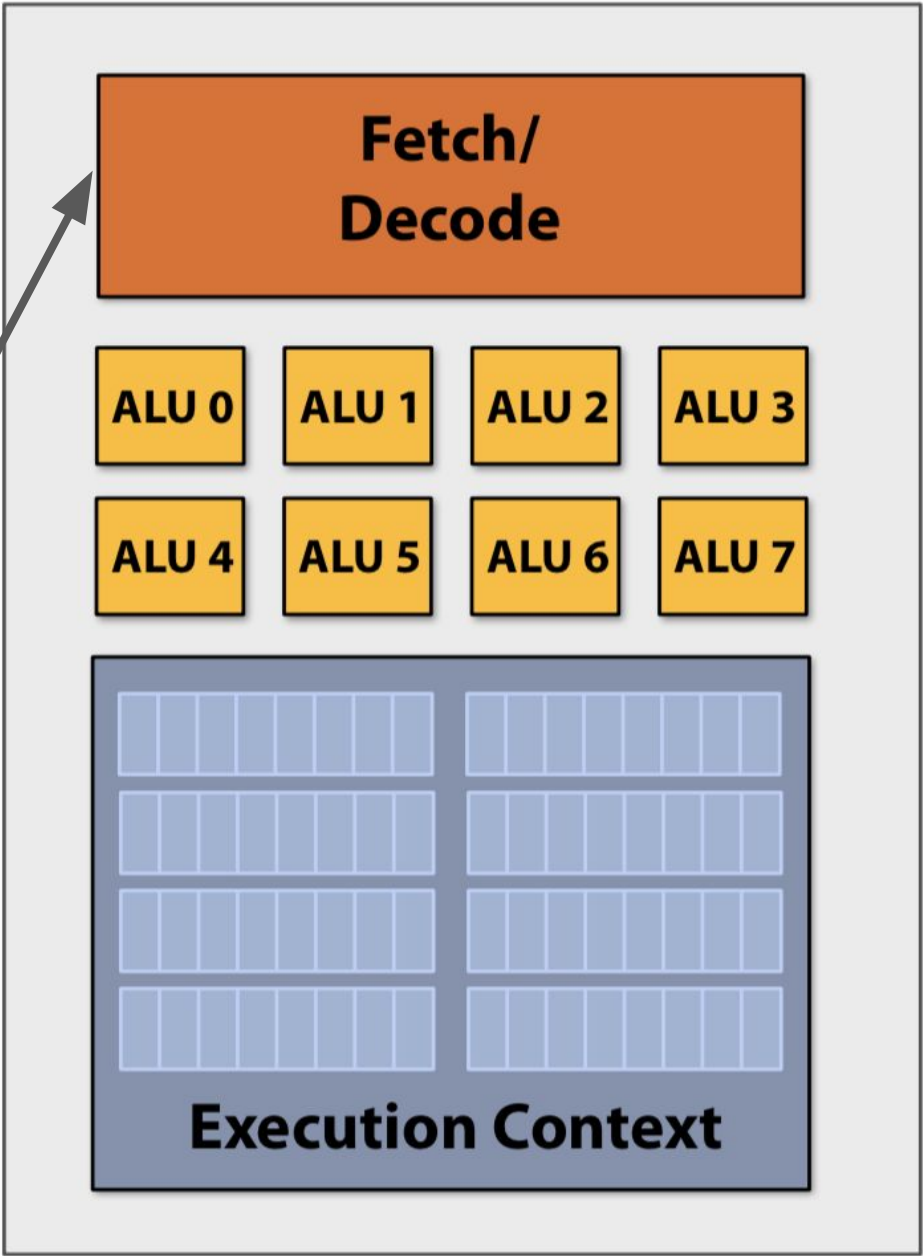


```
lddr r0, r1
lddr r2, r3
add r0, r2
```

```
vldr v0, v1
vldr v2, v2
vadd v0, v2
```

1 instr/cycle

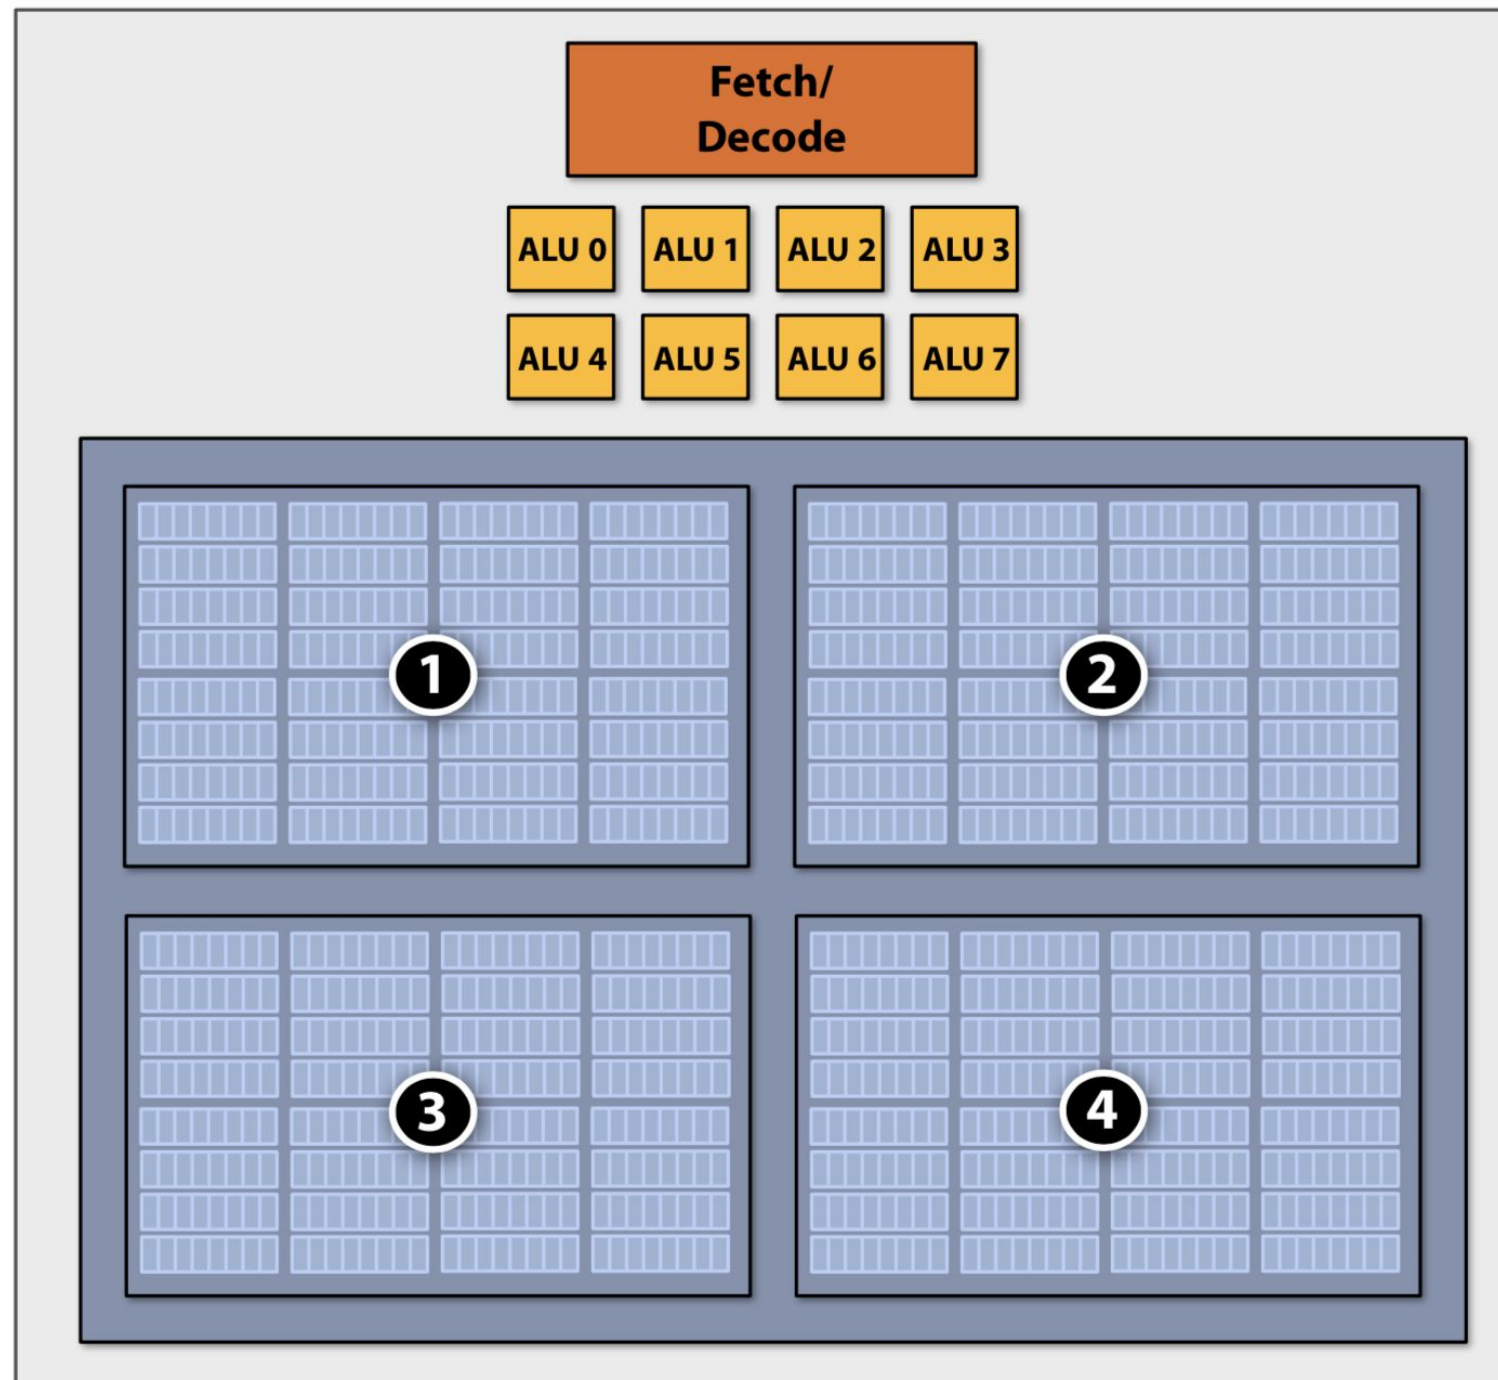
SIMD CPU



1 instr/cycle
8 computations/cycle

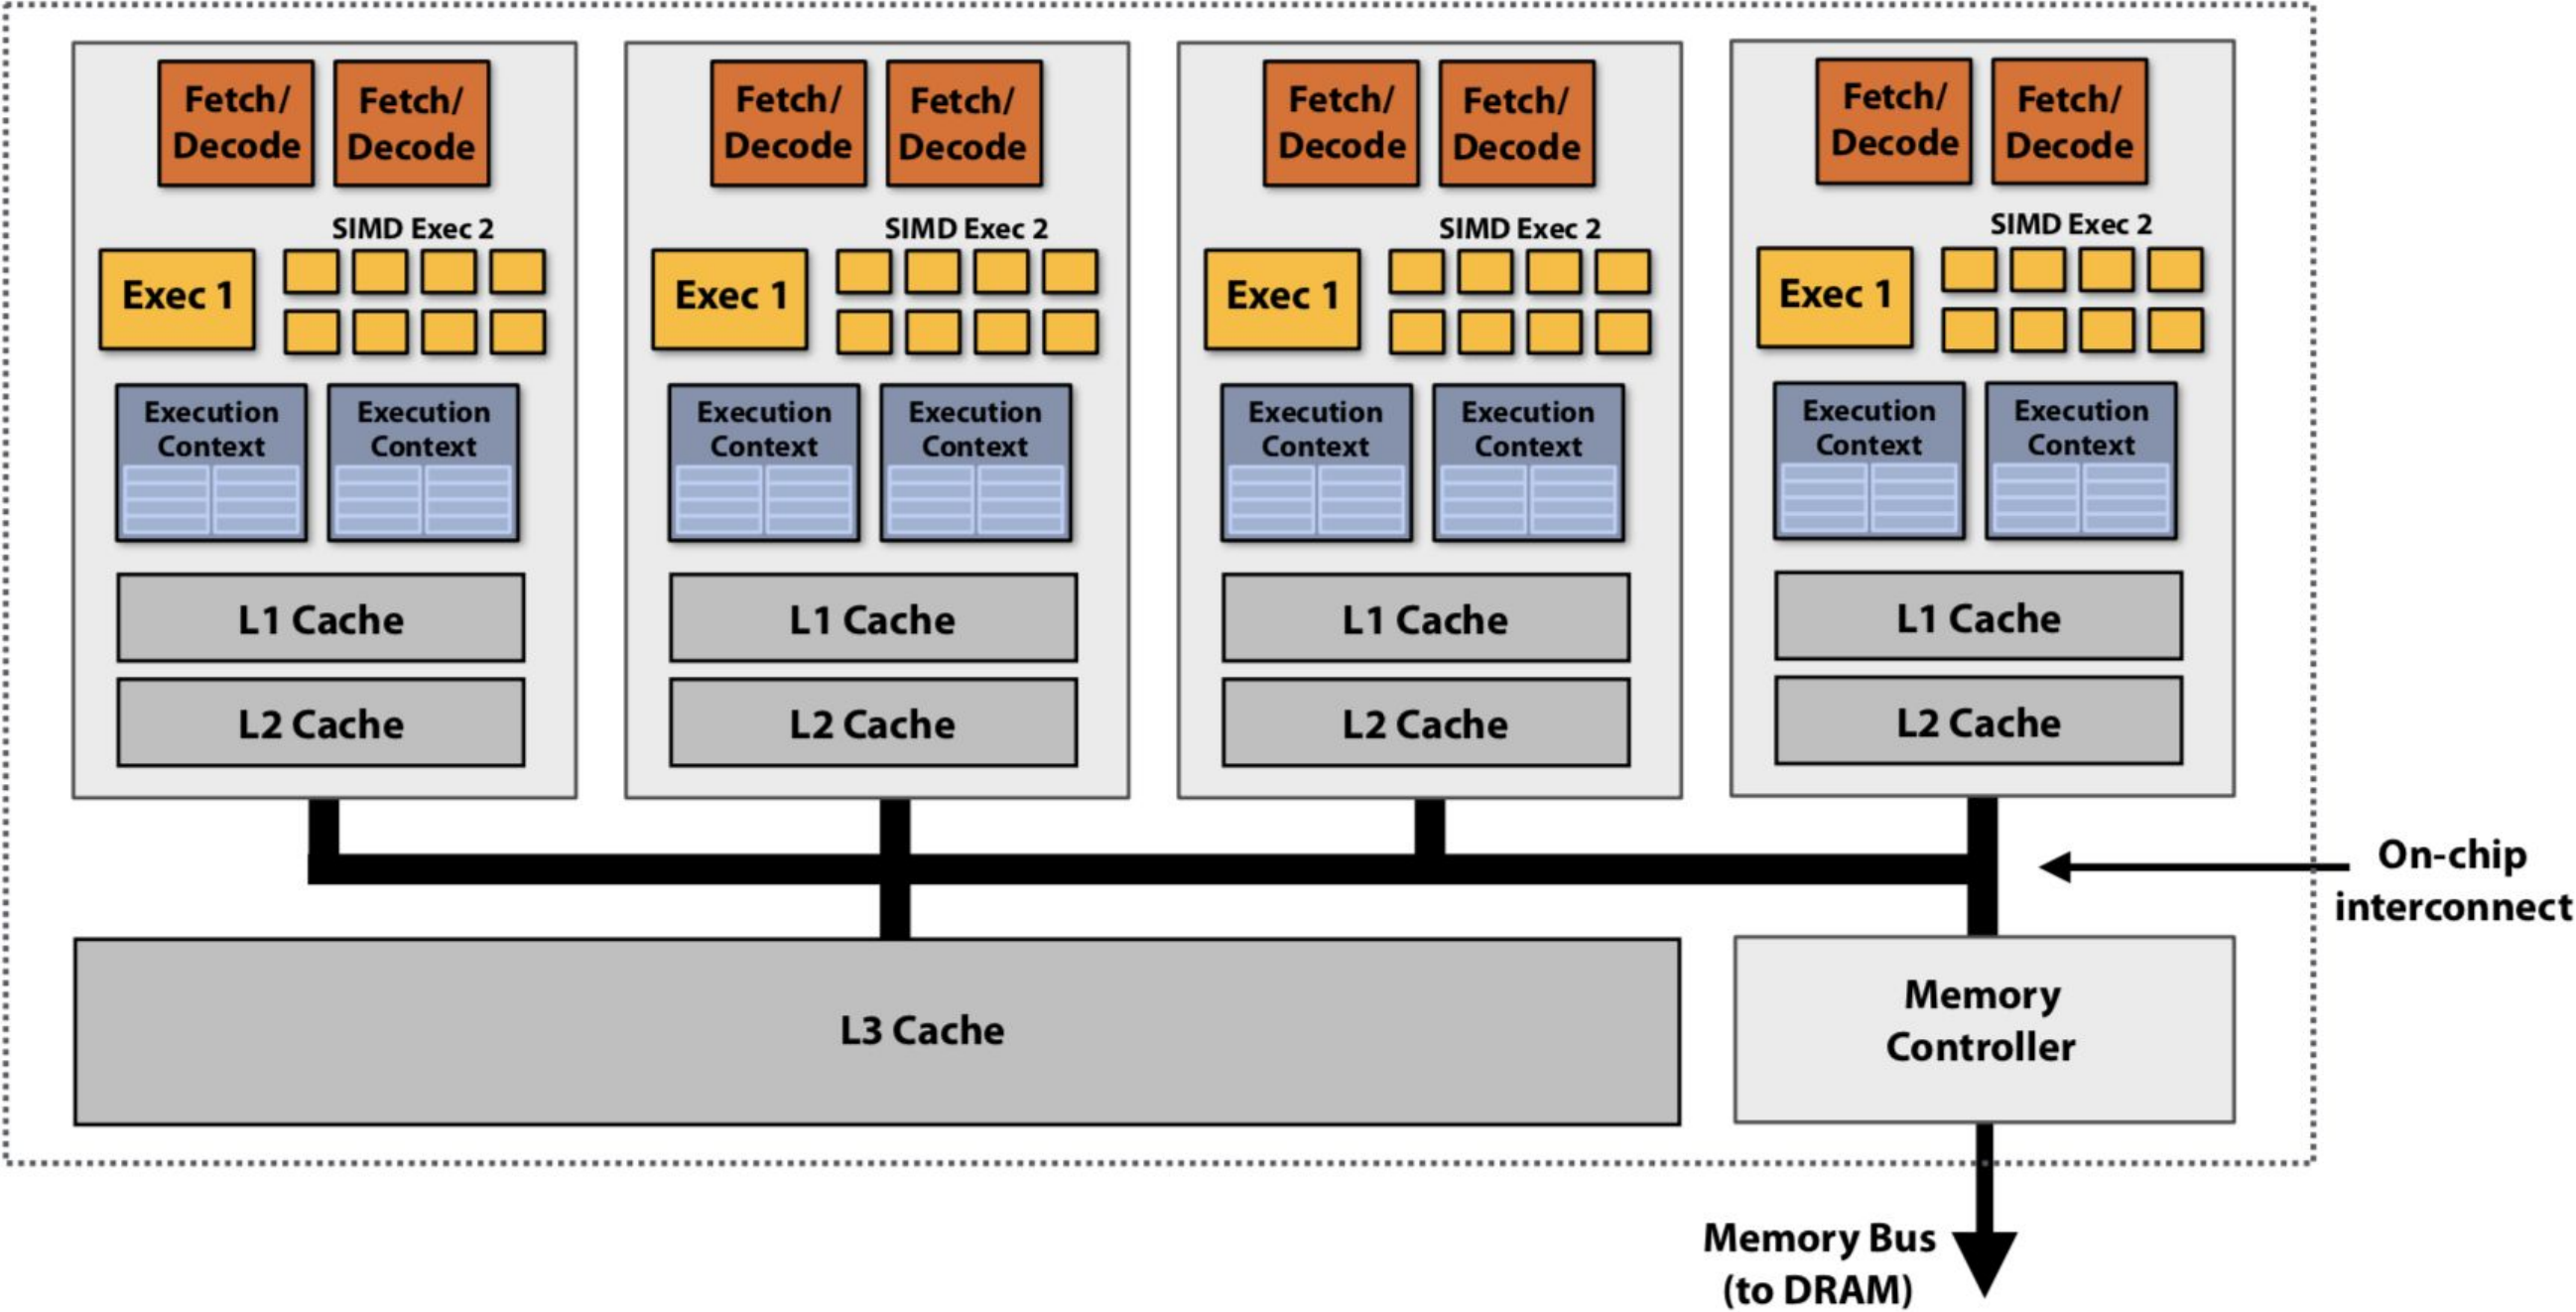
Thread Level Parallelism (TLP)

Multiple Execution Contexts = Hardware Multithreading

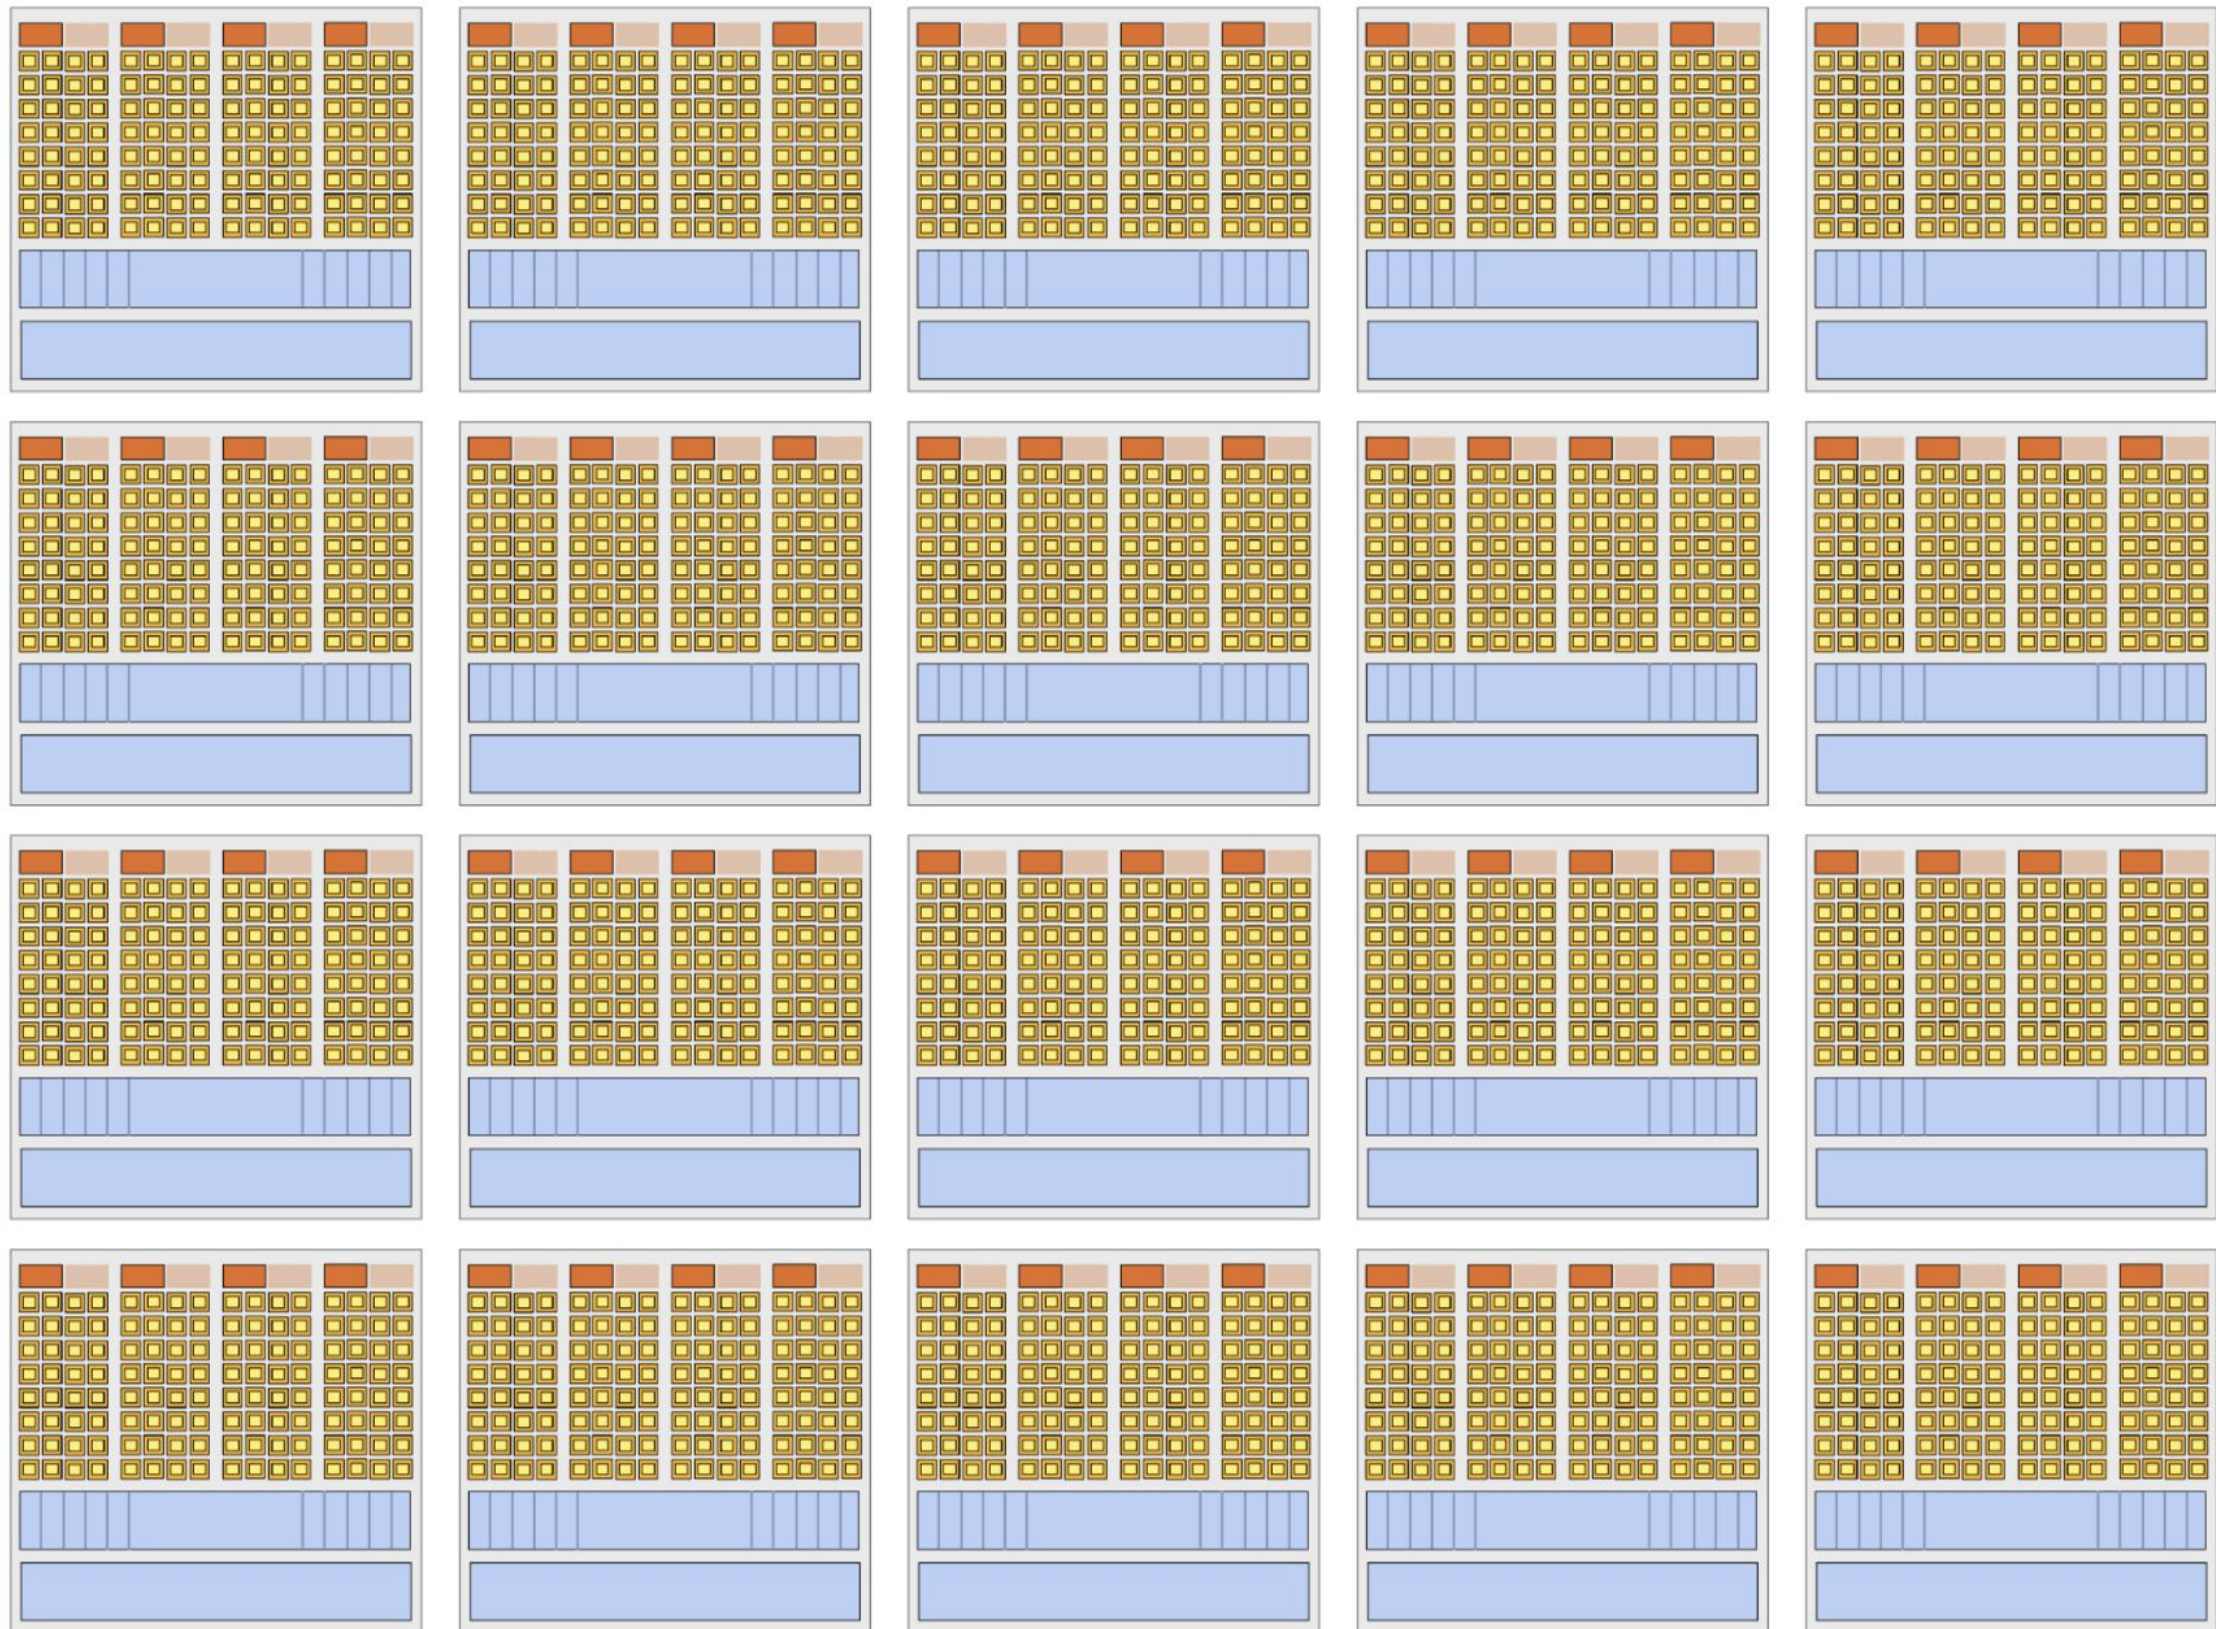


Multiple programs/threads can run together concurrently, but still have to share pipeline resources.

Modern Quad-Core CPU



Modern GPU (NVIDIA GTX 1080)



$20 \text{ cores} * 4 \text{ pipelines} * 32 \text{ element SIMD ALUs} = \mathbf{2,560 \text{ parallel computations}}$
 $* 2 \text{ FLOPs (mult-add)} * 1.6 \text{ Ghz} = \mathbf{8.1 \text{ TeraFLOPS}}$

Architecture Summary

Hardware makes optimizations for instruction latency & throughput.

Memory latency is very significant.

Single-thread performance is not increasing drastically on chips anymore.

Multicore chips and application-specific processors are rising.

Software/algorithm needs to match the architecture (or vice versa).